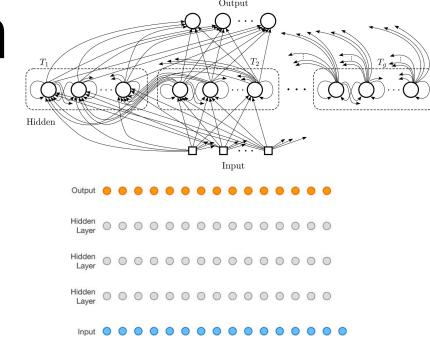


Recurrent Nets and Sequence Generation

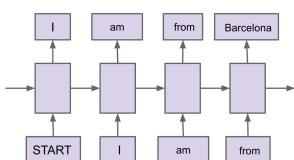
- Lecture topics:
 - Recurrent Neural Networks
 - Long-Short Term Memory (LSTM)
 - (Conditional) Sequence Generation

- Guest Lecturer: **Oriol Vinyals**

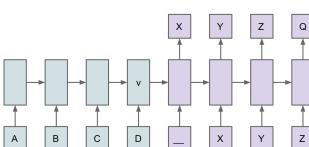
- Joined DeepMind in 2016.
- Worked in Google Brain from 2013 to 2016.
- PhD in Artificial Intelligence from UC Berkeley (2009-13). Supervisor: Darrell / Morgan.
- Current research topics: deep learning, sequence modeling, generative models, distillation, RL/Starcraft, one shot learning.



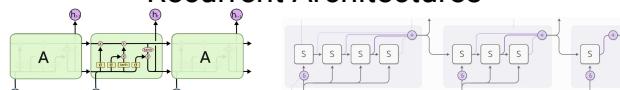
Sequence Prediction



Seq2Seq



Recurrent Architectures



Sequences and Recurrent Neural Nets

Oriol Vinyals // @OriolVinyalsML
Research Scientist
DeepMind
UCL, Feb 2018

Roadmap

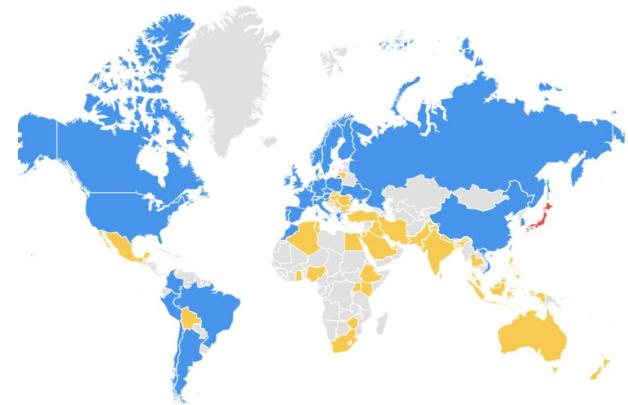
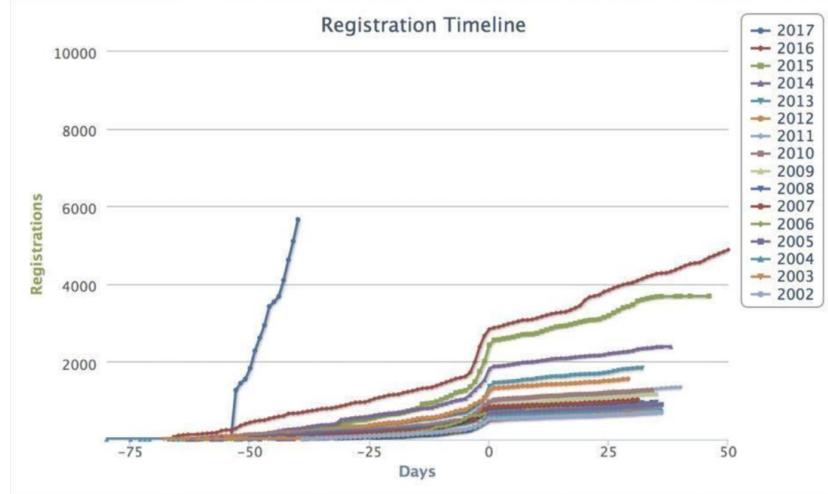
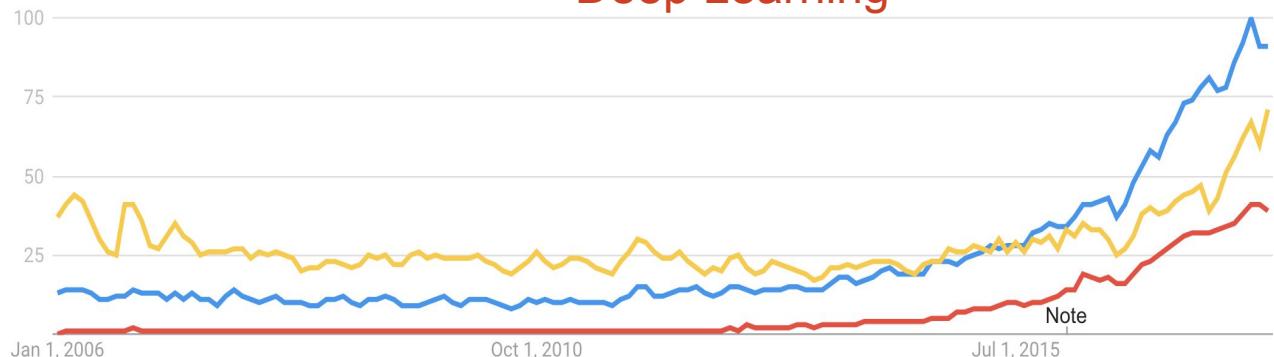
- I. Deep Learning in (almost) one slide
- II. Why should I care about Sequences?
- III. Recurrent Neural Networks Fundamentals
 - A. The chain rule (loss)
 - B. Vanishing gradients (optimization)
 - C. LSTMs / GRUs (architecture)
- IV. Recurrent Networks as Sequence Decoders
 - A. Sequence-To-Sequence Learning & Applications
- V. Beyond RNNs

Part I: Deep Learning Overview

Lots of attention

- Startups / VCs
- Big companies
- Open source efforts
- Press (noise...)
- Universities

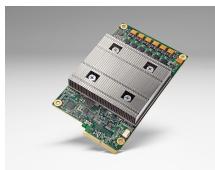
Machine Learning
Artificial Intelligence
Deep Learning



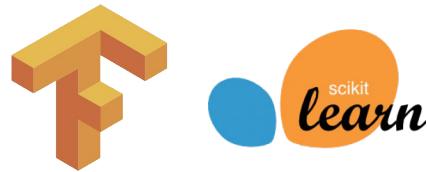
Deep Learning: Zooming Out



Platforms



Frameworks



Caffe K Keras

Microsoft CNTK mxnet theano

PYTORCH Caffe²



Datasets

0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9



Caltech 101 MGENET



WMT Workshop 2014

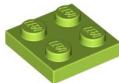
Deep Learning: Zooming In

Non-Linearity



Relu
Sigmoid
Tanh
GRU
LSTM
Linear
...

Connectivity Pattern



Fully connected
Convolutional
Dilated
Recurrent
Recursive
Skip / Residual
Random
...

Optimizer



SGD
Momentum
RMSProp
Adagrad
Adam
Second Order (KFac)
...

Loss



Cross Entropy
Adversarial
Variational
Max. Likelihood
Sparse
L2 Reg
REINFORCE
...

Hyper Parameters

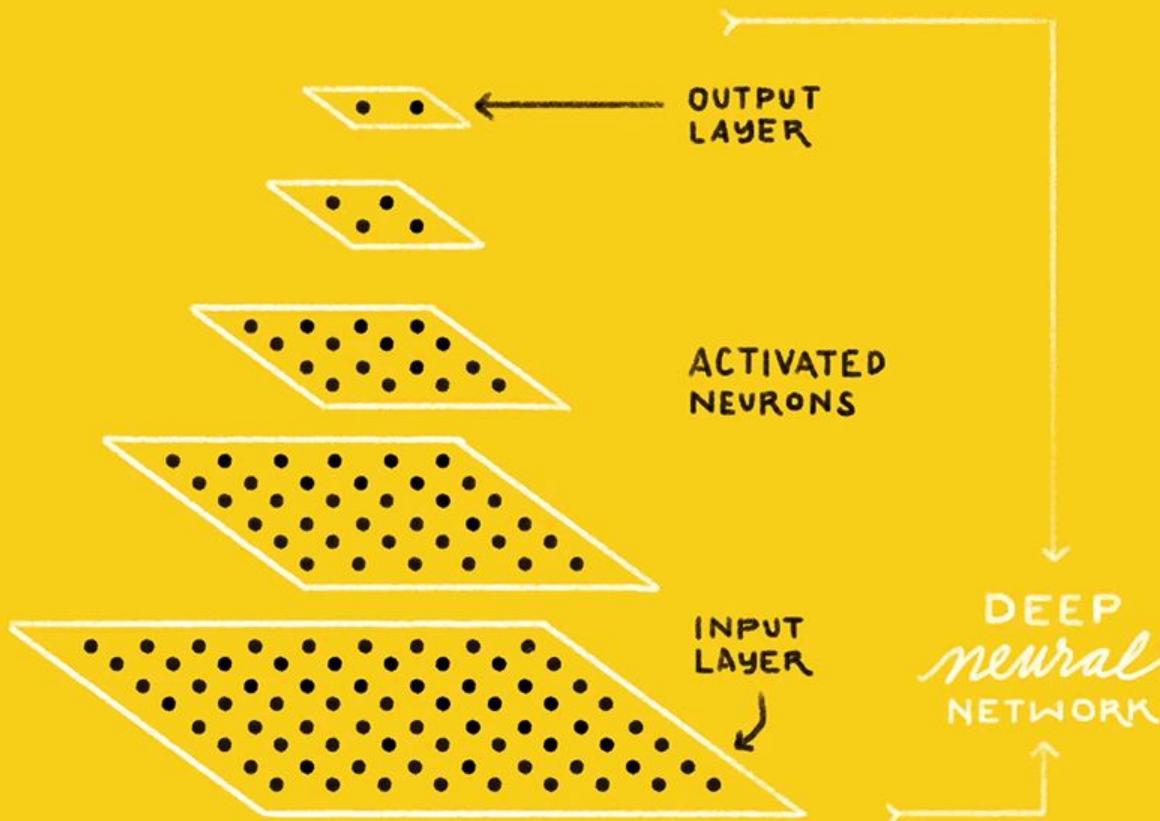
Learning Rate
Decay
Layer Size
Batch Size
Dropout Rate
Weight init
Data augmentation
Gradient Clipping
Beta
Momentum



IS THIS A
CAT or DOG?



CAT DOG



A simple recipe

1. Lots of data

A simple recipe

1. Lots of data
2. Deep, BIG DNNs/CNNs/RNNs

A simple recipe

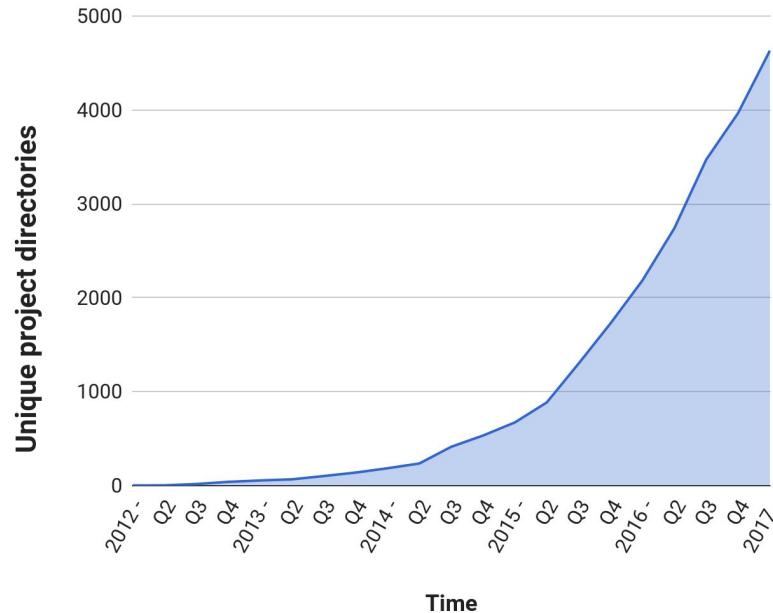
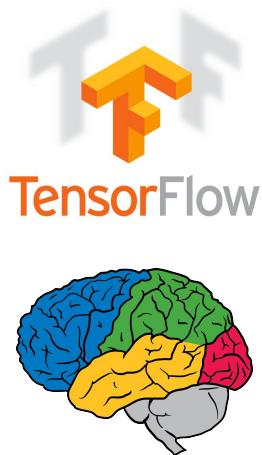
1. Lots of data
2. Deep, BIG DNNs/CNNs/RNNs
3. Lots of GPUs

A simple recipe

1. Lots of data
2. Deep, BIG DNNs/CNNs/RNNs
3. Lots of GPUs
4. PROFIT!

Growing Use of Deep Learning at Google

“Google will soon be a big LSTM”, Jürgen Schmidhuber

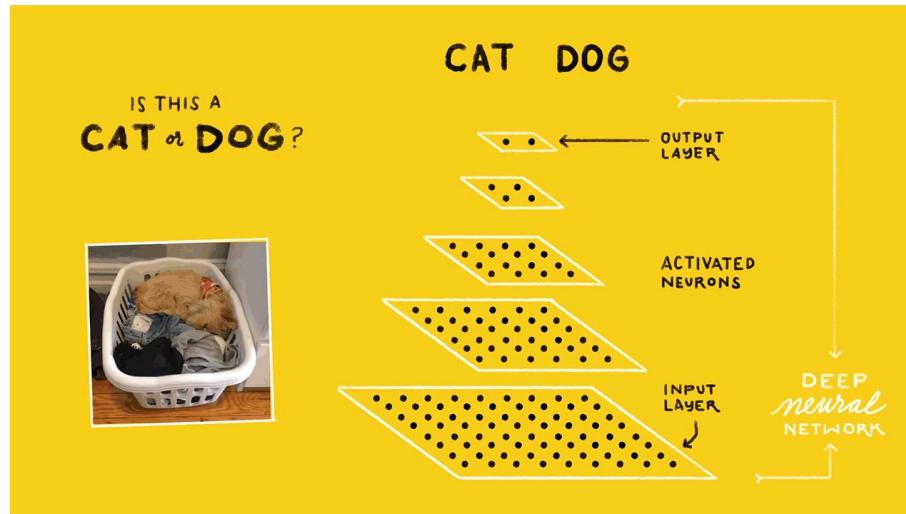


Across many products/areas:

- Android
- Apps
- drug discovery
- Gmail
- Image understanding
- Maps
- Photos
- Robotics research
- Search
- Speech
- Translate
- YouTube
- ... many others ...

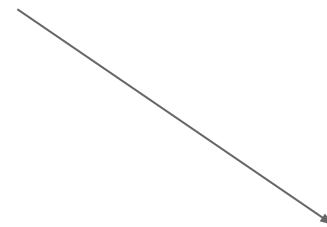
Are we done yet?

- Neural Nets struggle with sequences
- Inputs and outputs must be fixed-sized vectors



Part II: Why Sequences?

[Part, II:; Why, Sequences?]



Part II: Why Sequences?

[Part, II:, Why, Sequences?]

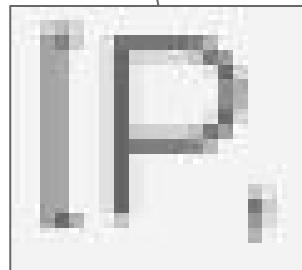
[P, a, r, t]

Part II: Why Sequences?

[Part, II; Why, Sequences?]

[P, a, r, t]

Part II: Why Sequences?



Part II: Why Sequences?

[Part, II:, Why, Sequences?]

Part II: Why Sequences?

[Part, II:, Why, Sequences?]

[P, a, r, t]

Part II: Why Sequences?

[Part, II:, Why, Sequences?]

[P, a, r, t]

Part II: Why Sequences?

[P,

Sequences

- Words, Letters

50 years ago, the fathers of artificial intelligence convinced everybody that logic was the key to intelligence. Somehow we had to get computers to do logical reasoning. The alternative approach, which they thought was crazy, was to forget logic and try and understand how networks of brain cells learn things. Curiously, two people who rejected the logic based approach to AI were Turing and Von Neumann. If either of them had lived I think things would have turned out differently... now neural networks are everywhere and the crazy approach is winning.

Geoff Hinton

- Speech



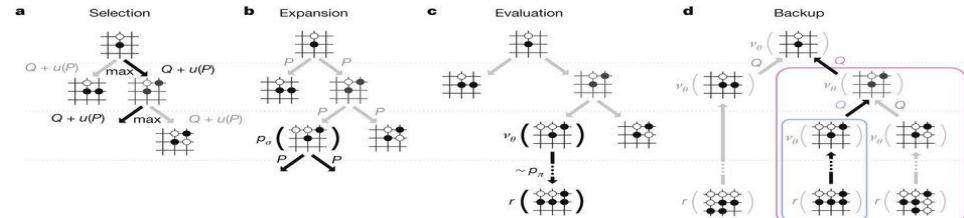
- Images, Videos, Touch



- Programs

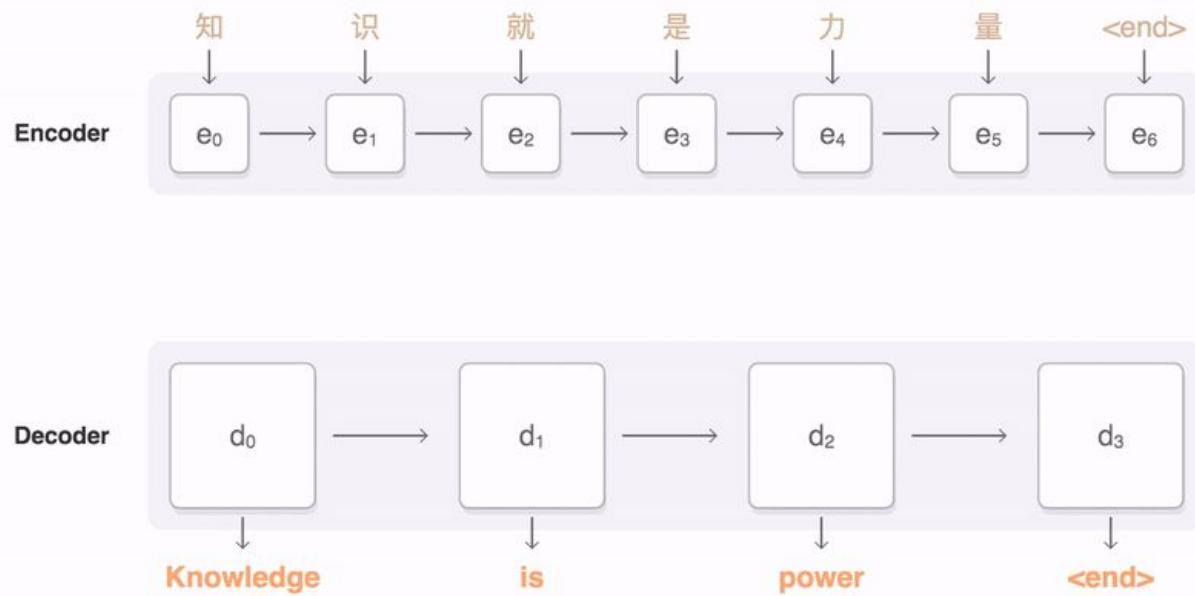
```
while (*d++ = *s++);
```

- Sequential Decision Making (RL)



Mix and Match: Google Neural MT

Wu et al, 2016 (Kalchbrenner et al, 2013; Sutskever et al, 2014;
Cho et al, 2014; Bahdanau et al, 2014; ...)



Math / Code [Karpathy, 2015]

Proof. Omitted. \square

Lemma 0.1. Let \mathcal{C} be a set of the construction.

Let \mathcal{C} be a gerber covering. Let \mathcal{F} be a quasi-coherent sheaves of \mathcal{O} -modules. We have to show that

$$\mathcal{O}_{\mathcal{O}_X} = \mathcal{O}_X(\mathcal{L})$$

Proof. This is an algebraic space with the composition of sheaves \mathcal{F} on $X_{\text{étale}}$ we have

$$\mathcal{O}_X(\mathcal{F}) = \{\text{morph}_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F})\}$$

where \mathcal{G} defines an isomorphism $\mathcal{F} \rightarrow \mathcal{F}$ of \mathcal{O} -modules. \square

Lemma 0.2. This is an integer \mathcal{Z} is injective.

Proof. See Spaces, Lemma ??.

Lemma 0.3. Let S be a scheme. Let X be a scheme and X is an affine open covering. Let $\mathcal{U} \subset \mathcal{X}$ be a canonical and locally of finite type. Let X be a scheme. Let X be a scheme which is equal to the formal complex.

The following to the construction of the lemma follows.

Let X be a scheme. Let X be a scheme covering. Let

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over S and Y .

Proof. Let X be a nonzero scheme of X . Let X be an algebraic space. Let \mathcal{F} be a quasi-coherent sheaf of \mathcal{O}_X -modules. The following are equivalent

- (1) \mathcal{F} is an algebraic space over S .
- (2) If X is an affine open covering.

Consider a common structure on X and X the functor $\mathcal{O}_X(U)$ which is locally of finite type. \square

```
static void num_serial_settings(struct tty_struct *tty)
{
    if (tty == tty)
        disable_single_st_p(dev);
    pci_disable_spool(port);
    return 0;
}

static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000fffff8) & 0x0000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &offset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

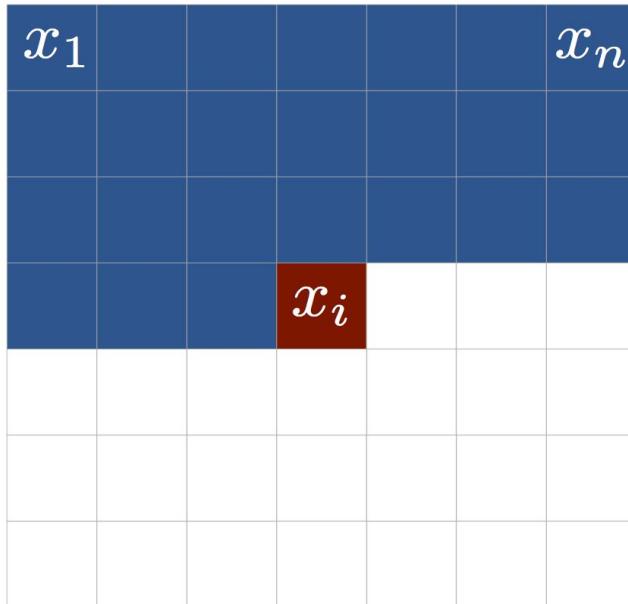


Human: A young girl asleep on the sofa cuddling a stuffed bear.

NIC: A close up of a child holding a stuffed animal.

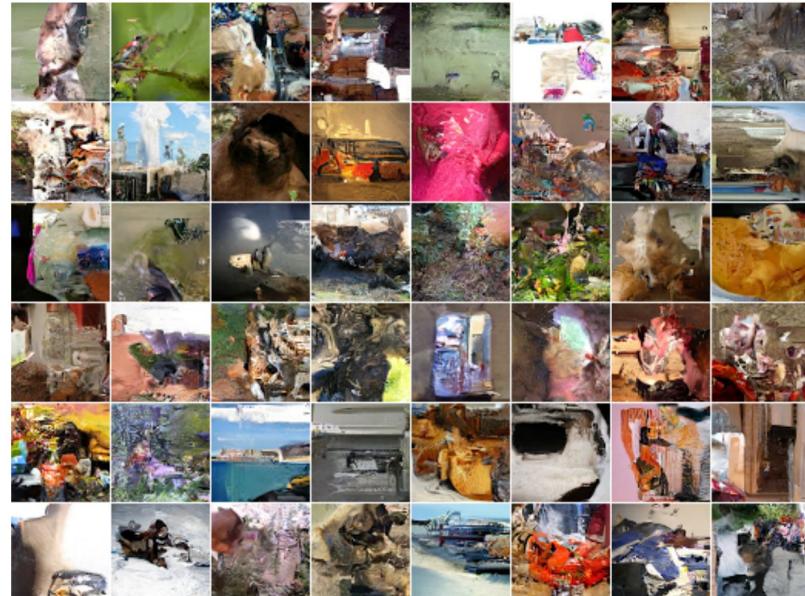
NIC: A baby is asleep next to a teddy bear.

Pixels as a sequence



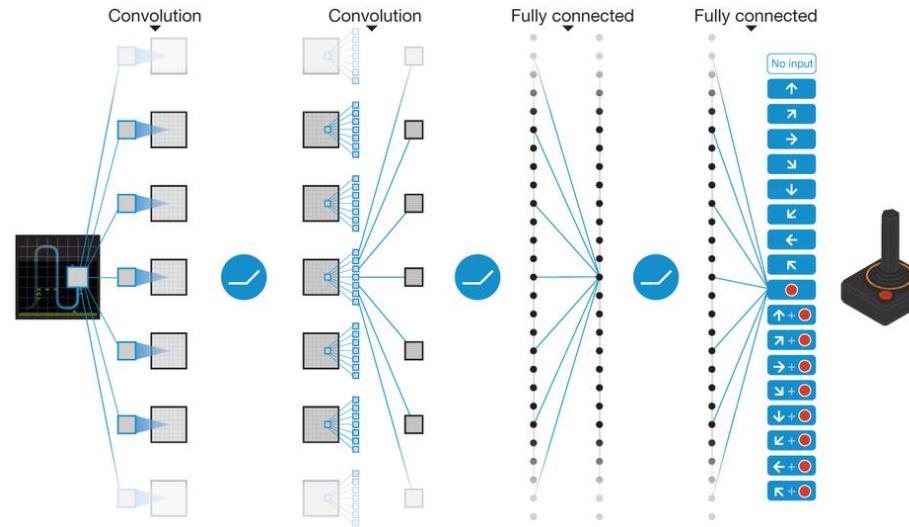
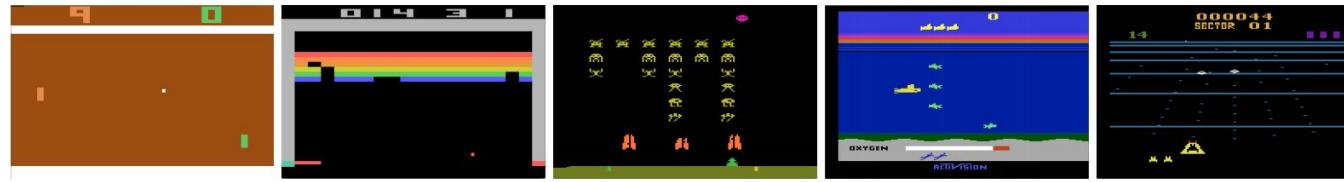
$$p(\mathbf{x}) = \prod_{i=1}^{n^2} p(x_i | x_1, \dots, x_{i-1})$$

The Chain Rule (more on it later)



Deep Reinforcement Learning

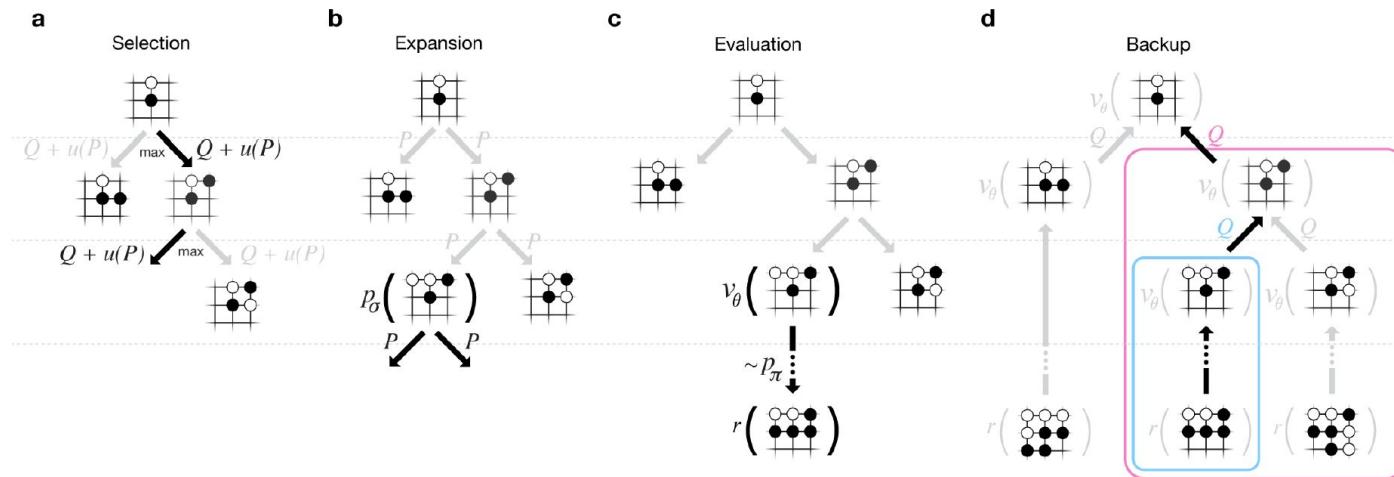
Human-level control through deep reinforcement learning, Mnih et al, Nature 2015



Deep Reinforcement Learning

Mastering the game of Go with deep neural networks and tree search.

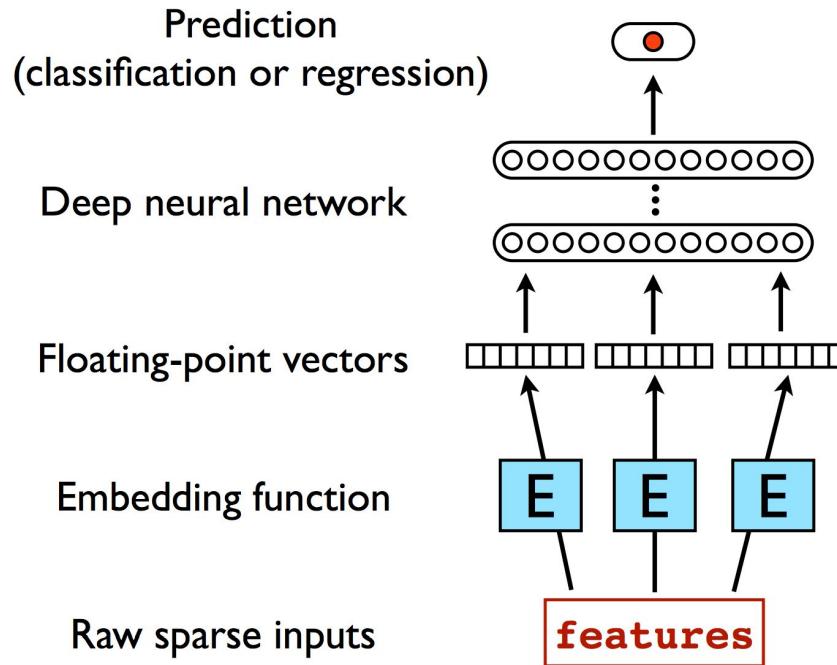
Silver et al, Nature 2016



Part III: Recurrent Neural Networks

Fundamentals

Word Vectors Aside



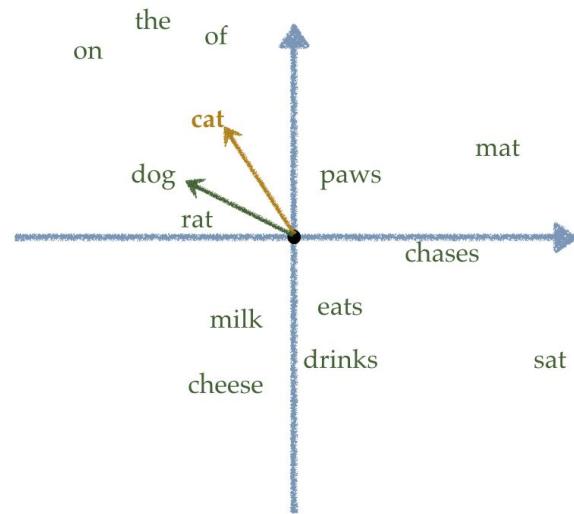
Embedding == Sparse matrix multiplication

Let $v = [0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \dots \ 0 \ 0 \ 0]$

each position is a word ($\text{length}(v) == \text{vocab size}$)

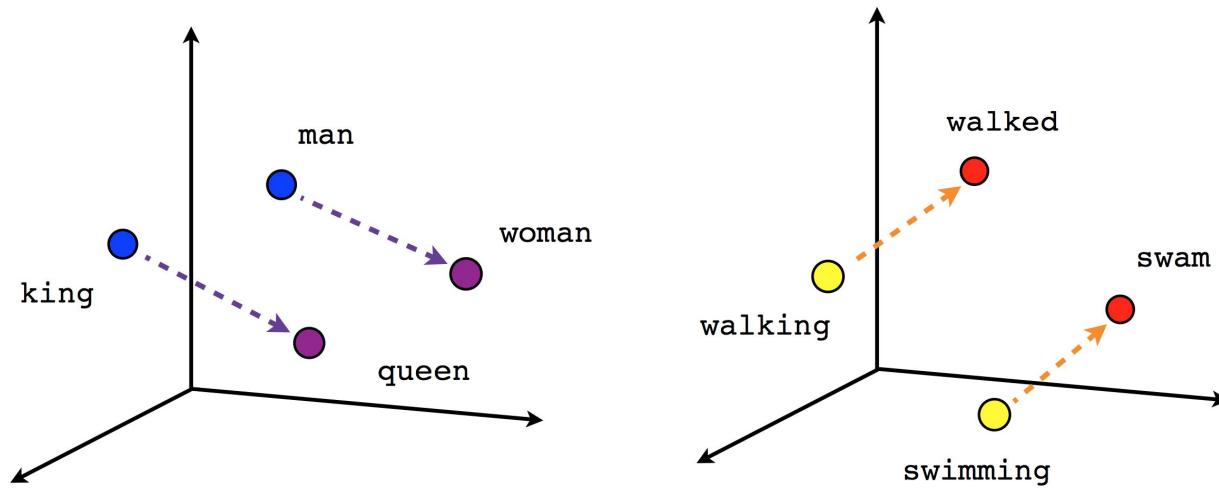
To embed a word we simply do:

$$e = Wv$$



where W is a $D \times \text{vocab size}$ matrix (one D dim. vector for each word)

Vector arithmetic



Solve analogies with vector arithmetic!

$$V(\text{queen}) - V(\text{king}) \approx V(\text{woman}) - V(\text{man})$$

$$V(\text{queen}) \approx V(\text{king}) + (V(\text{woman}) - V(\text{man}))$$

Part III.A The Chain Rule

Loss and Architecture basics

Basics

Supervised Learning: many input / output examples (\mathbf{x}, \mathbf{y})

Find mapping f s.t.

$$\mathbf{y} \approx f(\mathbf{x})$$

Define

$$\mathcal{L}(\boldsymbol{\theta}) = \sum_{i=1}^N \ell(f_{\boldsymbol{\theta}}, \{\mathbf{x}_i, \mathbf{y}_i\})$$

and learn by optimizing

$$\boldsymbol{\theta} = \arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta})$$

Motivating Example: Language

KING LEAR:

O, if you were a feeble sight, the courtesy of your law,
Your sight and several breath, will wear the gods
With his heads, and my hands are wonder'd at the deeds,
So drop upon your lordship's head, and your opinion
Shall be against your honour.

Who wrote these lines?

- A. William Shakespeare
- B. William Shakespeare's ghostwriter
- C. Ben Johnson
- D. Molière (translation)
- E. Andrej Karpathy's recurrent neural network

n-grams

$$P(w_t | w_{t-1}, \dots, w_{t-n+1})$$

what to cook with broccoli and _

what to cook with broccoli and **beef**

what to cook with broccoli and **butter**

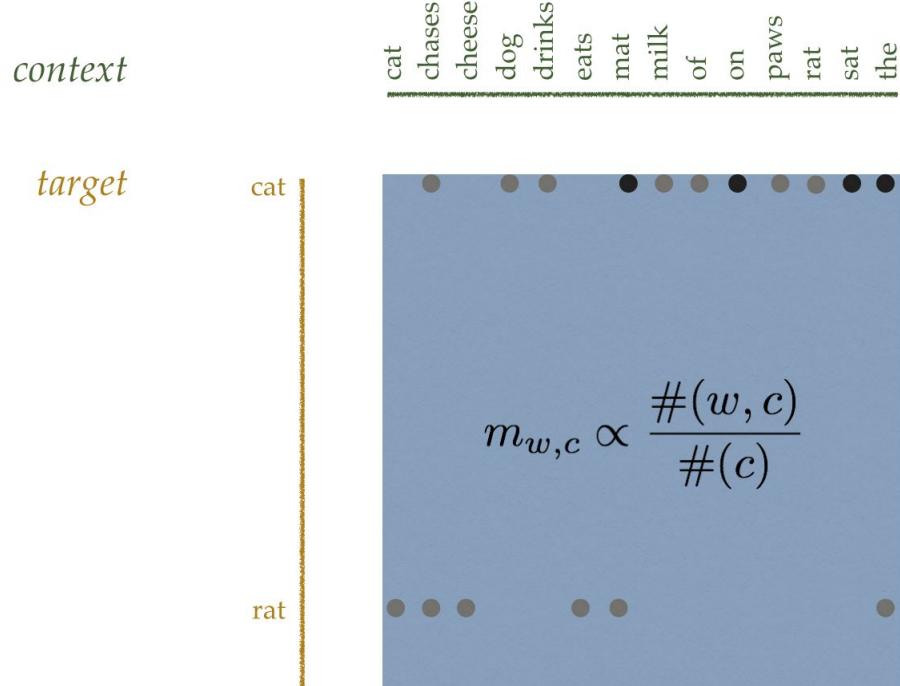
what to cook with broccoli and **blenders**

what to cook with broccoli and **boomboxes**

Modelling “P”

<i>context</i>						<i>target</i>	$P(w_t w_{t-1}, w_{t-2}, \dots w_{t-5})$
the	cat	sat	on	the	mat		0.15
w_{t-5}	w_{t-4}	w_{t-3}	w_{t-2}	w_{t-1}	w_t		
the	cat	sat	on	the	rug		0.12
the	cat	sat	on	the	hat		0.09
the	cat	sat	on	the	dog		0.01
the	cat	sat	on	the	the		0
the	cat	sat	on	the	sat		0
the	cat	sat	on	the	robot		?
the	cat	sat	on	the	printer		?

2-grams



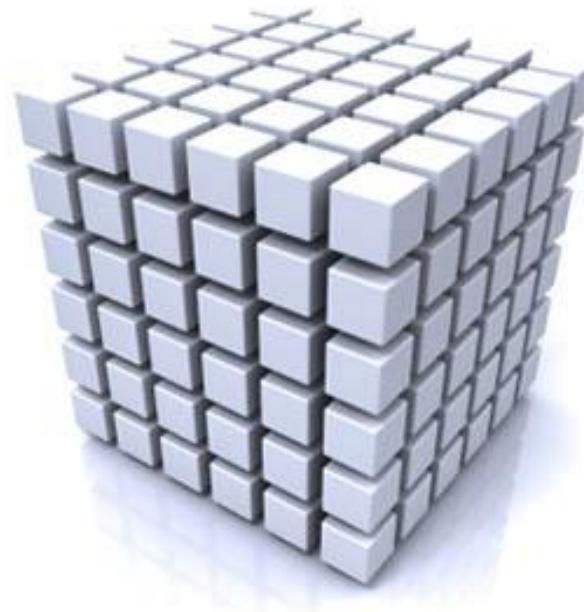
the **cat** sat on the mat
the **cat** drinks milk
the dog chases the **cat**
the paws of the **cat**

the cat chases the **rat**
the **rat** eats cheese
the **rat** eats the mat

n-grams

Vocabulary: 10K

Table size: 10000^n



n-grams

$$P(w_1, w_2, \dots, w_{T-1}, w_T) \approx \prod_{t=1}^T P(w_t | w_{t-1}, \dots, w_{t-n+1})$$

the	cat	sat	on	the	mat	$P(w_1)$
the	cat	sat	on	the	mat	$P(w_2 w_1)$
the	cat	sat	on	the	mat	$P(w_3 w_2, w_1)$
the	cat	sat	on	the	mat	$P(w_4 w_3, w_2)$
the	cat	sat	on	the	mat	$P(w_5 w_4, w_3)$
the	cat	sat	on	the	mat	$P(w_6 w_5, w_4)$

The Chain Rule

$$P(w_1, w_2, \dots, w_{T-1}, w_T) = \prod_{t=1}^T P(w_t | w_{t-1}, w_{t-2}, \dots, w_1)$$

the	cat	sat	on	the	mat	$P(w_1)$
the	cat	sat	on	the	mat	$P(w_2 w_1)$
the	cat	sat	on	the	mat	$P(w_3 w_2, w_1)$
the	cat	sat	on	the	mat	$P(w_4 w_3, w_2, w_1)$
the	cat	sat	on	the	mat	$P(w_5 w_4, w_3, w_2, w_1)$
the	cat	sat	on	the	mat	$P(w_6 w_5, w_4, w_3, w_2, w_1)$

Proof

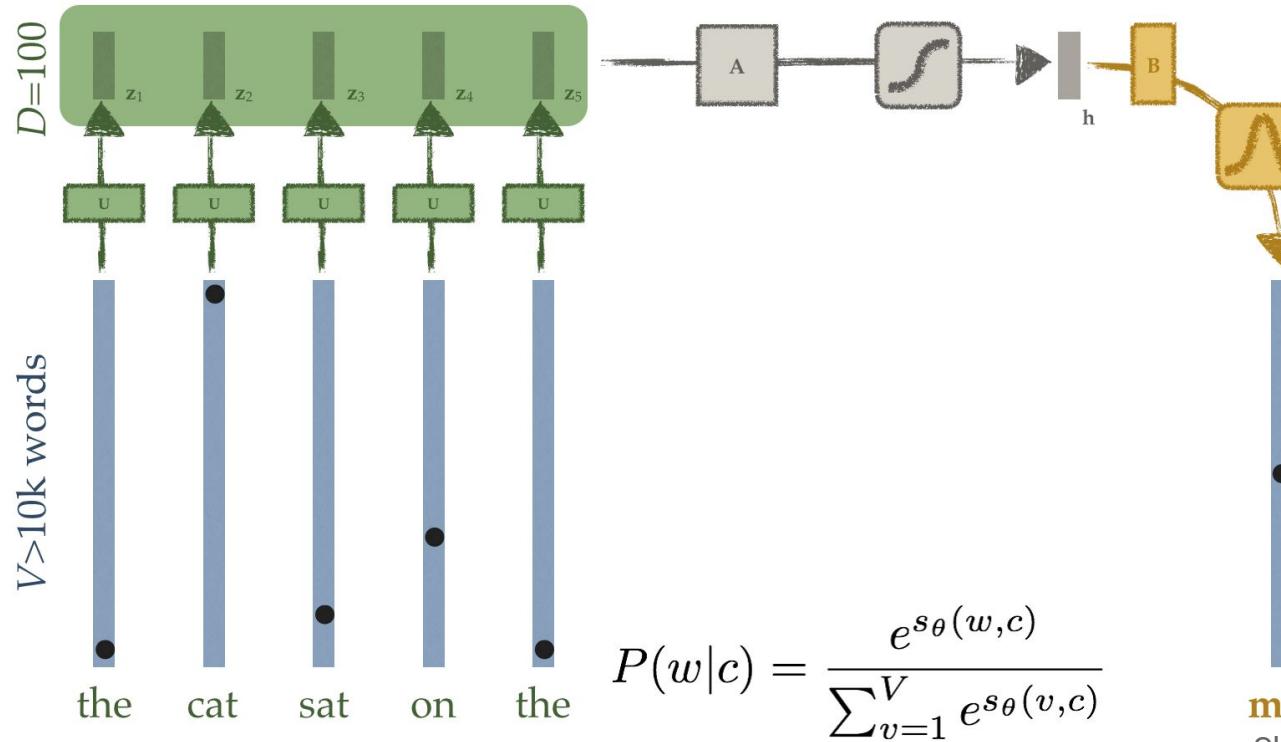
$$\sum_{w_1, \dots, w_T} p(w_1, \dots, w_T) = 1$$

$$\begin{aligned} & \sum_{w_1, \dots, w_T} \prod_{t=1}^T p(w_t | w_1, \dots, w_{t-1}) = \\ & \sum_{w_1, \dots, w_T} p(w_T | w_1, \dots, w_{T-1}) \prod_{t=1}^{T-1} p(w_t | w_1, \dots, w_{t-1}) = \\ & \sum_{w_1, \dots, w_{T-1}} p(w_{T-1} | w_1, \dots, w_{T-2}) \prod_{t=1}^{T-2} p(w_t | w_1, \dots, w_{t-1}) = \\ & \quad \dots \\ & \sum_{w_1, \dots, w_T} \prod_{t=1}^T p(w_t | w_1, \dots, w_{t-1}) = 1 \end{aligned}$$

A Key Insight: vectorizing context

[Yoshua Bengio et al. (2001, 2003), "A Neural Probabilistic Language Model", *JMLR*;
Andriy Mnih and Geoff Hinton, "Three new graphical models for statistical language modeling", *ICML*]

$$p(w_t | w_1, \dots, w_{t-1}) = p_\theta(w_t | f_\theta(w_1, \dots, w_{t-1}))$$

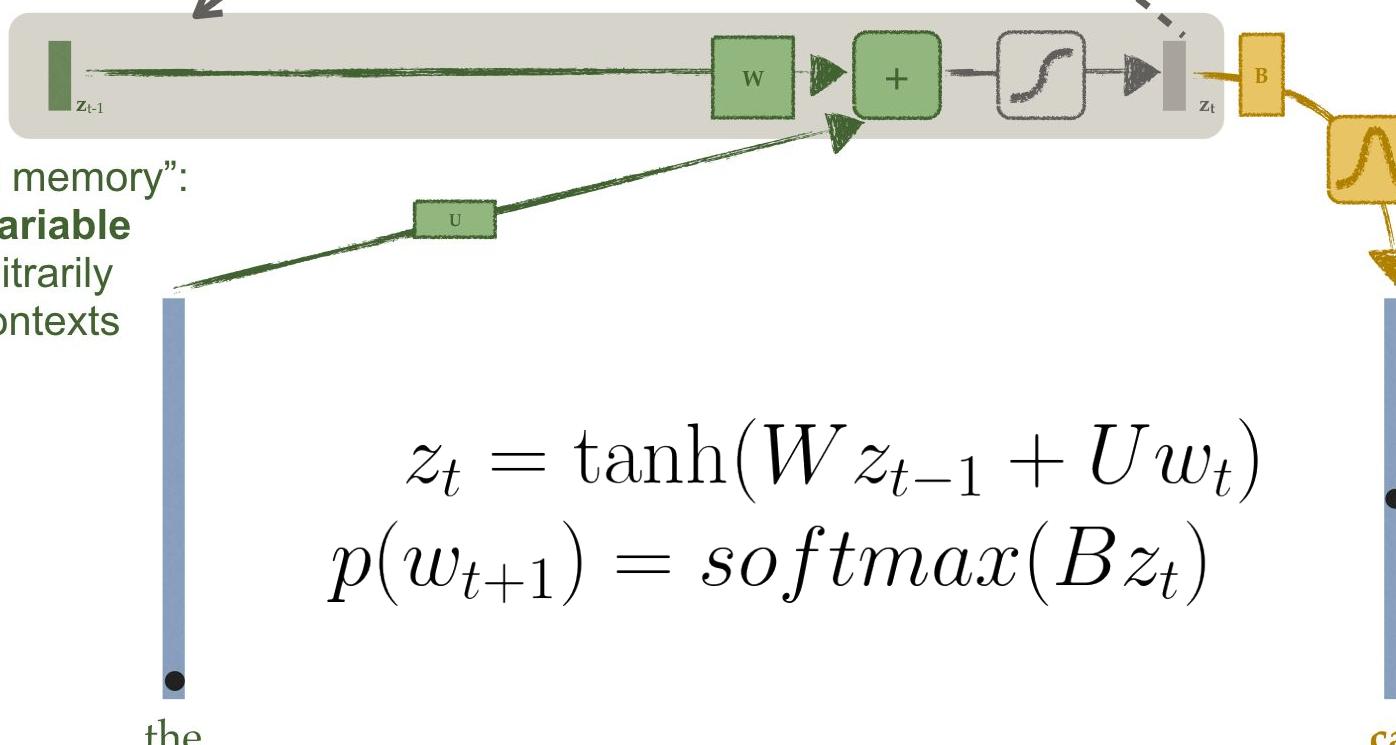


mat

Slide Credit: Piotr Mirowski

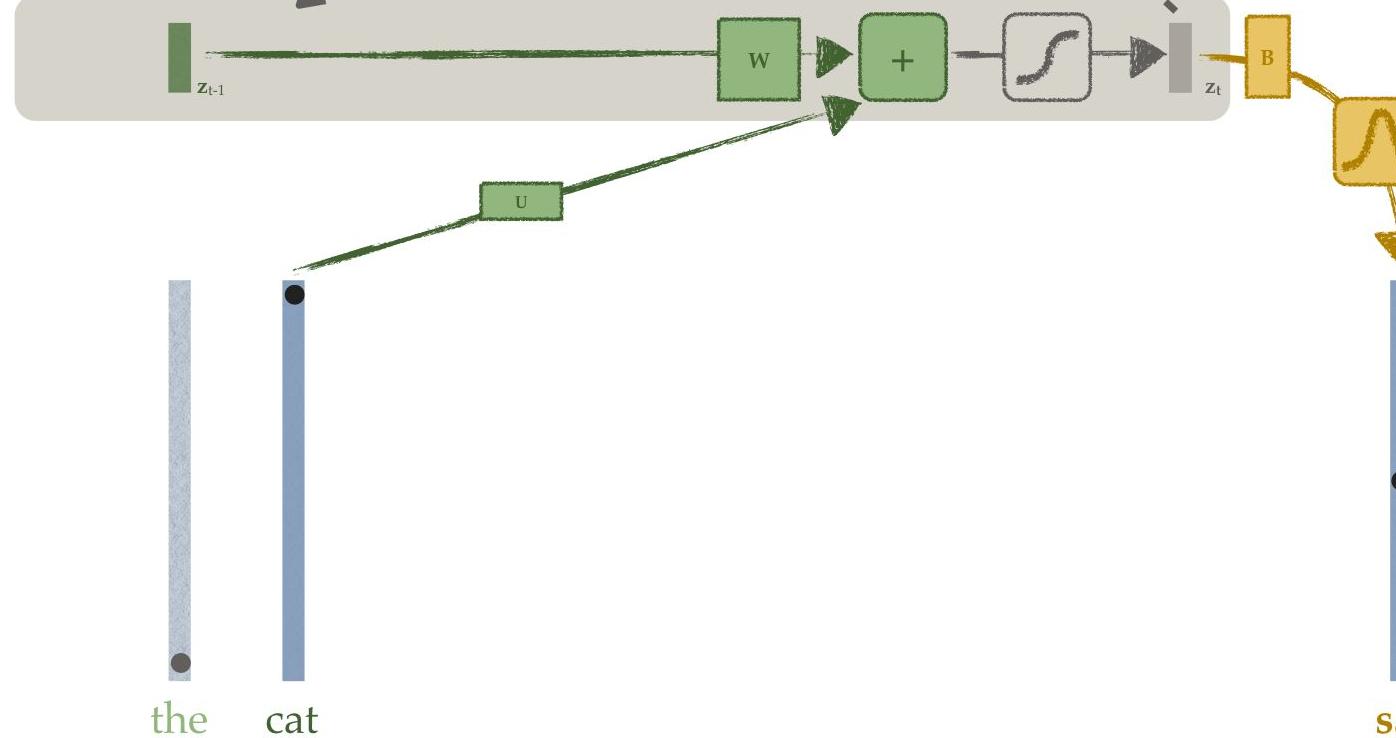
Recurrent Neural Network Language Models

[Jeffrey L Elman (1991) "Distributed representations, simple recurrent networks and grammatical structure", *Machine Learning*; Tomas Mikolov et al. (2010) "Recurrent neural network based language model", *INTERSPEECH*]

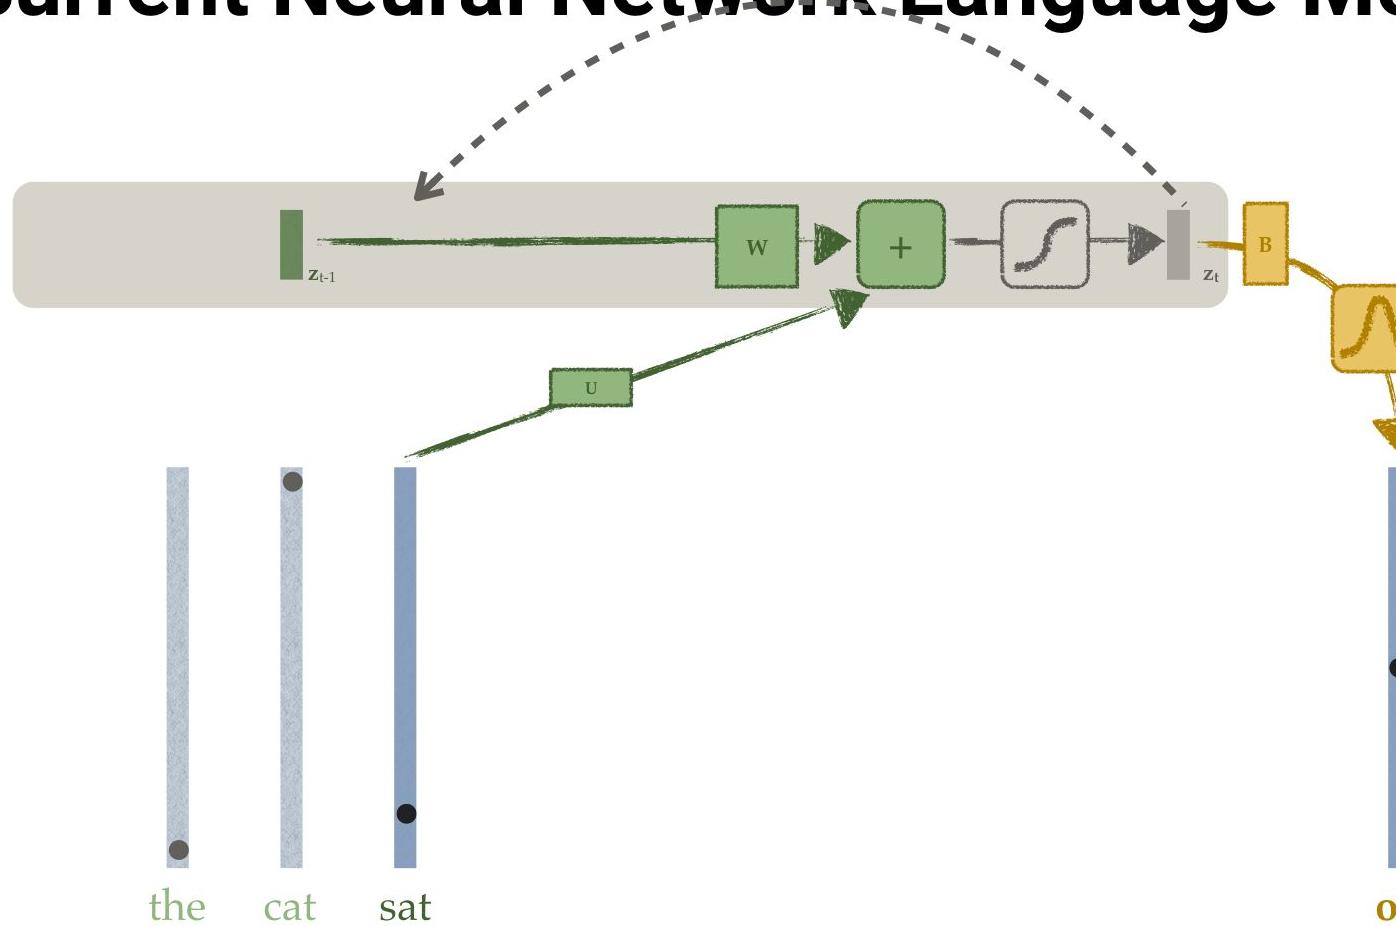


$$\begin{aligned} z_t &= \tanh(Wz_{t-1} + Uw_t) \\ p(w_{t+1}) &= softmax(Bz_t) \end{aligned}$$

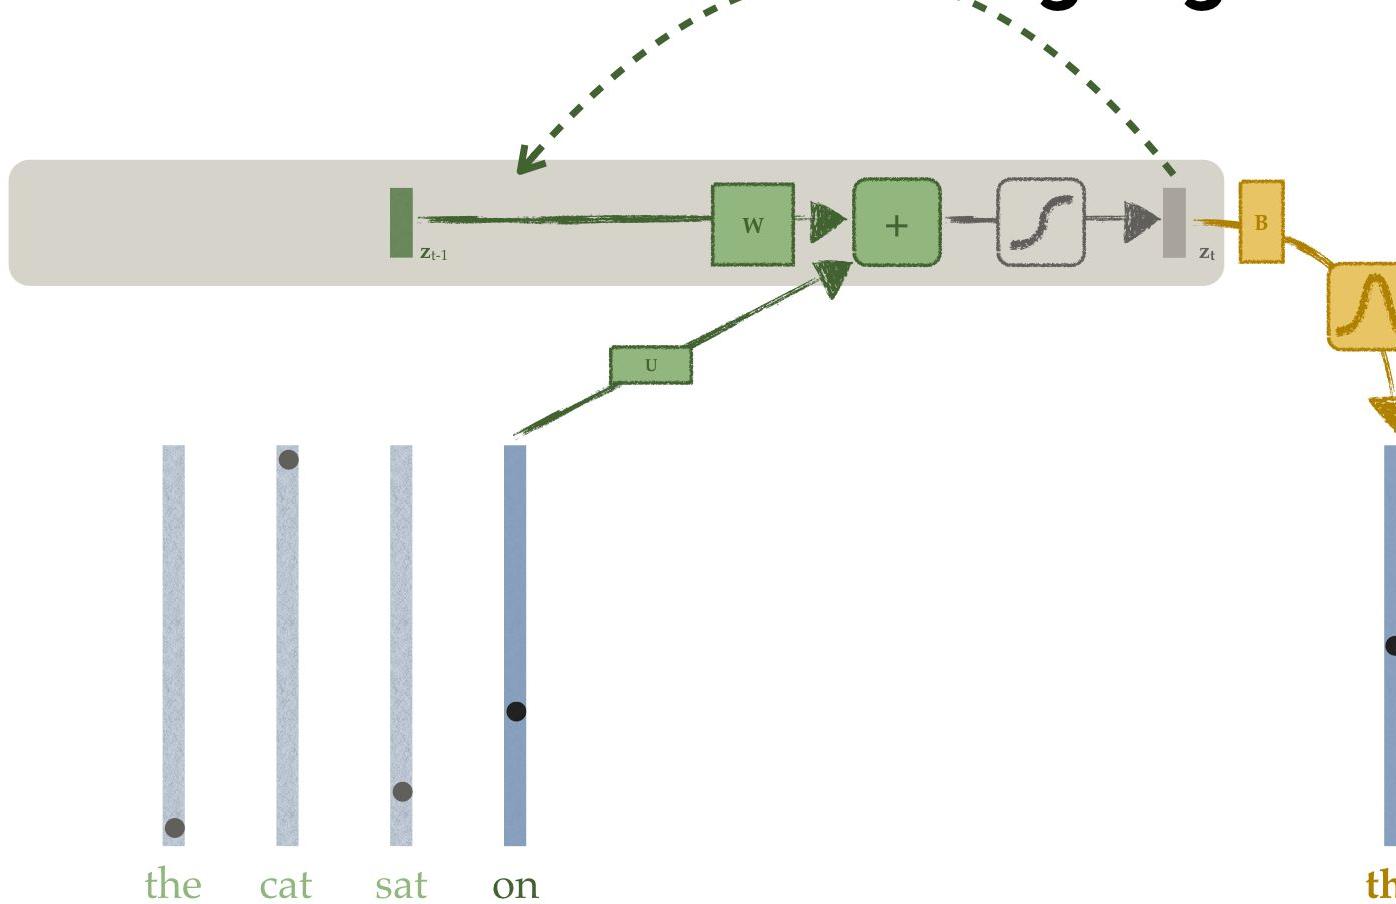
Recurrent Neural Network Language Models



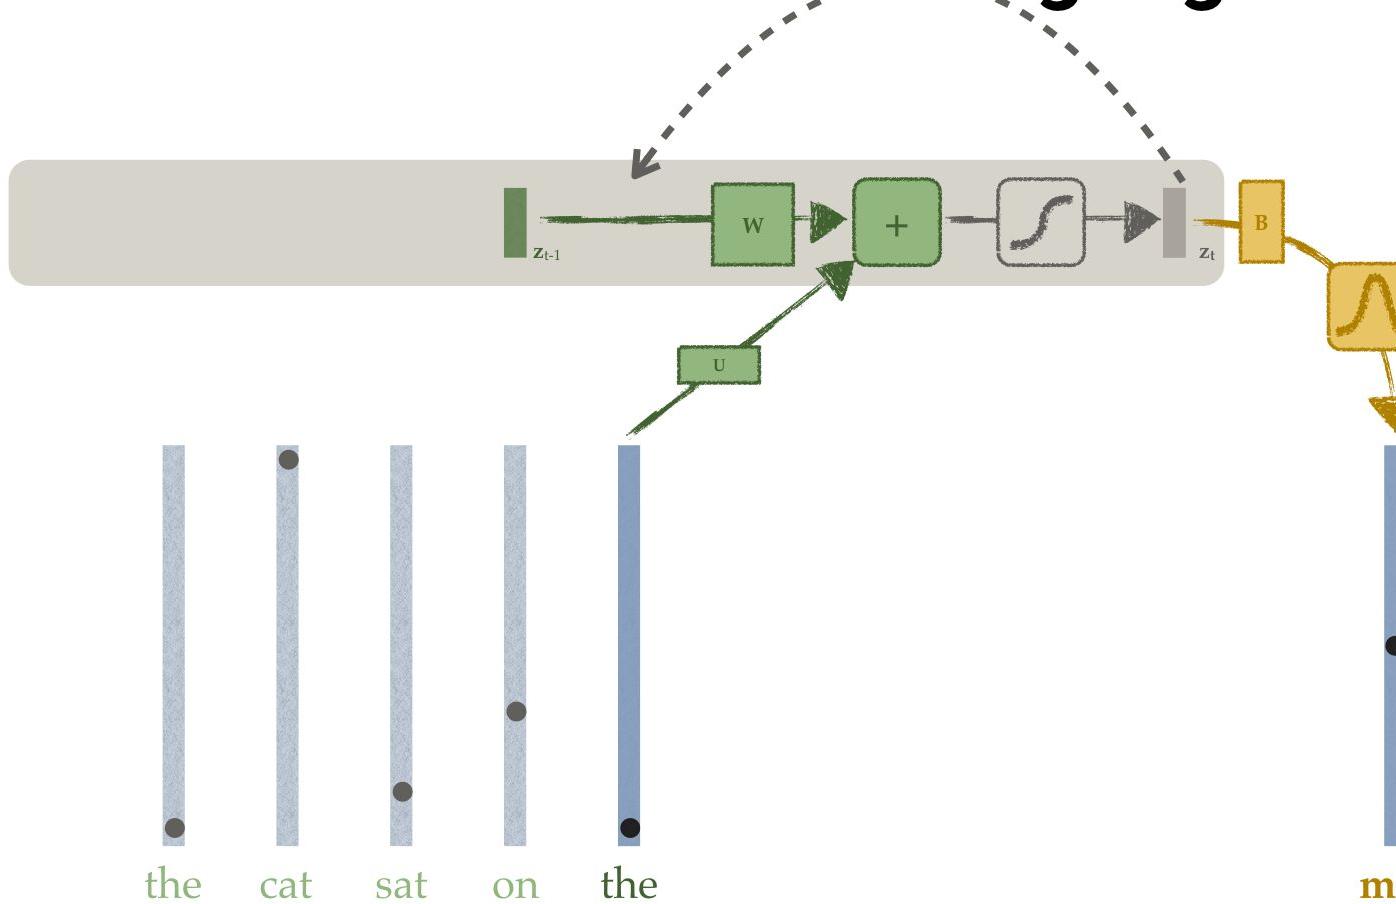
Recurrent Neural Network Language Models



Recurrent Neural Network Language Models



Recurrent Neural Network Language Models



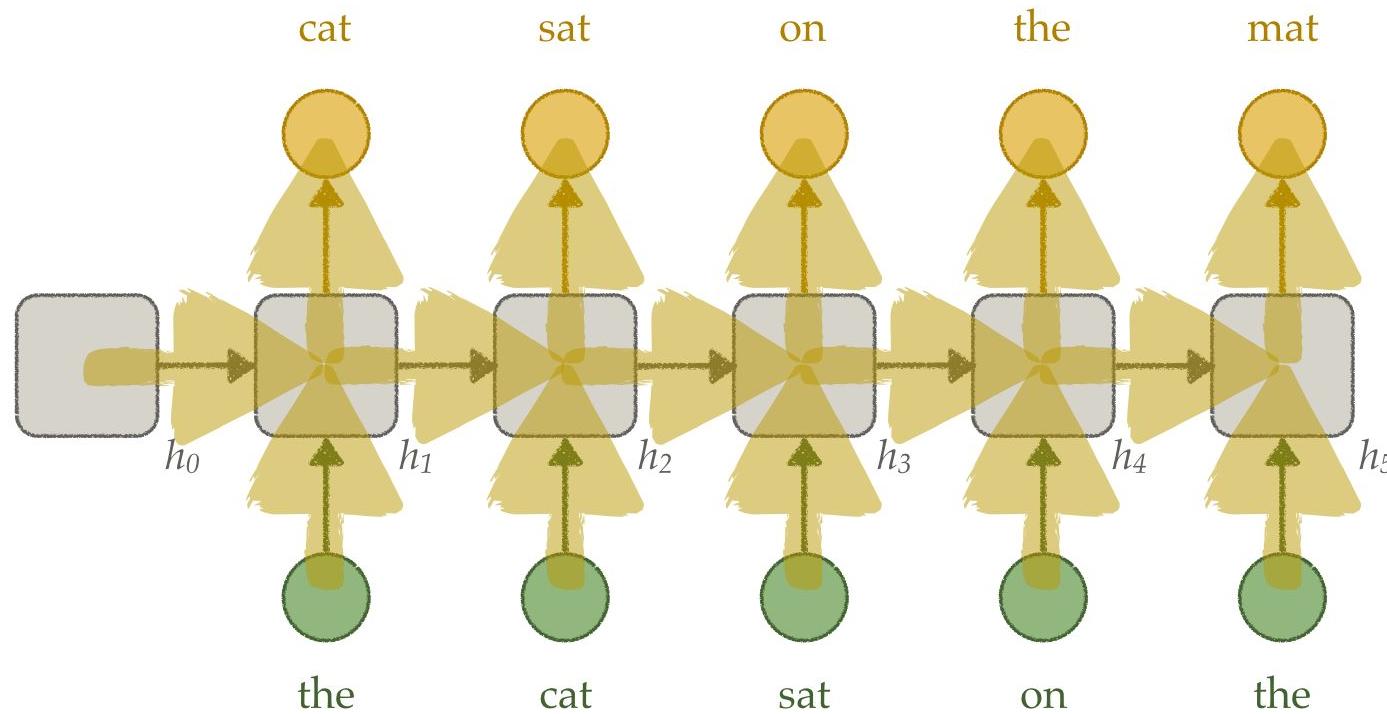
mat

Slide Credit: Piotr Mirowski

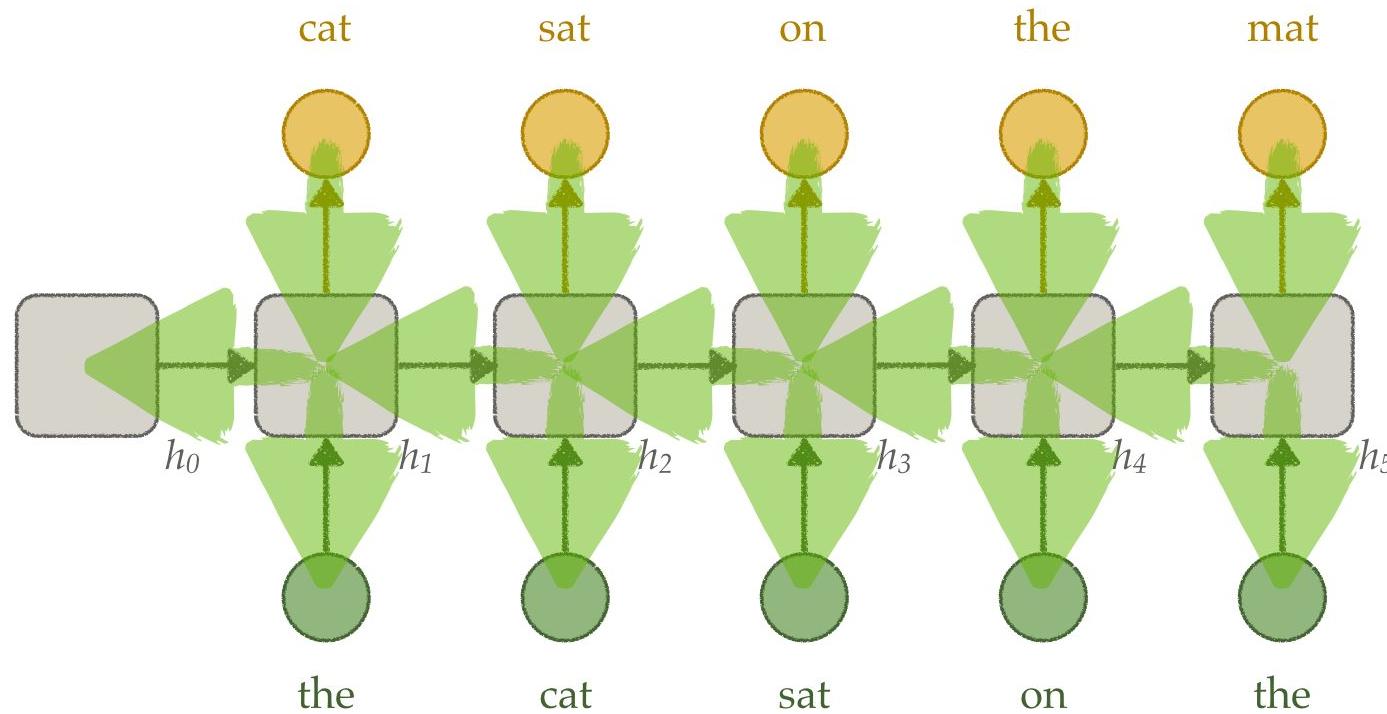
What do we Optimize?

$$\theta^* = \arg \max_{\theta} E_{w \sim data} \log P_{\theta}(w_1, \dots, w_T)$$

Recurrent Neural Network Language Models



Recurrent Neural Network Language Models



Deriving BPTT

$$z_t = \tanh(W z_{t-1} + U w_t)$$

$$p(w_{t+1}) = \text{softmax}(B z_t)$$

$$L(W, U, B) = - \sum_{t=1}^T \log p_{W,U,B}(w_t | w_1, \dots, w_{t-1})$$

$$L_t = - \log p(w_t | w_1, \dots, w_{t-1})$$

Deriving BPTT - B (softmax)

$$L_t = -\log p(w_t | w_1, \dots, w_{t-1})$$

$$\frac{\delta L_t}{\delta B} = \frac{\delta L_t}{\delta p_t} \frac{\delta p_t}{\delta l_t} \frac{\delta l_t}{\delta B}$$

$$\frac{\delta L_t}{\delta B} = (p_t - \text{label}_t) z_t^T$$

$$\frac{\delta L}{\delta B} = \sum_{t=1}^T (p_t - \text{label}_t) z_t^T$$

Deriving BPTT - W (recurrent)

$$z_t = \tanh(W z_{t-1} + U w_t)$$
$$p(w_{t+1}) = \text{softmax}(B z_t)$$

$$\frac{\delta L_t}{\delta W} = \frac{\delta L_t}{\delta p_t} \frac{\delta p_t}{\delta z_t} \frac{\delta z_t}{\delta W}$$

$$\frac{\delta L_t}{\delta W} = \sum_{k=1}^t \frac{\delta L_t}{\delta p_t} \frac{\delta p_t}{\delta z_t} \frac{\delta z_t}{\delta z_k} \frac{\delta z_k}{\delta W}$$

Part III.B Vanishing Gradients

Optimization Tricks

Vanishing (& exploding) Gradients

[Sepp Hochreiter (1991) "Untersuchungen zu dynamischen neuronalen Netzen", *Diploma TUM*;

Yoshua Bengio et al. (1994) "Learning Long-Term Dependencies with Gradient Descent is Difficult", *IEEE Transactions on Neural Networks*]

- n-grams have VERY strong vanishing gradients :)
- RNNs *can* encode long term, but *learning* is hard
- Does long-term ALWAYS matter?
 - Language has strong correlations short term (but also long!)
 - Seq2seq models *must* remember (later)
 - Learning programs

A simple example

$$h_t = w * h_{t-1}$$

$$h_t = w^t * h_0$$

$h_t \rightarrow \infty$ if $|w| > 1$

$h_t \rightarrow 0$ if $|w| < 1$

A simple example

$$h_t = w * h_{t-1}$$

$$h_t = w^t * h_0$$

$h_t \rightarrow \infty$ if $|w| > 1$

$h_t \rightarrow 0$ if $|w| < 1$

$$h_t = \tanh(w * h_{t-1})$$

$h_t \rightarrow$ bounded for all $w :$)

$$\delta h_t / \delta h_{t-T} = (1 - \tanh^2(w * h_{t-1})) * w * \delta h_{t-1} / \delta h_{t-T}$$

A simple example

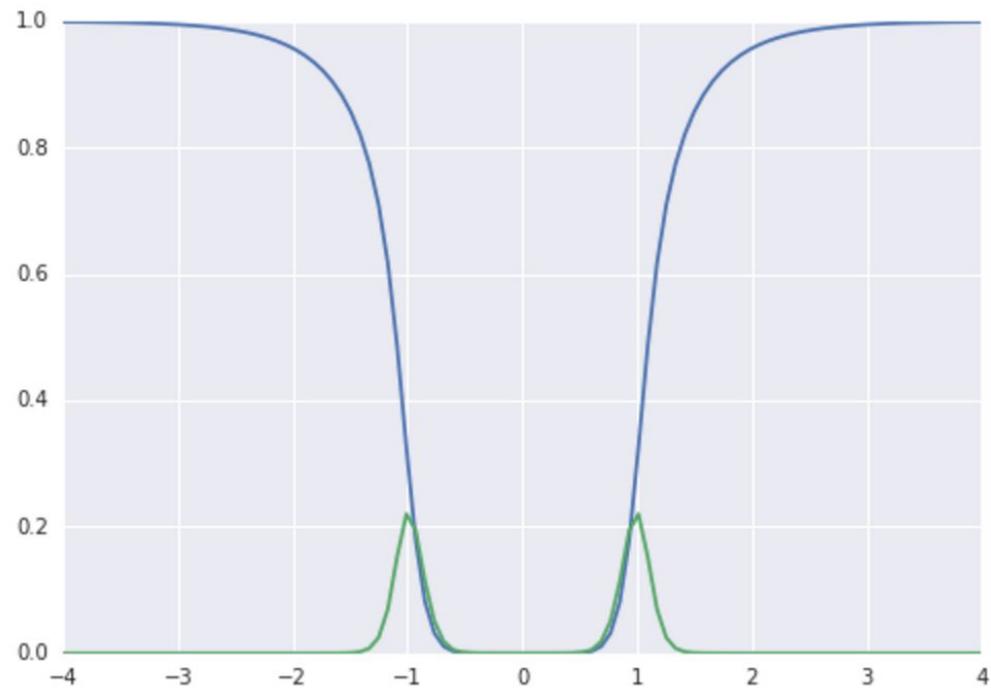
```
import numpy as np
from matplotlib import pyplot

def forward_backward_prop(w, T):
    hs = [0.5]
    for _ in range(T):
        hs.append(np.tanh(w*hs[-1]))
    dh = 1
    for t in range(T):
        dh = (1-hs[-1-t]*hs[-1-t])*w*dh
    return hs[-1], dh

T = 10
wlim = 4

ws = np.linspace(-wlim, wlim, 100)
res = []
for w in ws:
    res.append(forward_backward_prop(w,T))

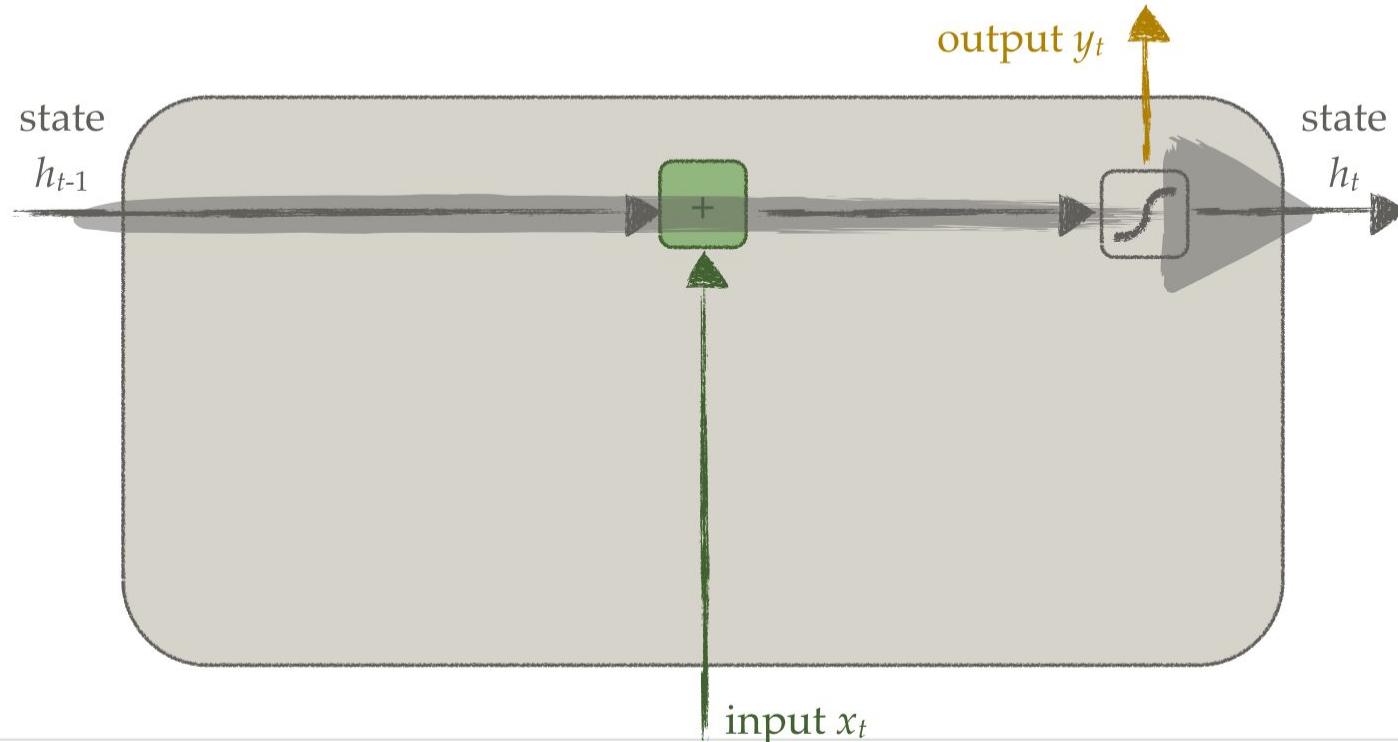
pyplot.plot(ws, [r[0] for r in res])
pyplot.plot(ws, [r[1] for r in res])
```



Part III.C LSTMs (& GRUs)

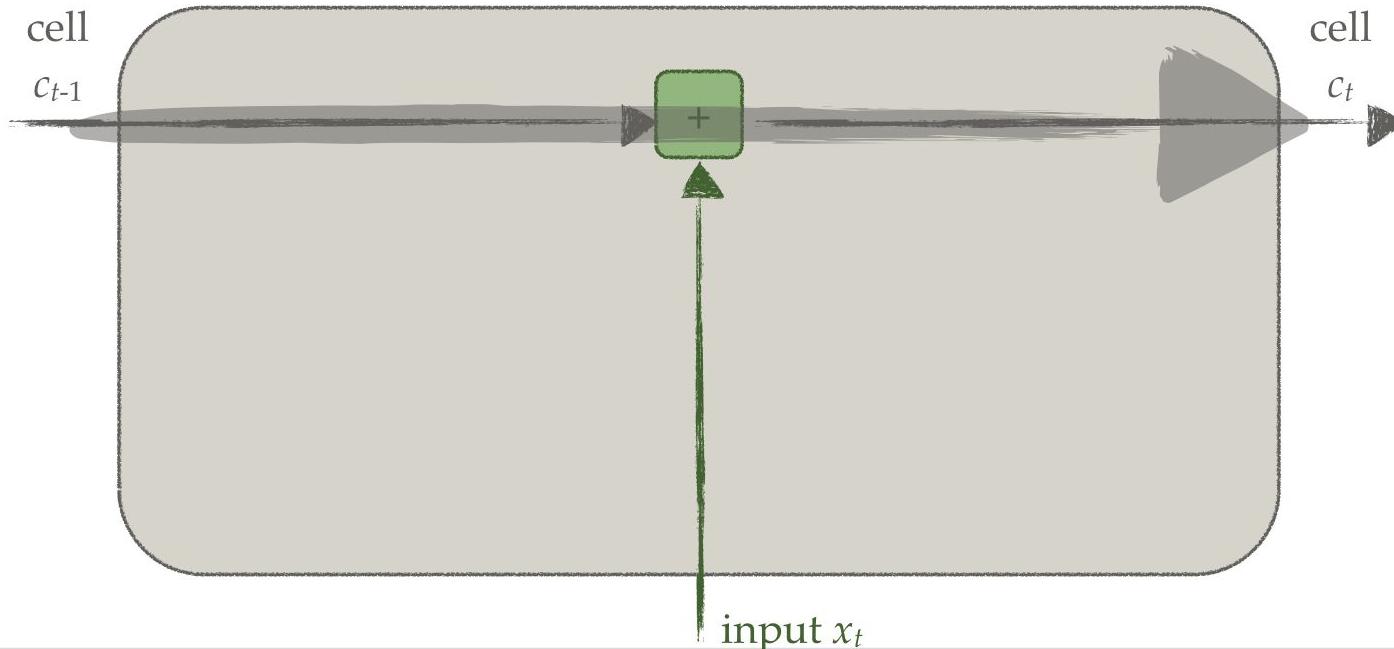
The Model Architecture

LSTMs



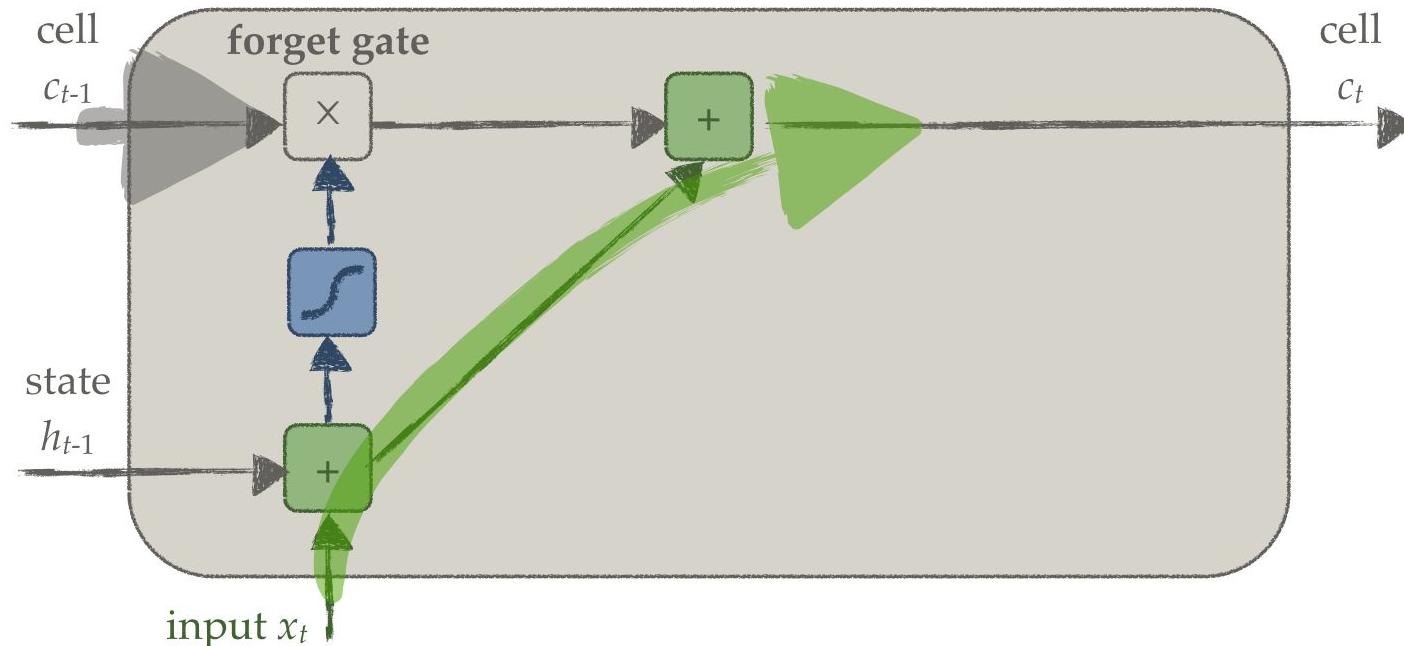
Requirement #1: linear cell

[Sepp Hochreiter and Jürgen Schmidhuber (1997) “Long Short-Term Memory”, *Neural Computation*;
Alex Graves (2013a) “Generating sequences with recurrent neural networks”, arXiv 1308.0850]

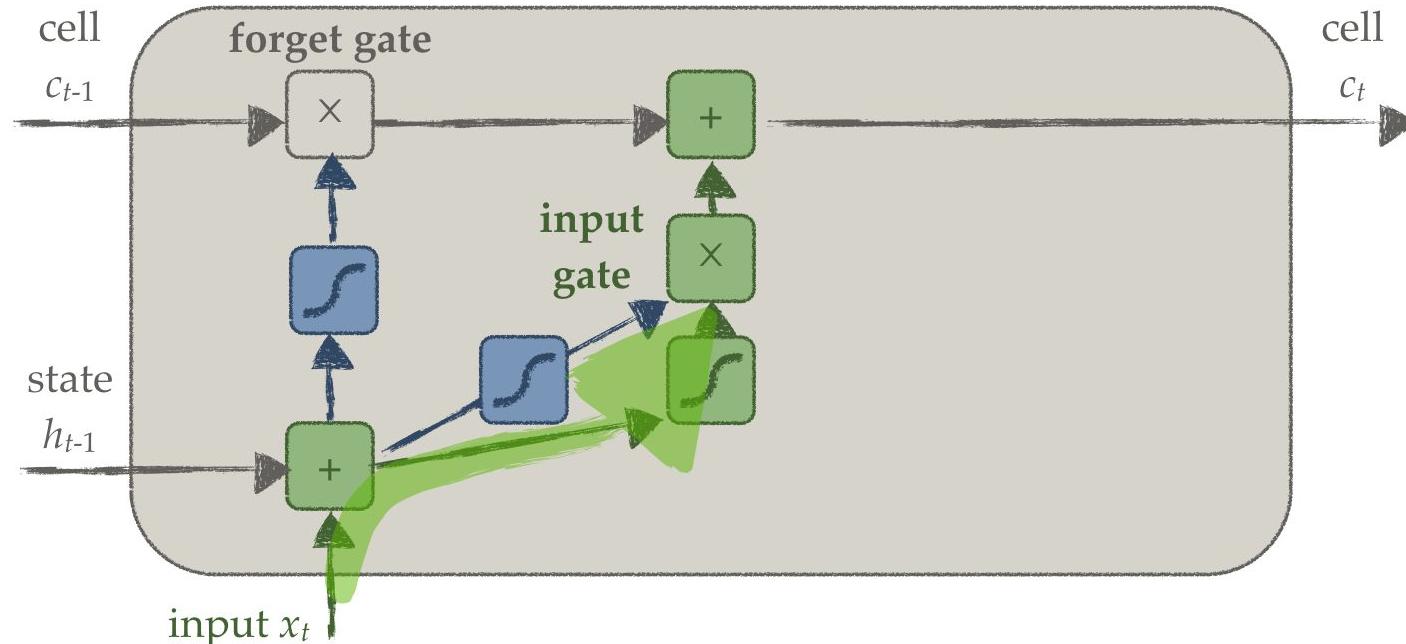


Requirement #2: forget information

[Sepp Hochreiter and Jürgen Schmidhuber (1997) “Long Short-Term Memory”, *Neural Computation*; Alex Graves (2013a) “Generating sequences with recurrent neural networks”, arXiv 1308.0850]



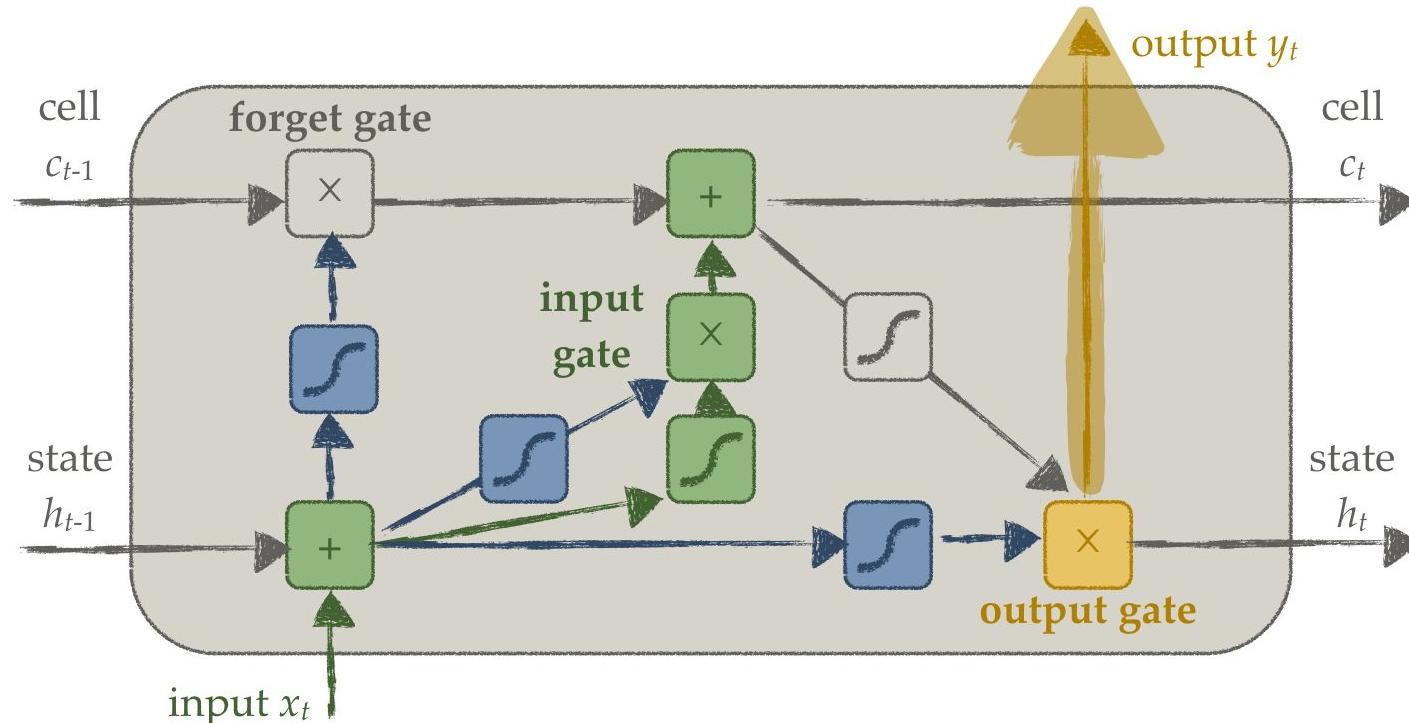
Requirement #3: ignore inputs



[Hochreiter and Schmidhuber (1997); Graves (2013a)]

Learning Sequences – Piotr Mirowski

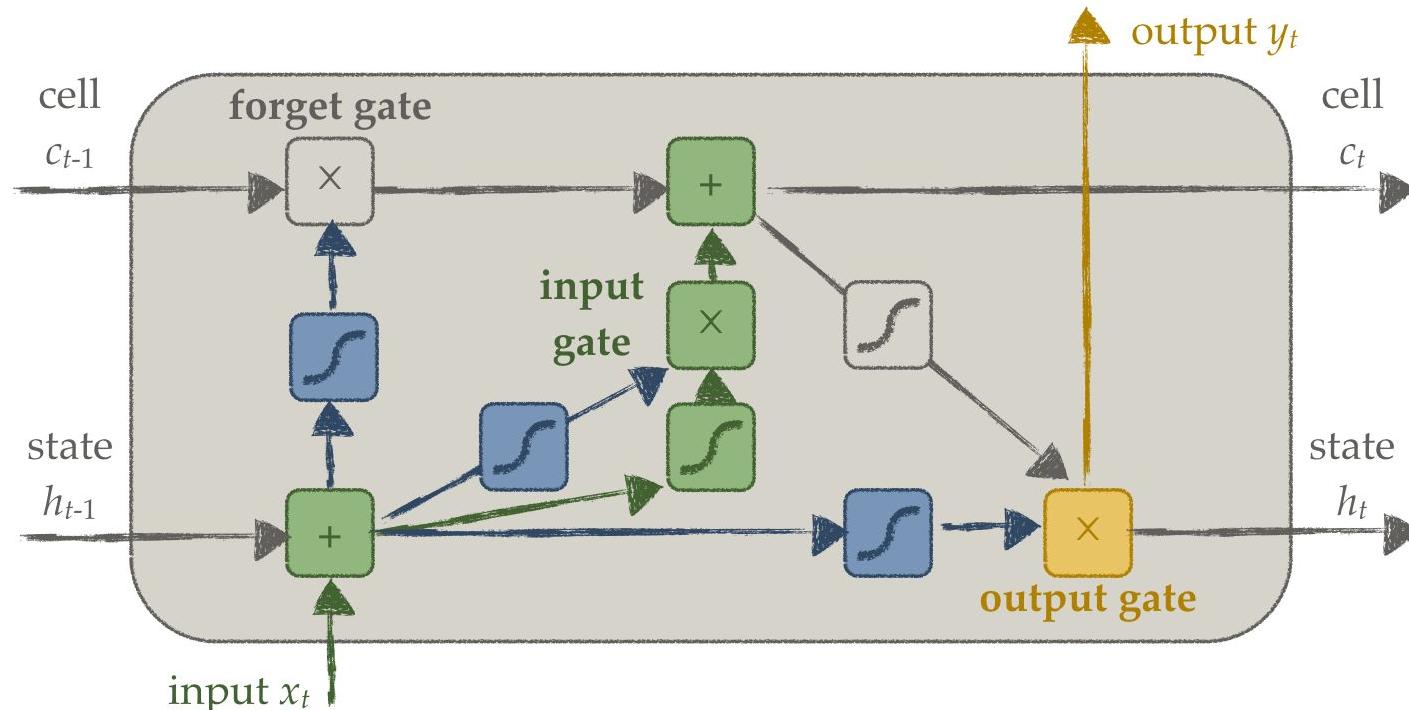
Requirement #4: control outputs



[Hochreiter and Schmidhuber (1997); Graves (2013a)]

Learning Sequences – Piotr Mirowski

Long Short-Term Memory (LSTM)



[Hochreiter and Schmidhuber (1997); Graves (2013a)]

Learning Sequences – Piotr Mirowski

LSTM [Hochreiter et al, 1997][Gers et al, 1999]

$$\begin{aligned}
 i_t &= W_{ix}x_t + W_{ih}h_{t-1} + b_i \\
 j_t &= W_{jx}x_t + W_{jh}h_{t-1} + b_j \\
 f_t &= W_{fx}x_t + W_{fh}h_{t-1} + b_f \\
 o_t &= W_{ox}x_t + W_{oh}h_{t-1} + b_o \\
 c_t &= \sigma(f_t) \odot c_{t-1} + \sigma(i_t) \odot \tanh(j_t) \\
 h_t &= \sigma(o_t) \odot \tanh(c_t)
 \end{aligned}$$

Enables long term dependencies to flow



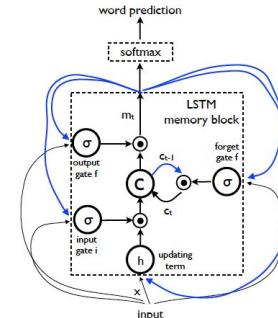
```

def __call__(self, inputs, state, scope=None):
    """Long short-term memory cell (LSTM)."""
    with vs.variable_scope(scope or type(self).__name__): # "BasicLSTMCell"
        # Parameters of gates are concatenated into one multiply for efficiency.
        c, h = array_ops.split(1, 2, state)
        concat = linear([inputs, h], 4 * self._num_units, True)

        # i = input_gate, j = new_input, f = forget_gate, o = output_gate
        i, j, f, o = array_ops.split(1, 4, concat)

        new_c = c * sigmoid(f + self._forget_bias) + sigmoid(i) * tanh(j)
        new_h = tanh(new_c) * sigmoid(o)

    return new_h, array_ops.concat(1, [new_c, new_h])
  
```



GRU, aka simplified LSTM [Cho et al, 2014]

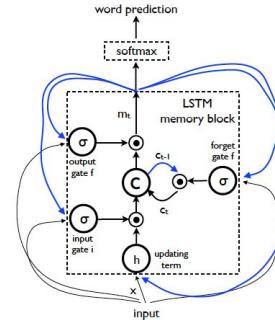
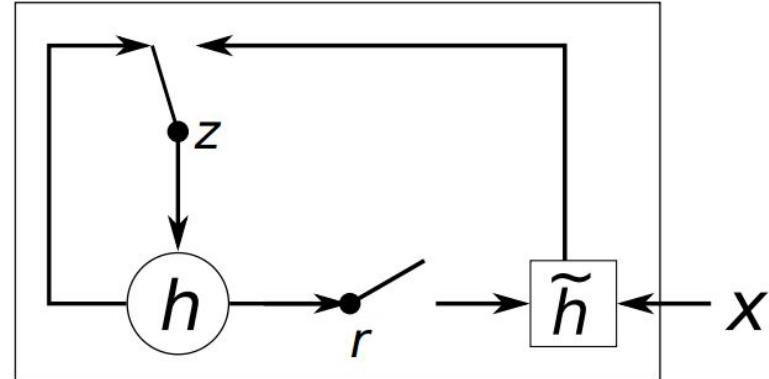
$$z_t = \sigma(x_t U_z + h_{t-1} W_z + b_z)$$

$$r_t = \sigma(x_t U_r + h_{t-1} W_r + b_r)$$

$$\tilde{h} = \tanh(x_t U_h + (h_{t-1} \odot r_t) W_h + b_h)$$

$$h_t = (1 - z_t) \odot \tilde{h} + z_t \odot h_{t-1}$$

Enables long term dependencies to flow



How about exploding gradients? [Pascanu et al, 2013]

Algorithm 1 Pseudo-code for norm clipping

```
 $\hat{g} \leftarrow \frac{\partial \mathcal{E}}{\partial \theta}$ 
if  $\|\hat{g}\| \geq threshold$  then
     $\hat{g} \leftarrow \frac{threshold}{\|\hat{g}\|} \hat{g}$ 
end if
```

Summary

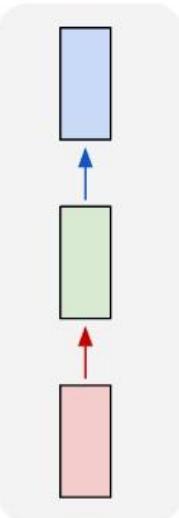
- Sequences are important, everything is a sequence
- Recurrent Nets are a way forward, but they pose some challenges
 - Chain Rule makes us not assume independence
 - Vectorizing context is key for feasibility / model size
 - Vanishing/Exploding gradients
 - LSTMs is better at this, but by no means “solves” all problems
 - Expensive to compute

Part IV: Recurrent Networks as Sequence Decoders

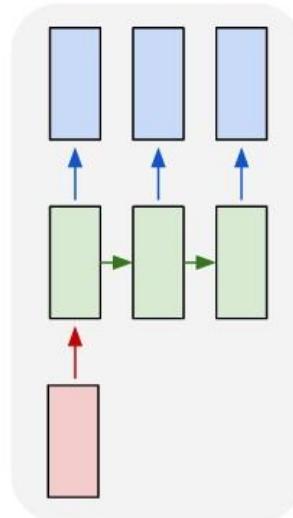
Sequence Models

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

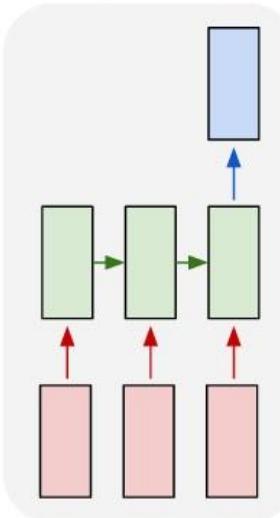
one to one



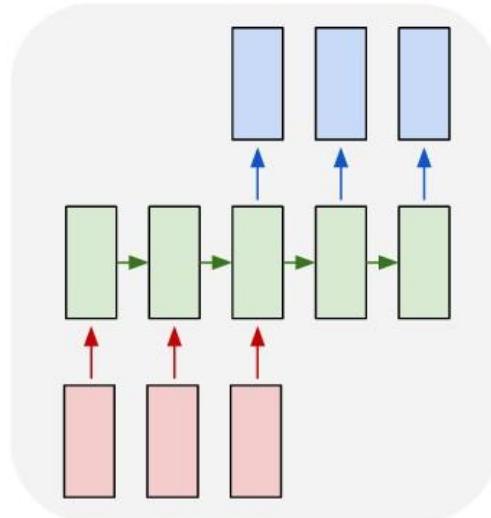
one to many



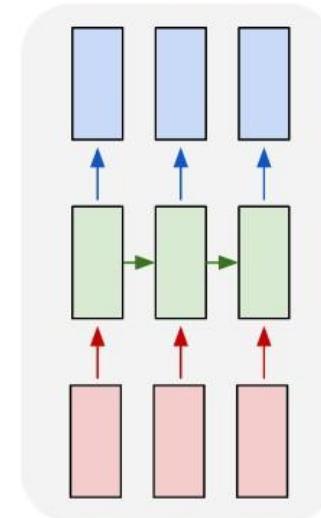
many to one



many to many



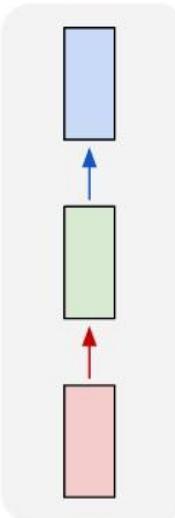
many to many



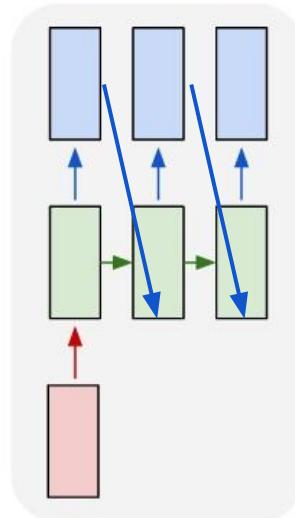
Sequence Models (Autoregressive)

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

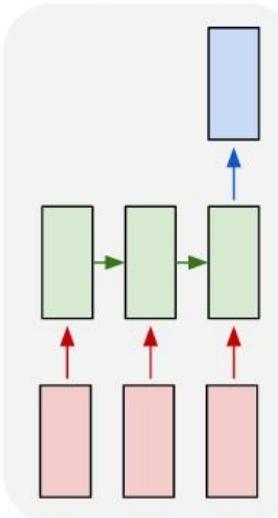
one to one



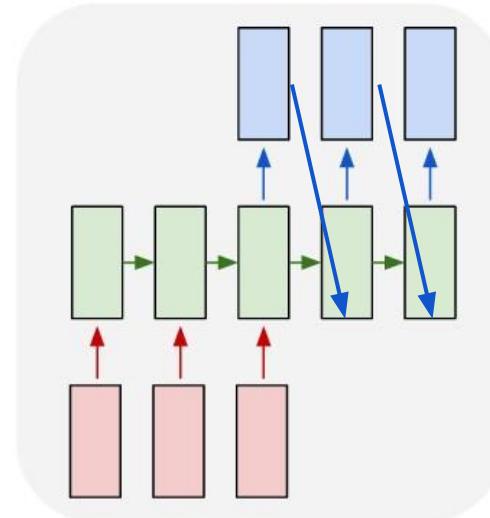
one to many



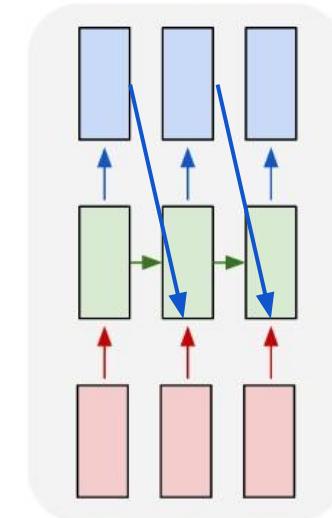
many to one



many to many



many to many



Mini-Turing Test

Some of the obese people lived five to eight years longer than others.

Abu Dhabi is going ahead to build solar city and no pollution city.

Or someone who exposes exactly the truth while lying.

VIERA , FLA . -- Sometimes, Rick Eckstein dreams about baseball swings.

For decades, the quintessentially New York city has elevated its streets to the status of an icon.

The lawsuit was captioned as United States ex rel.

Mini-Turing Test

Some of the obese people lived five to eight years longer than others.

Abu Dhabi is going ahead to build solar city and no pollution city.

Or someone who exposes exactly the truth while lying.

VIERA , FLA . -- Sometimes, Rick Eckstein dreams about baseball swings.

For decades, the quintessentially New York city has elevated its streets to the status of an icon.

The lawsuit was captioned as United States ex rel.

Mini-Turing Test

Some of the obese people lived five to eight years longer than others.

Abu Dhabi is going ahead to build solar city and no pollution city.

Or someone who exposes exactly the truth while lying.

VIERA , FLA . -- Sometimes, Rick Eckstein dreams about baseball swings.

For decades the quintessentially New York city has elevated its streets to the status of [an icon](https://github.com/tensorflow/models/tree/master/lm_1b).

The lawsuit was captioned as United States ex rel.

Results

Measure: perplexity (~branching factor (low == good))

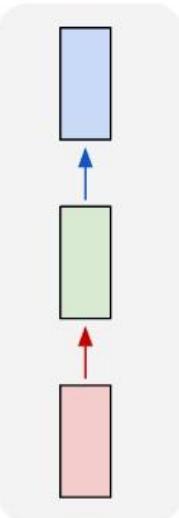
Model	Perplexity	#params (Billions)
5-gram (smoothed) (2013)	67.6	1.8
RNN + Maxent (2013)	51.3	20
Big Ensemble (2013)	45.3	LOTS
RNN (ICLR, 2015)	68.3	4.1
RNN + 5-gram (ICLR, 2015)	42.0	LOTS
Sparse NNMF (IS, 2015)	52.9	33 (!)
Ours (LSTM) (2016)	30.0	1.0
Ours (Ensemble) (2016)	23.7	LOTS

HUMANS: ???

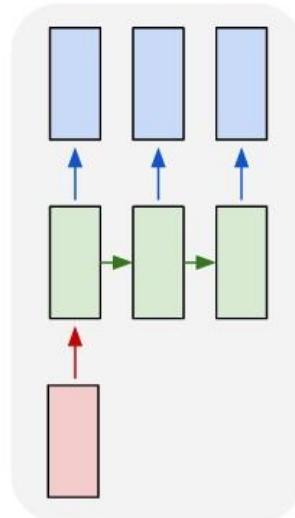
Sequence Models

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

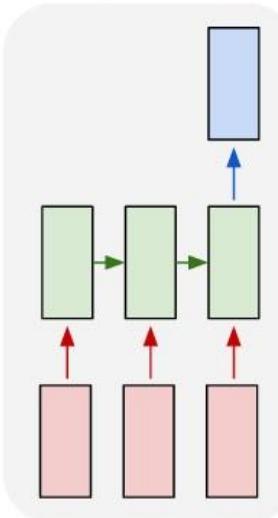
one to one



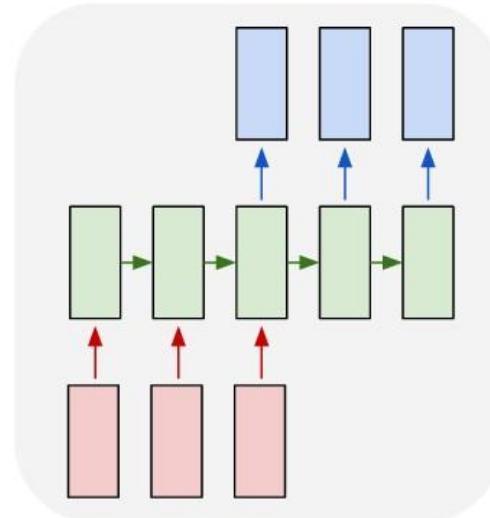
one to many



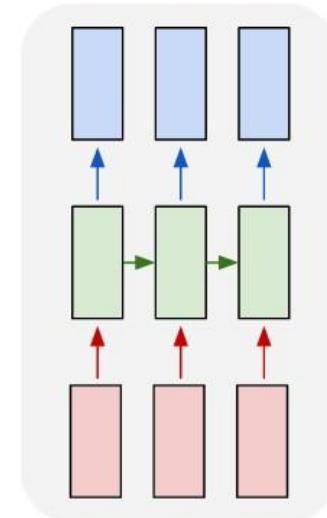
many to one



many to many



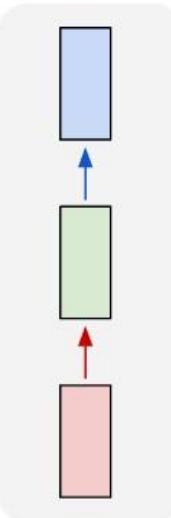
many to many



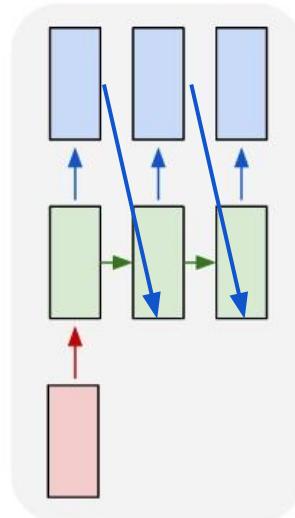
Sequence Models (Autoregressive)

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

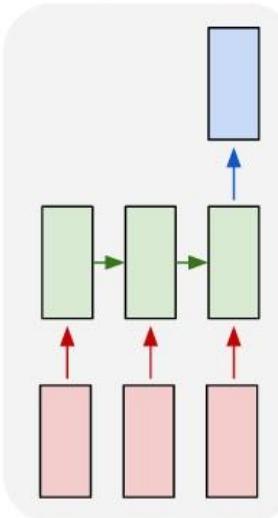
one to one



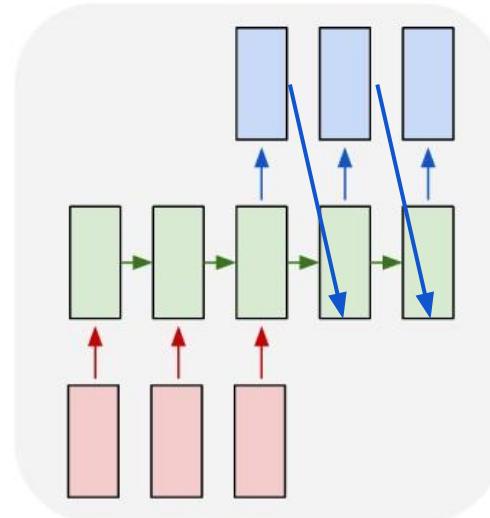
one to many



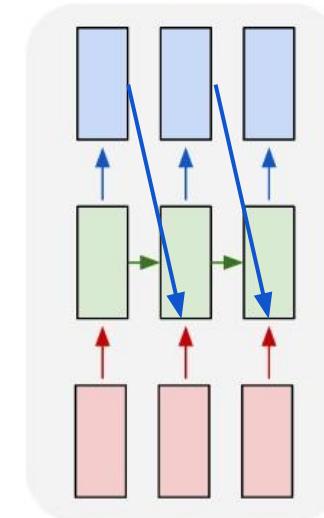
many to one



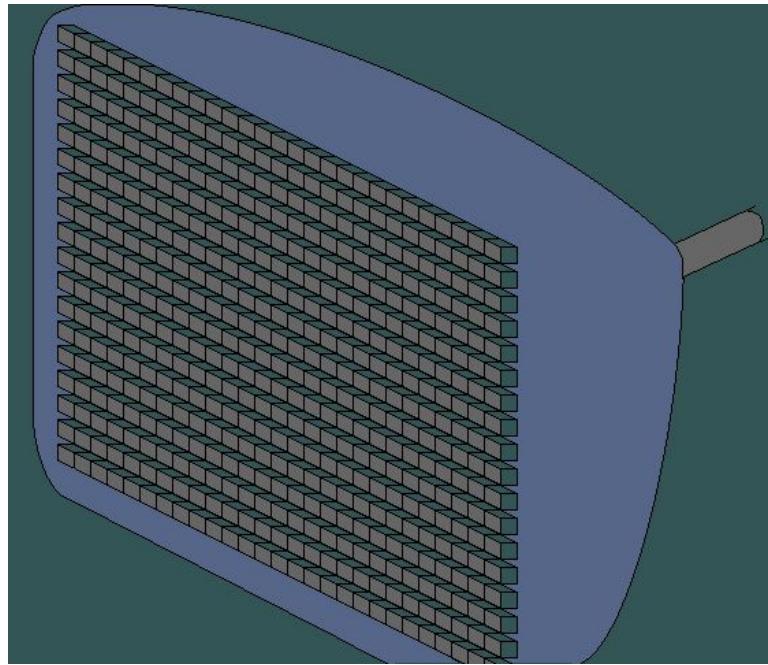
many to many



many to many

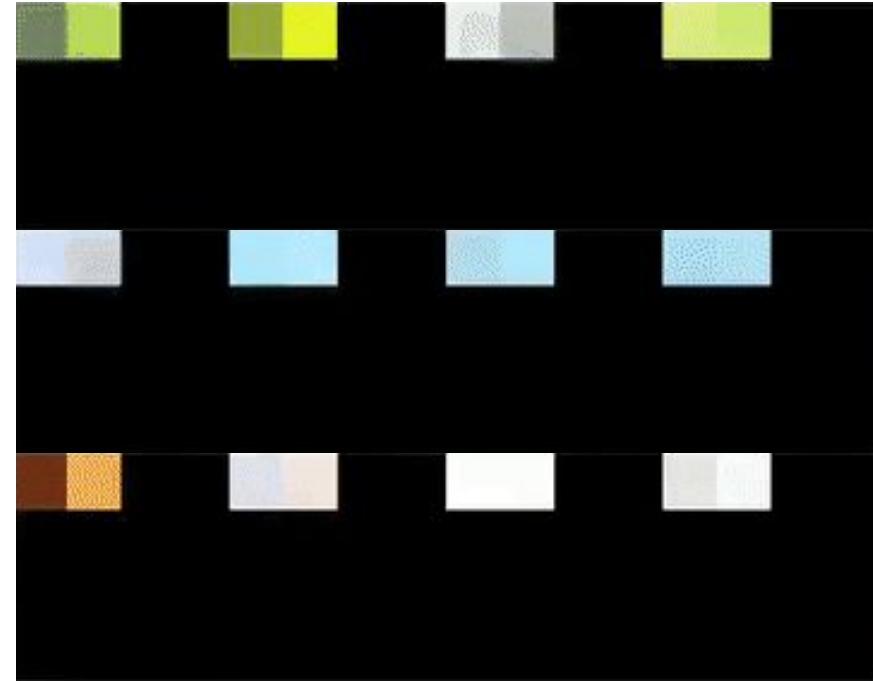


Modeling Images



Pixel-by-pixel

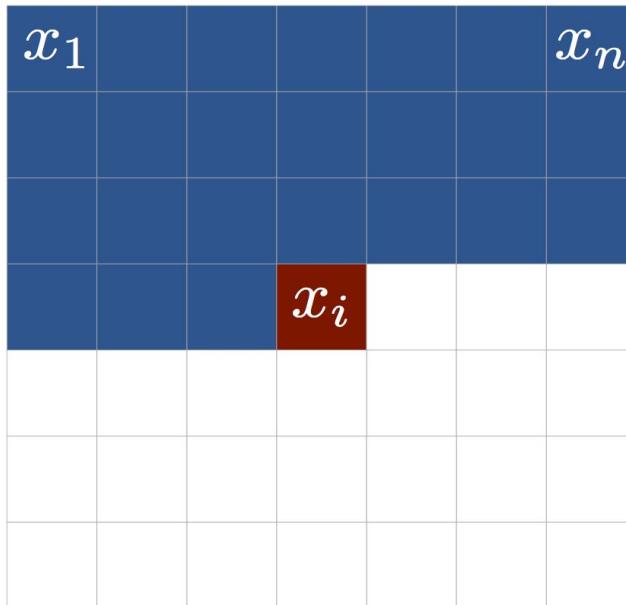
<https://giphy.com/gifs/television-13ep0e3Z06gHba>



Group-by-group

Reed et al. "Parallel Multiscale Autoregressive Density Estimation."

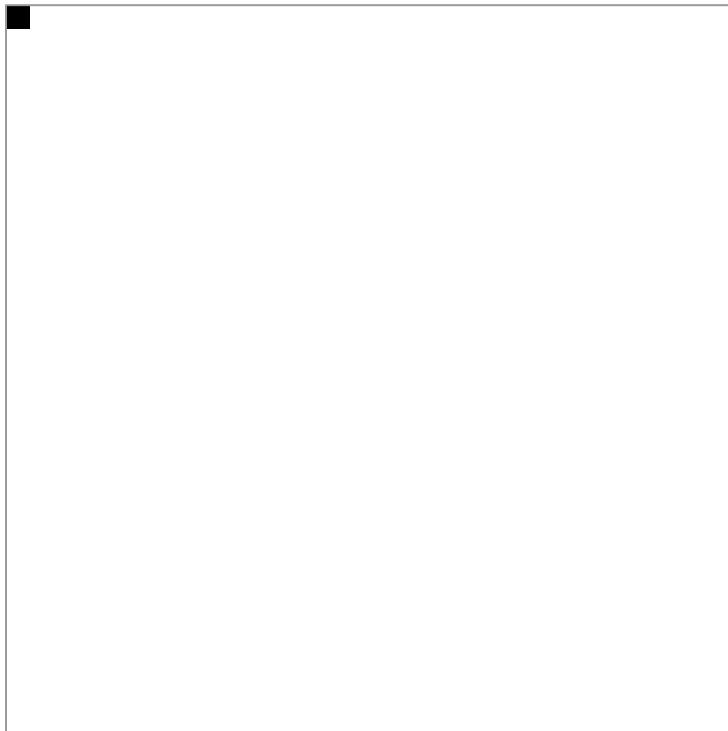
PixelRNN - Model



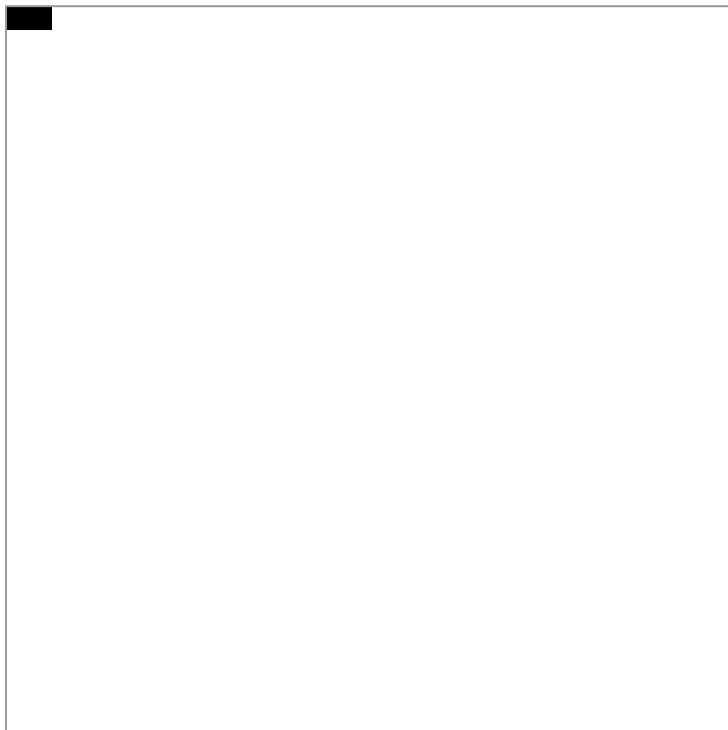
$$p(\mathbf{x}) = \prod_{i=1}^{n^2} p(x_i | x_1, \dots, x_{i-1})$$

- Fully visible
- Similar to language models with RNNs
- Model pixels with Softmax

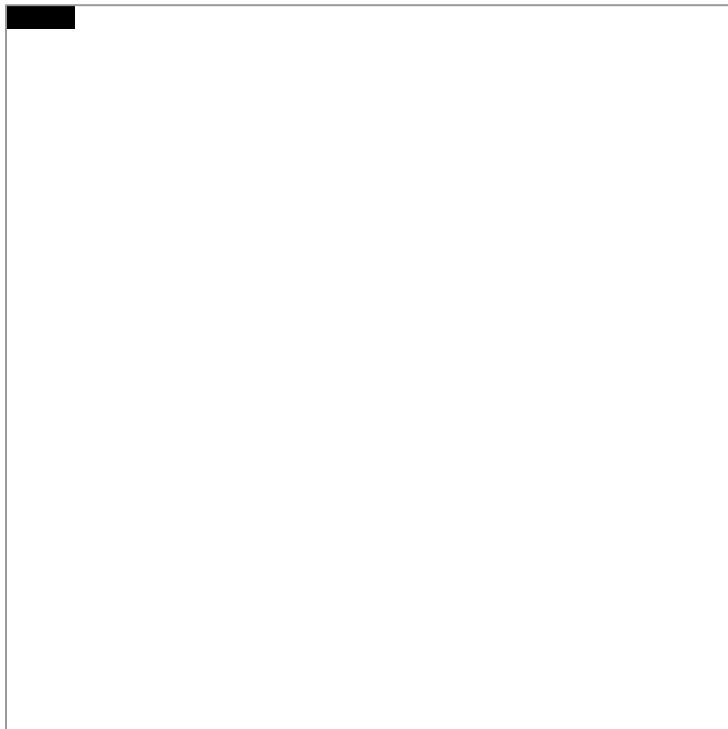
Softmax Sampling



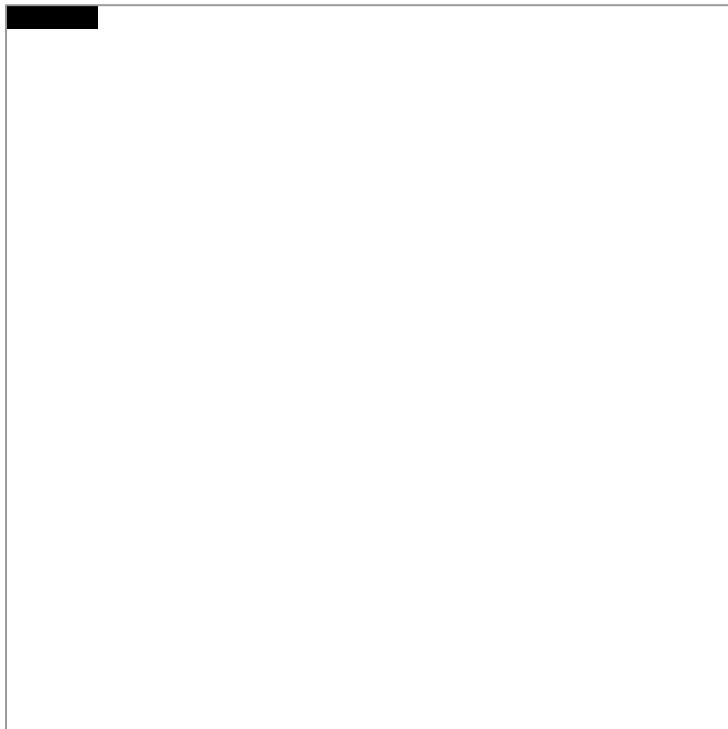
Softmax Sampling



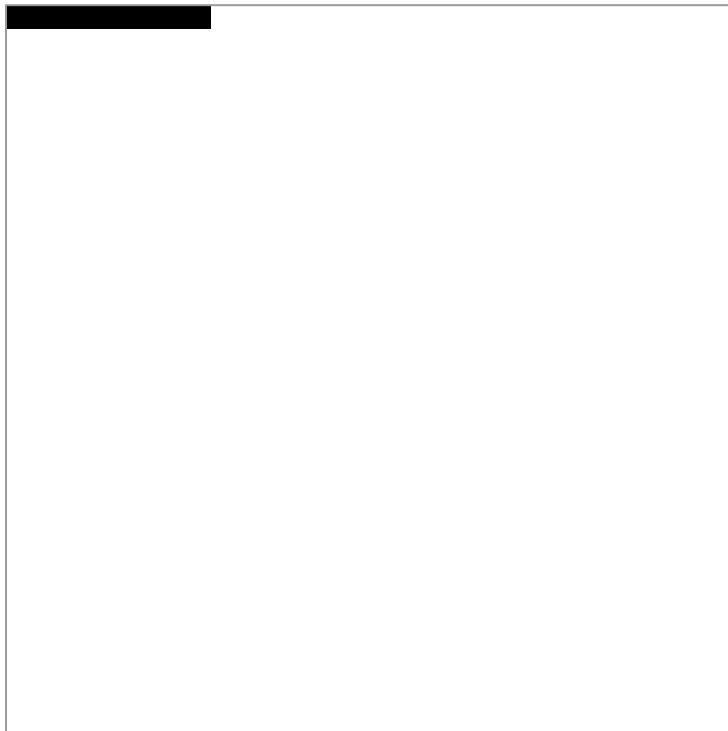
Softmax Sampling



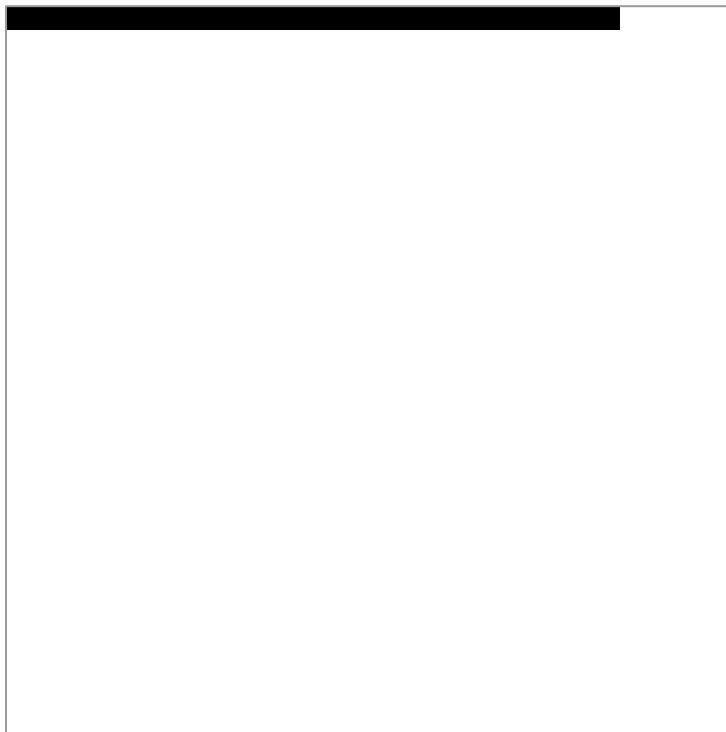
Softmax Sampling



Softmax Sampling



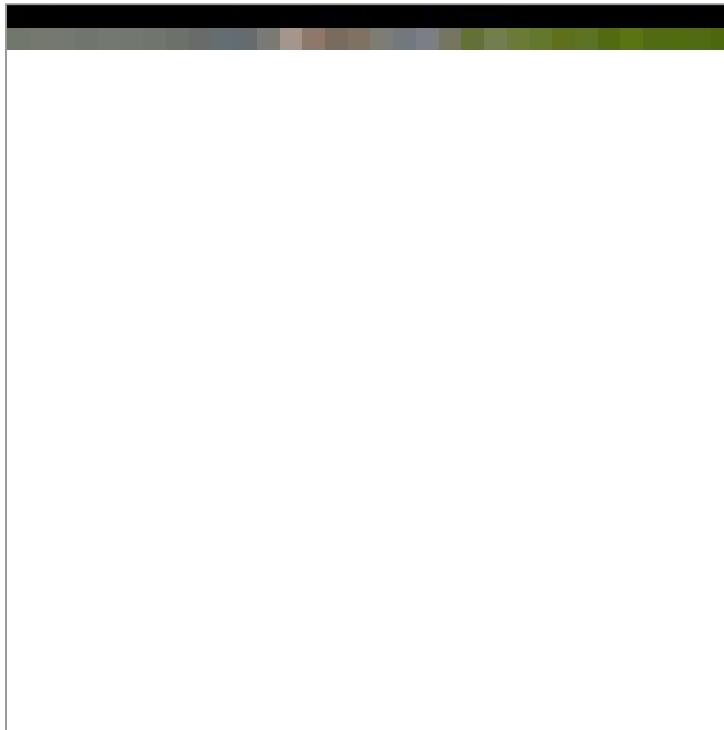
Softmax Sampling



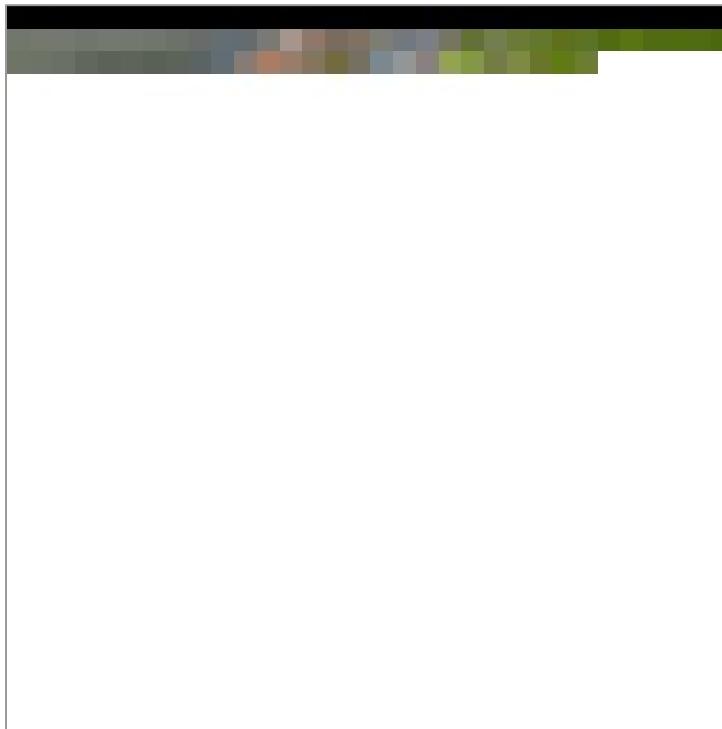
Softmax Sampling



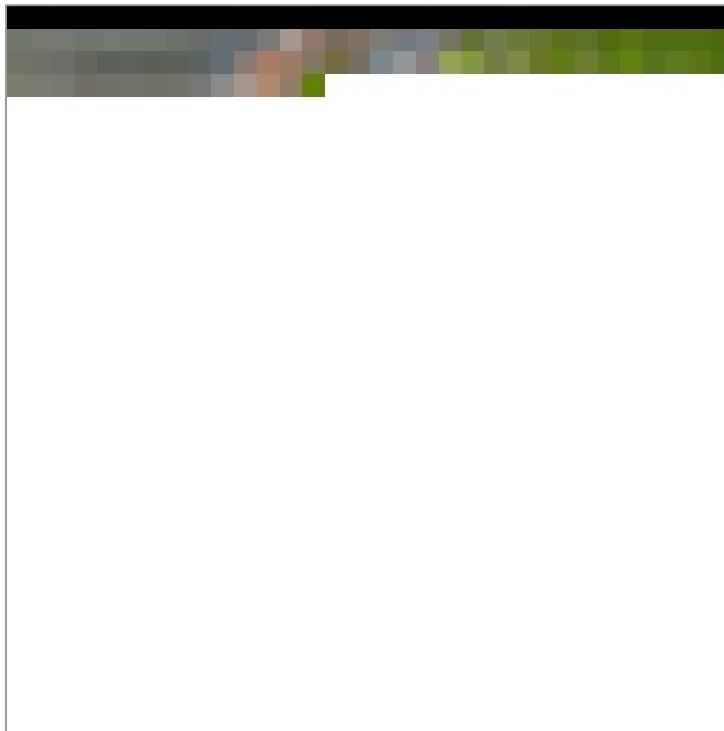
Softmax Sampling



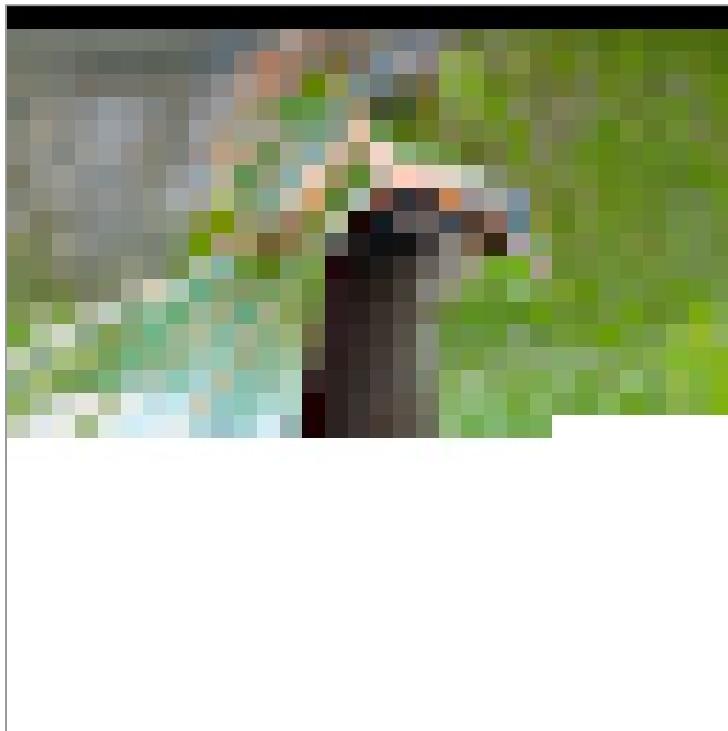
Softmax Sampling



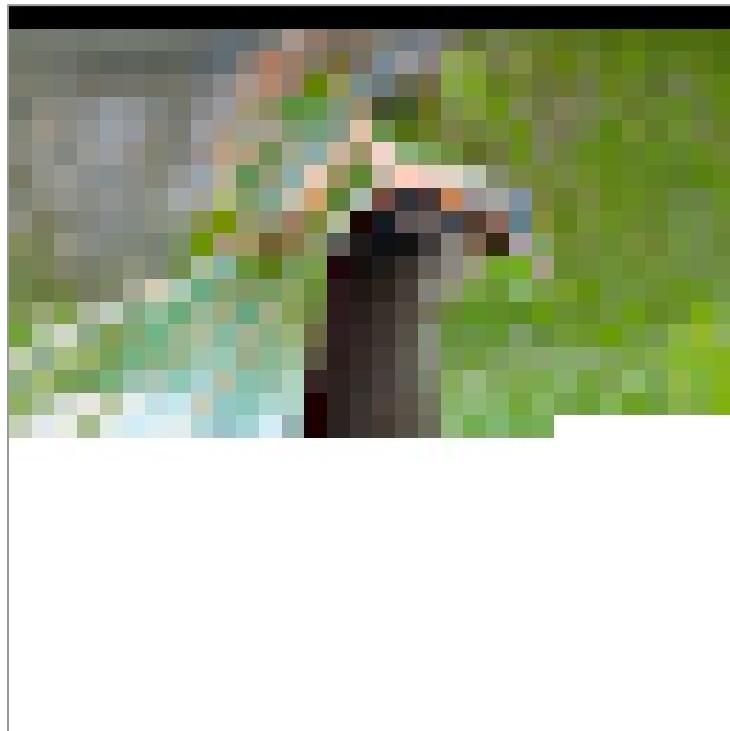
Softmax Sampling



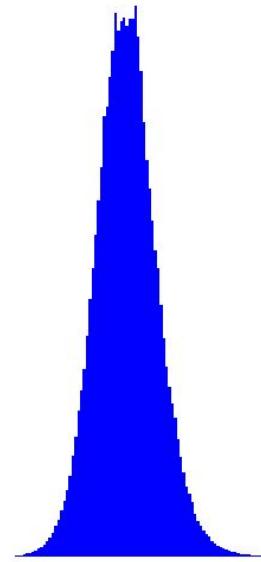
Softmax Sampling



Softmax Sampling

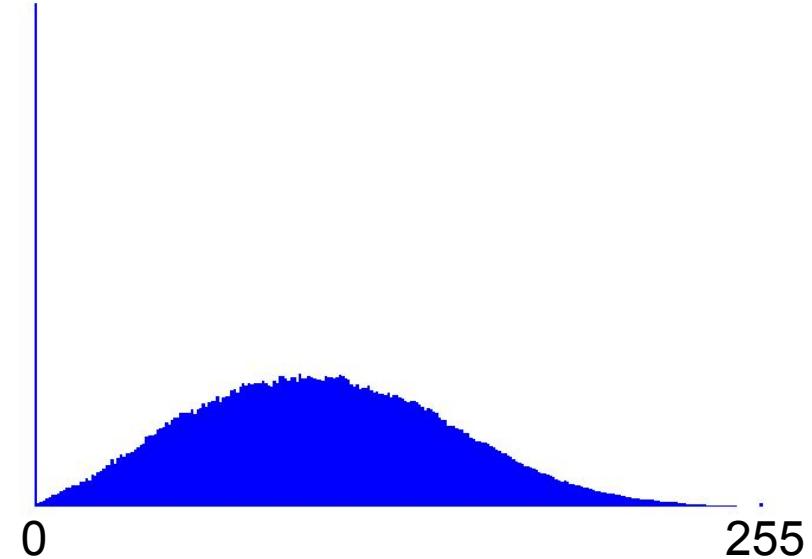
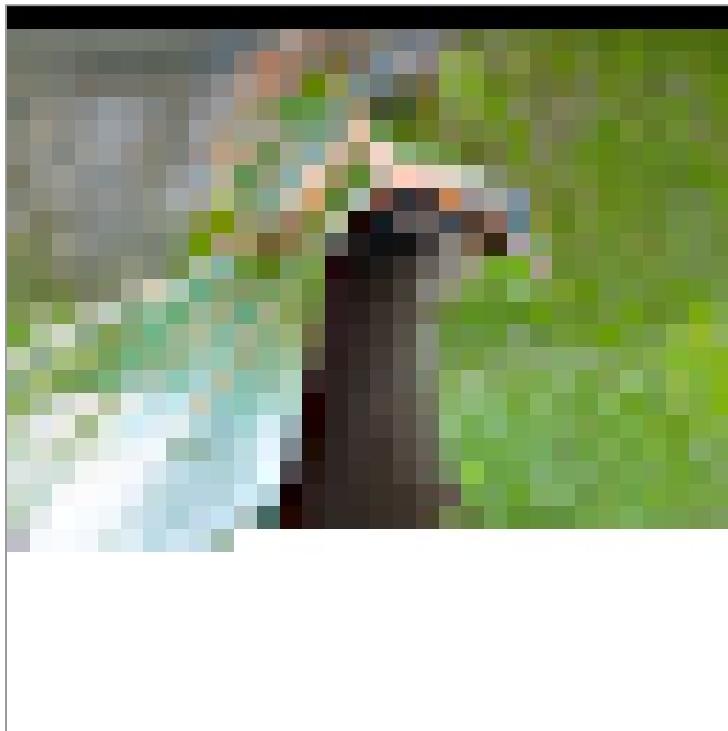


0

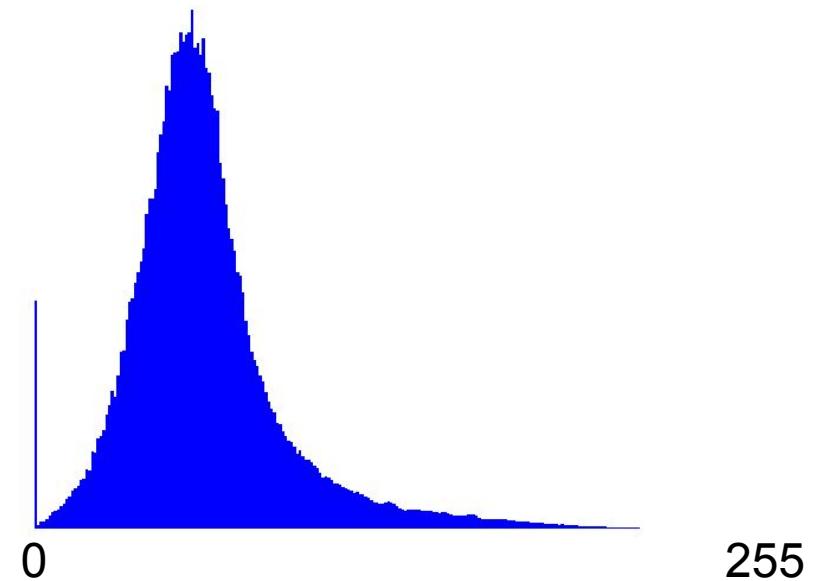
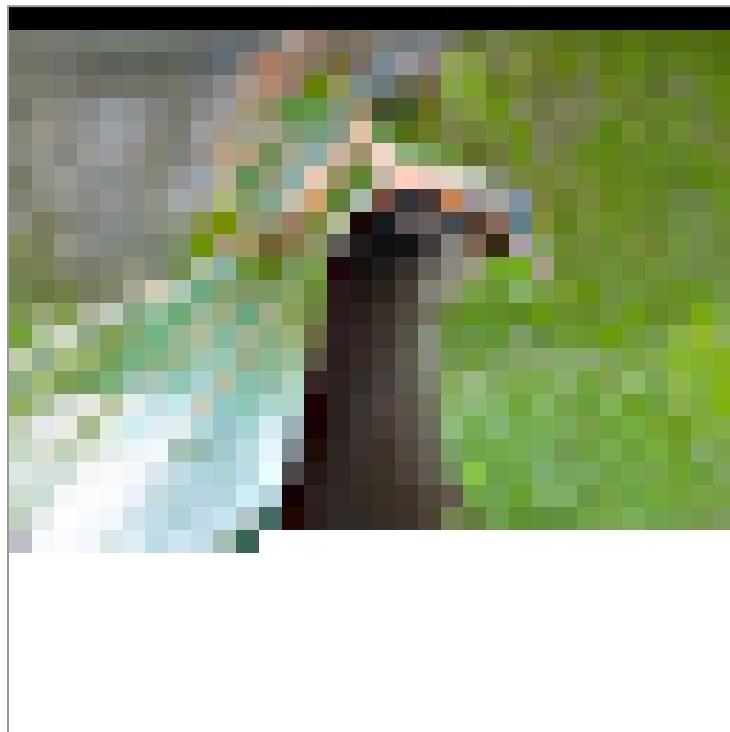


255

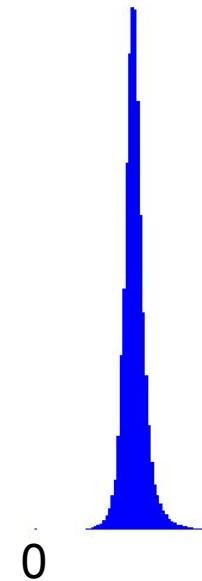
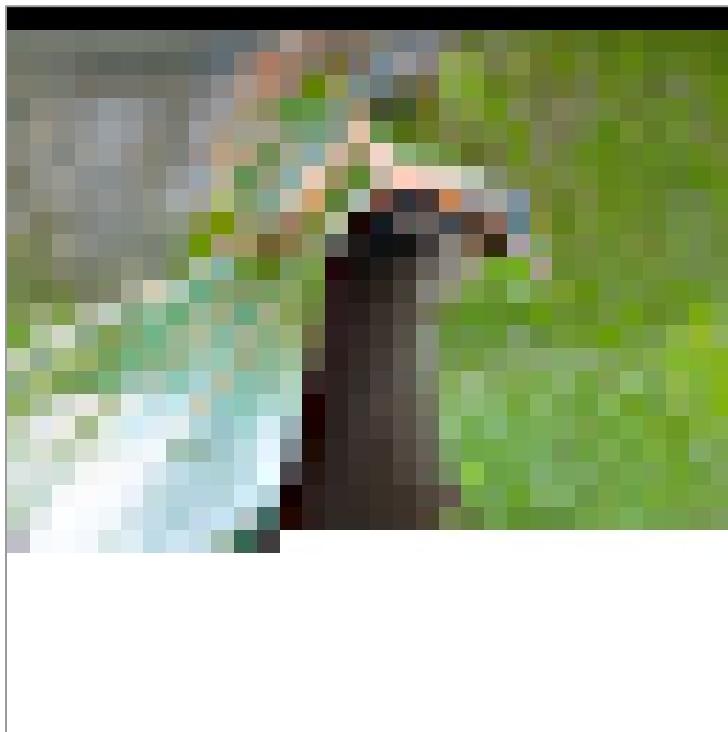
Softmax Sampling



Softmax Sampling



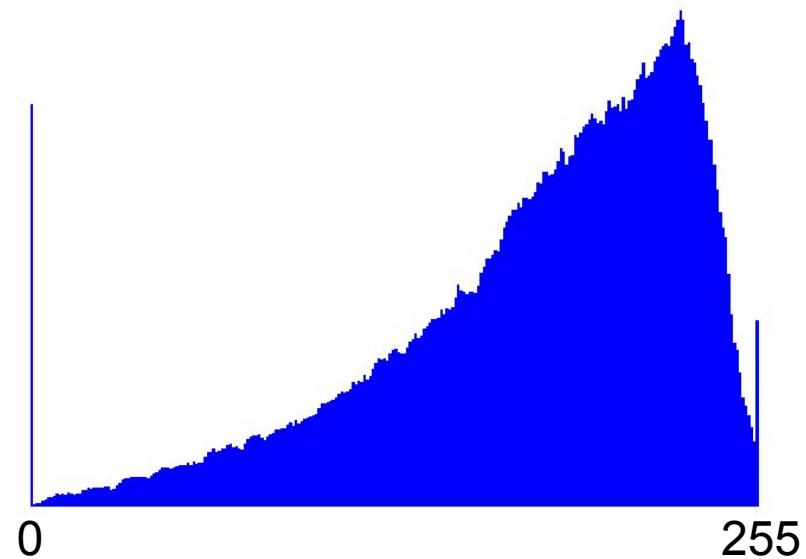
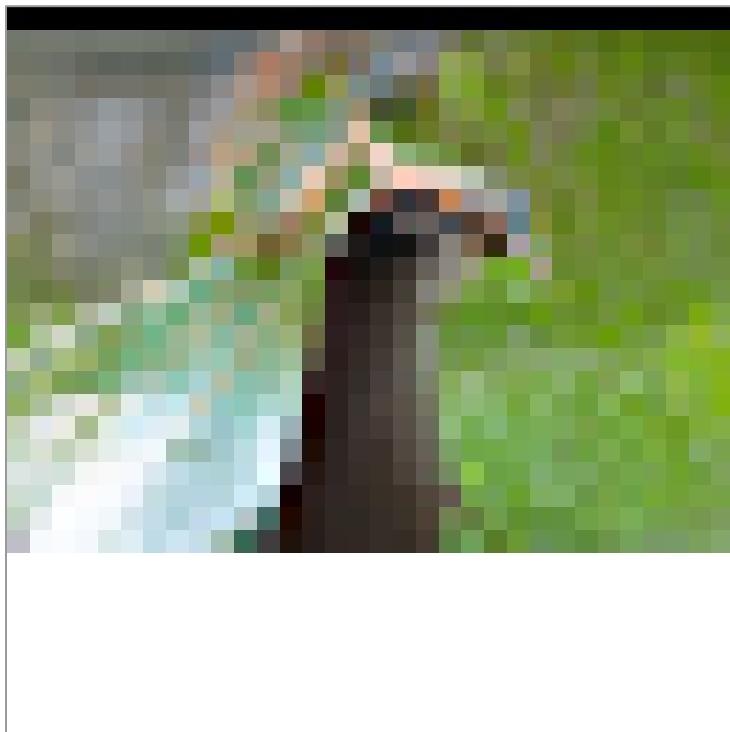
Softmax Sampling



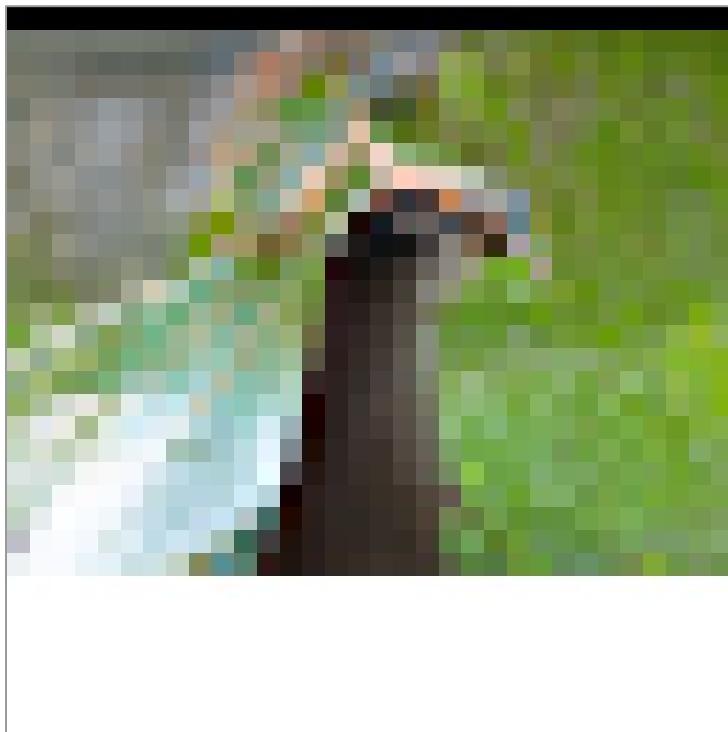
255

94

Softmax Sampling

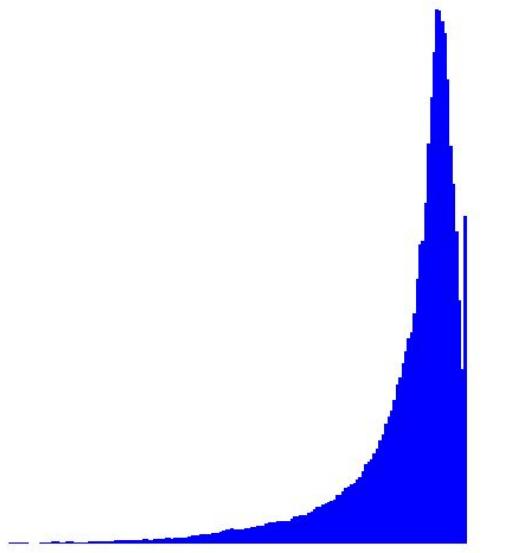


Softmax Sampling

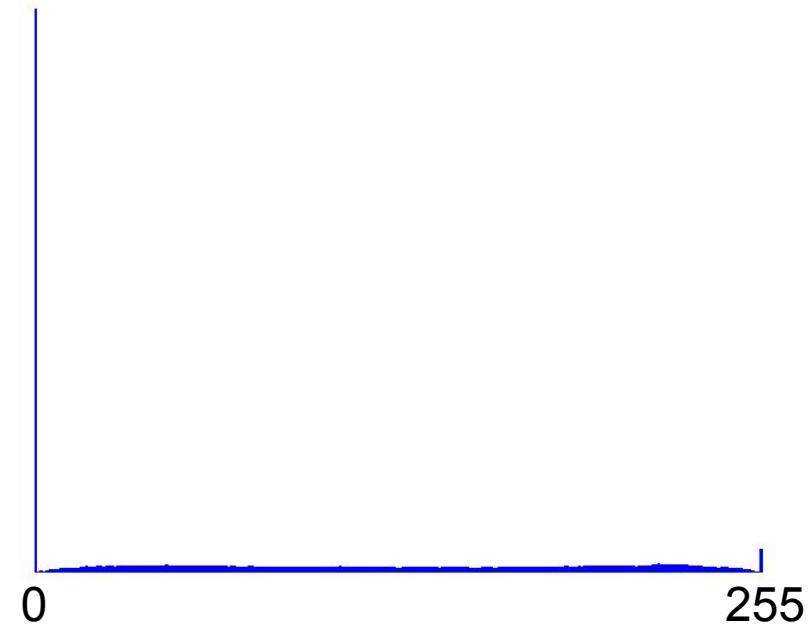
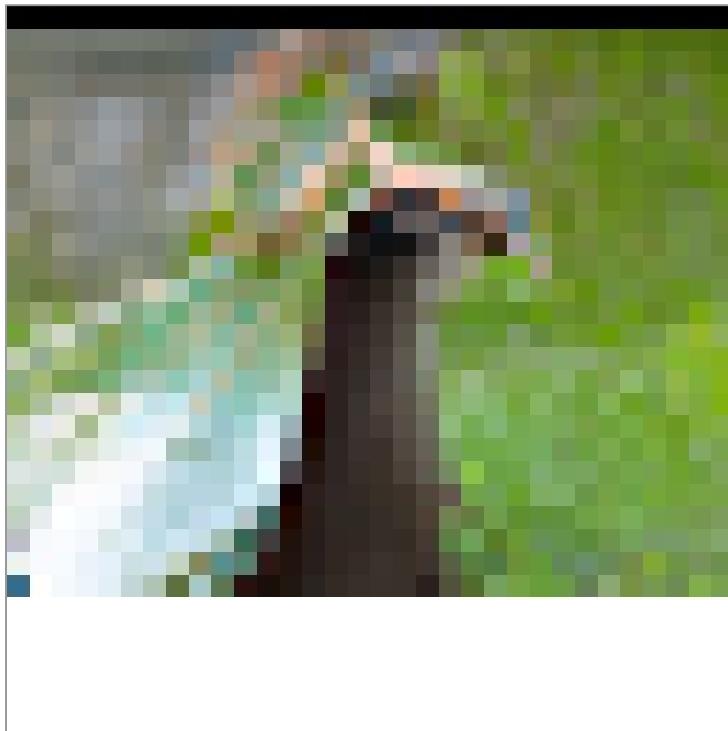


0

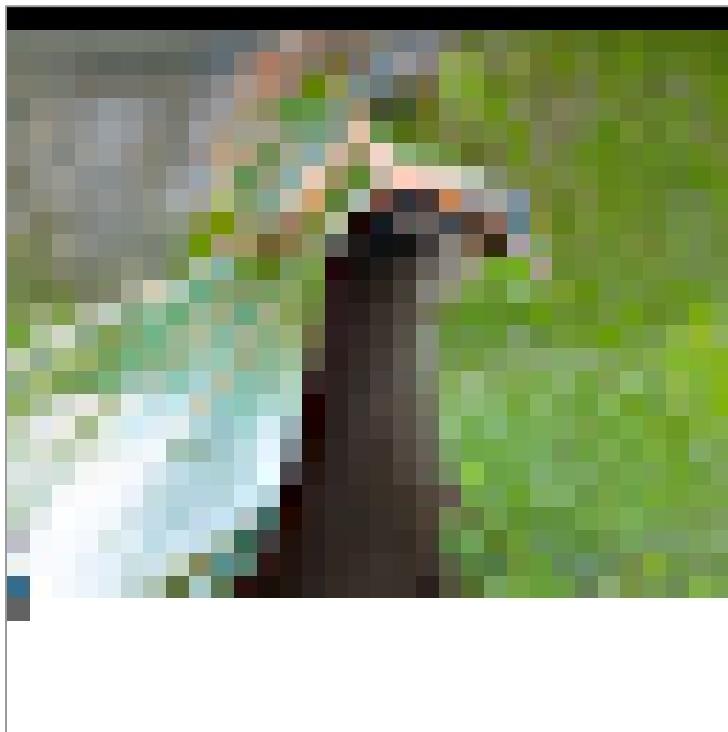
255



Softmax Sampling



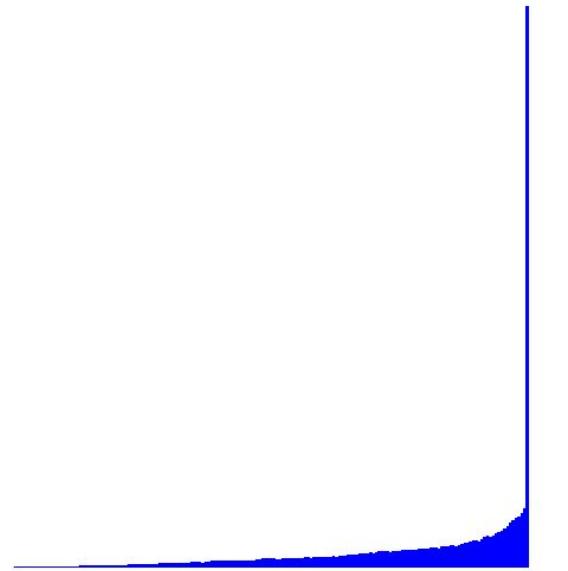
Softmax Sampling



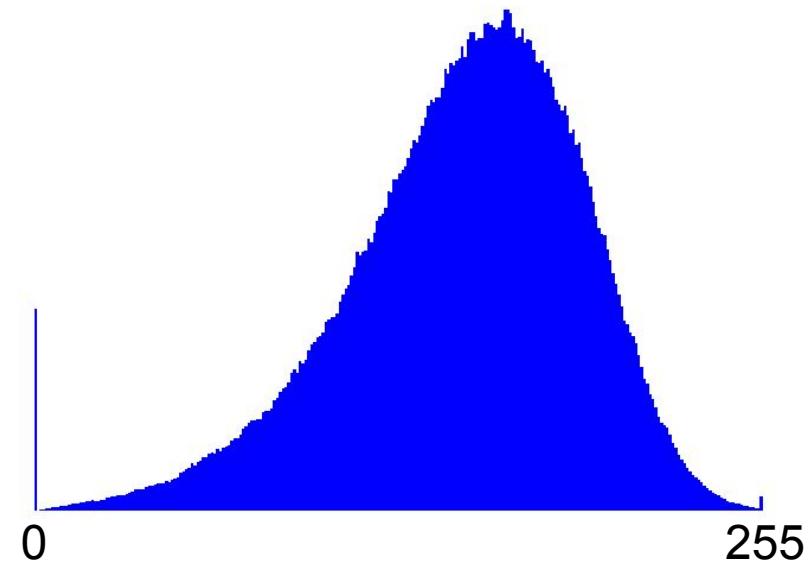
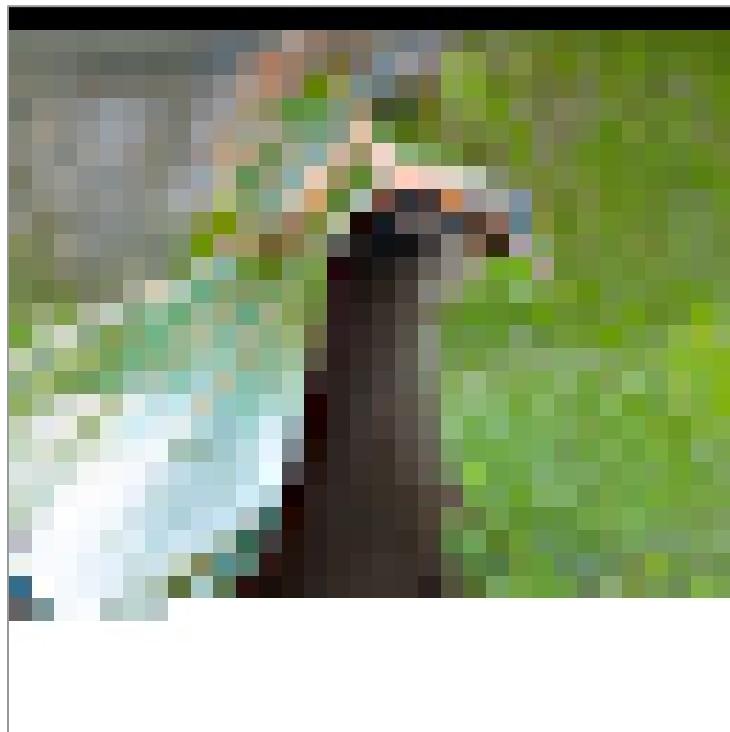
0

255

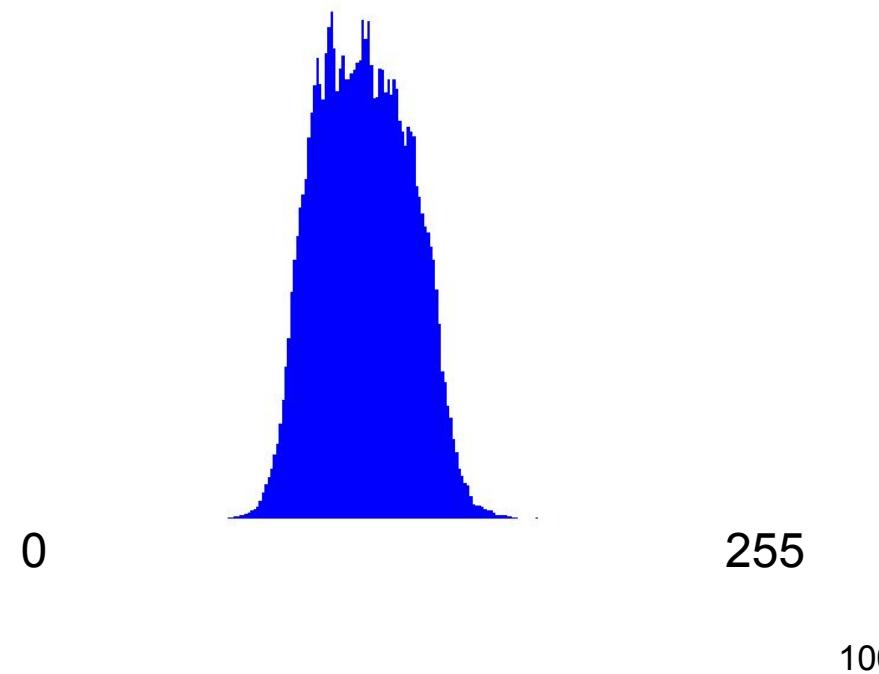
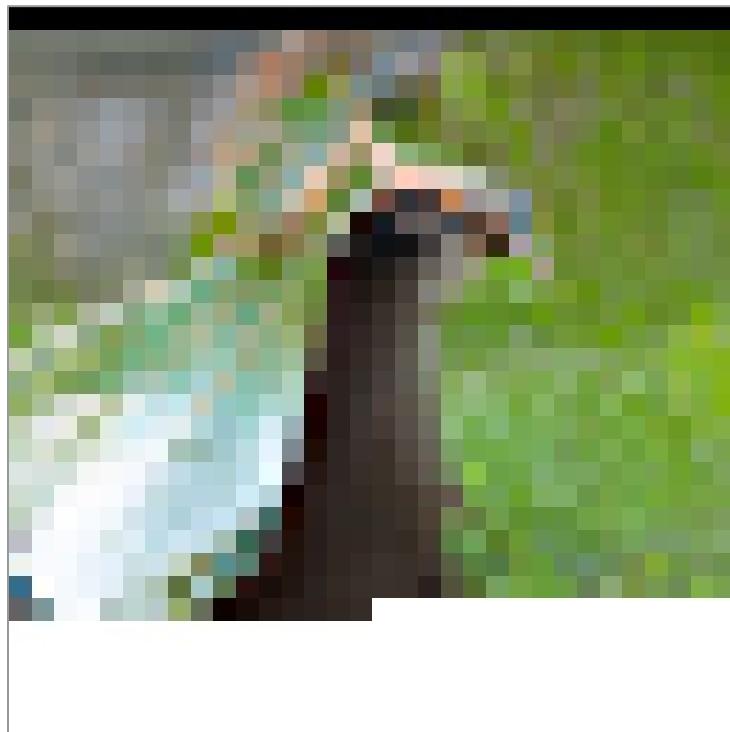
98



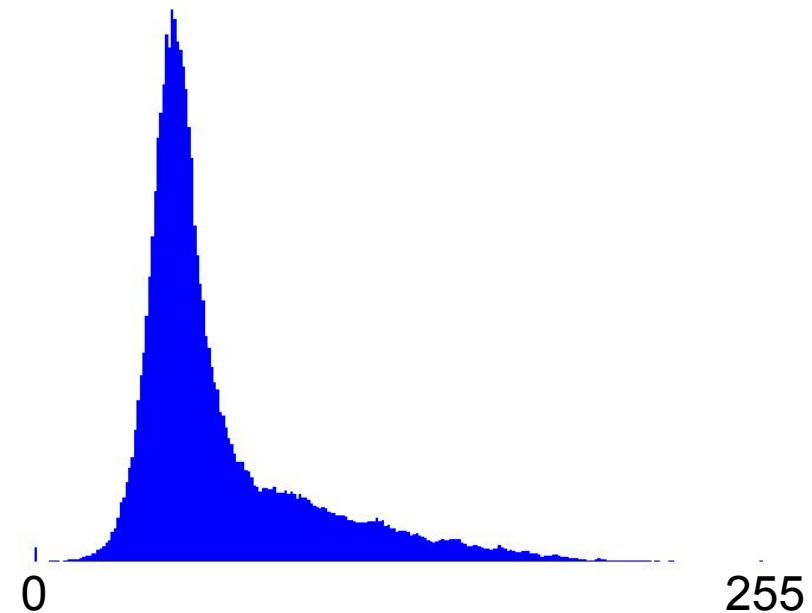
Softmax Sampling



Softmax Sampling

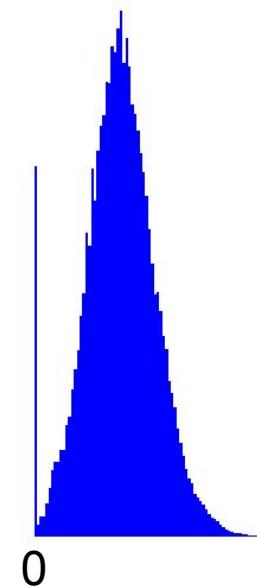


Softmax Sampling



101

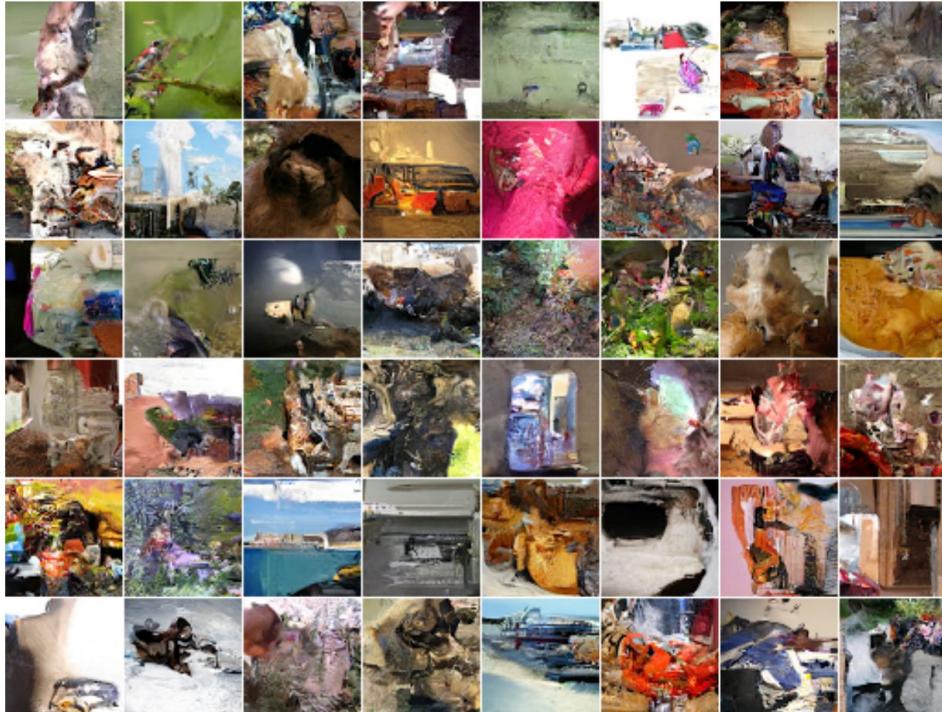
Softmax Sampling



102

Pixel RNN [van den Oord, Kalchbrenner, Kavukcuoglu, ICML16]

Sequence of Words == Sequence of Pixels

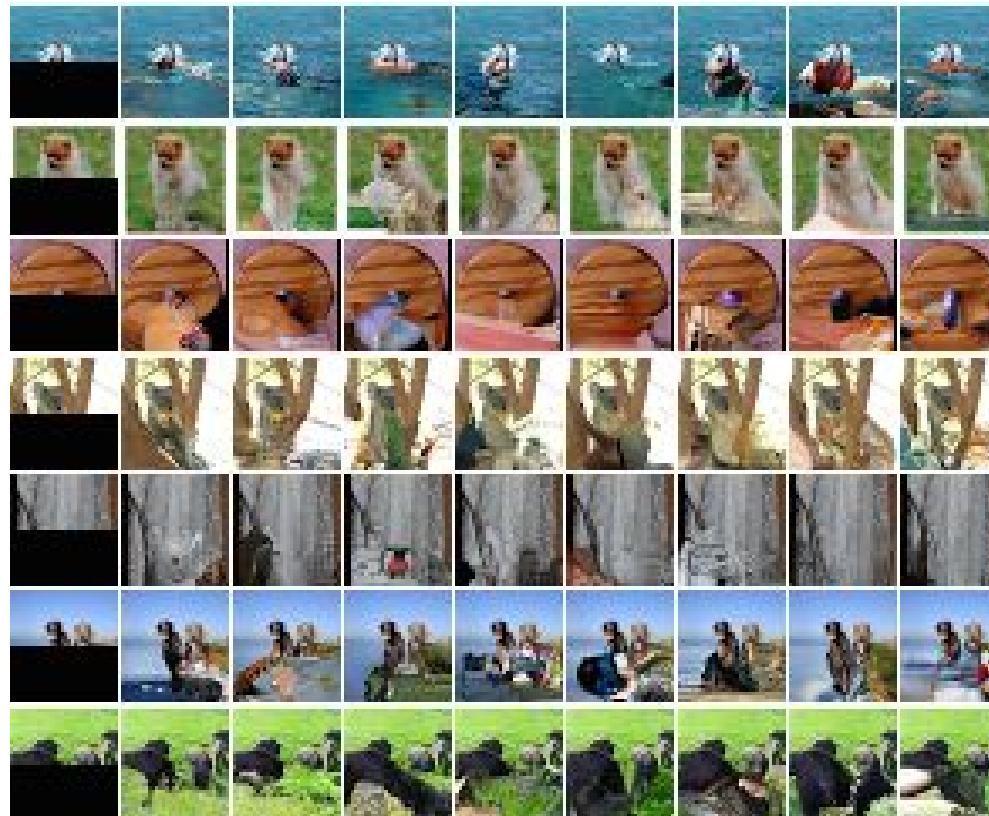


occluded



occluded

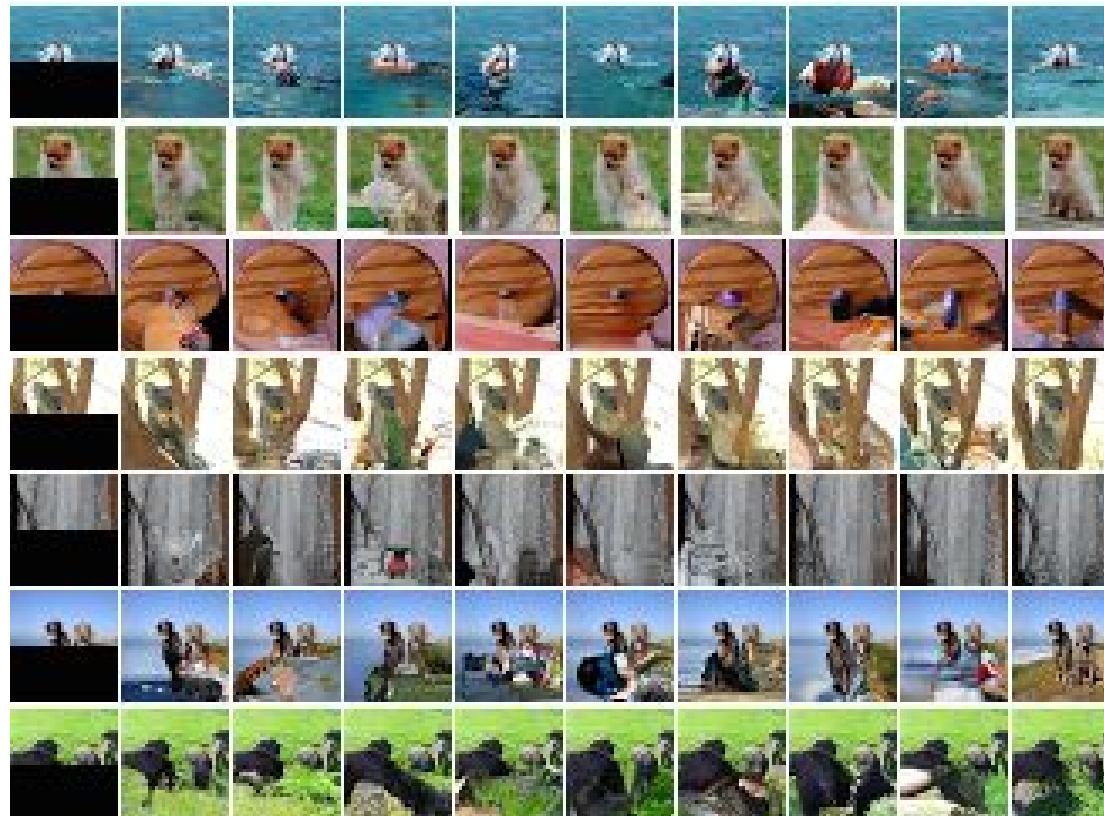
completions



occluded

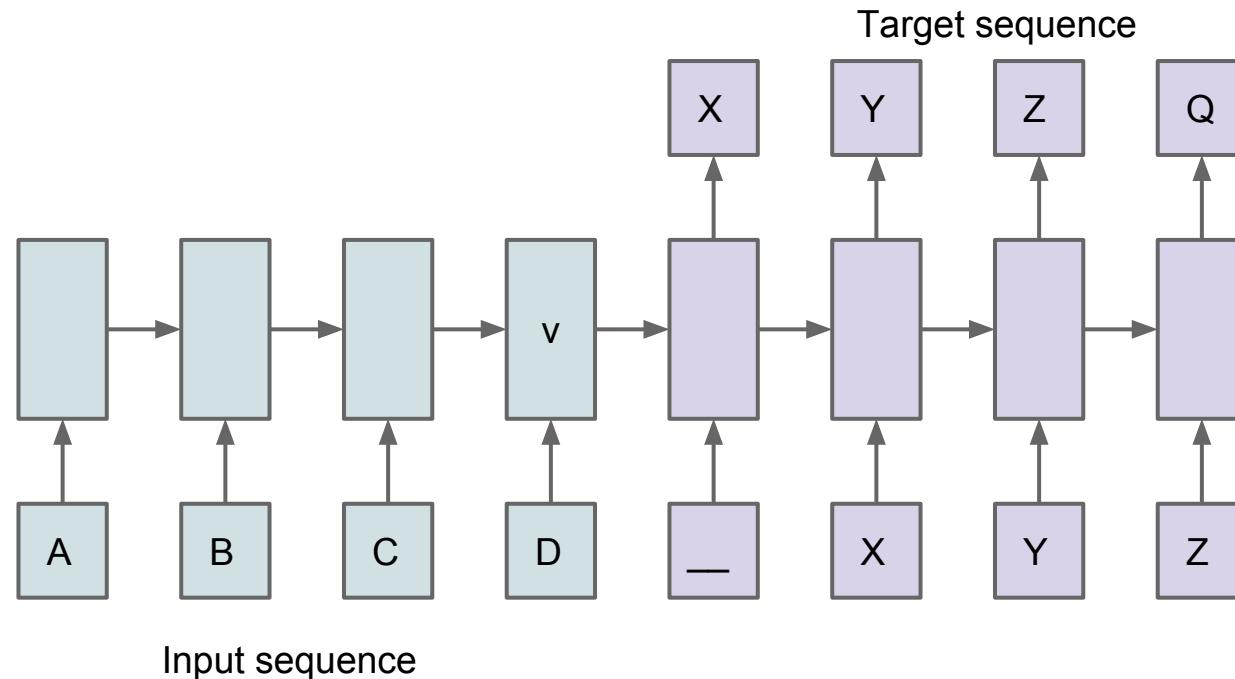
completions

original



Part IV.A: Sequence-to-Sequence & Applications

Main idea



$$P(y_1, \dots, y_{T'} | x_1, \dots, x_T) = \prod_{t=1}^{T'} p(y_t | v, y_1, \dots, y_{t-1})$$

Sequence-to-Sequence



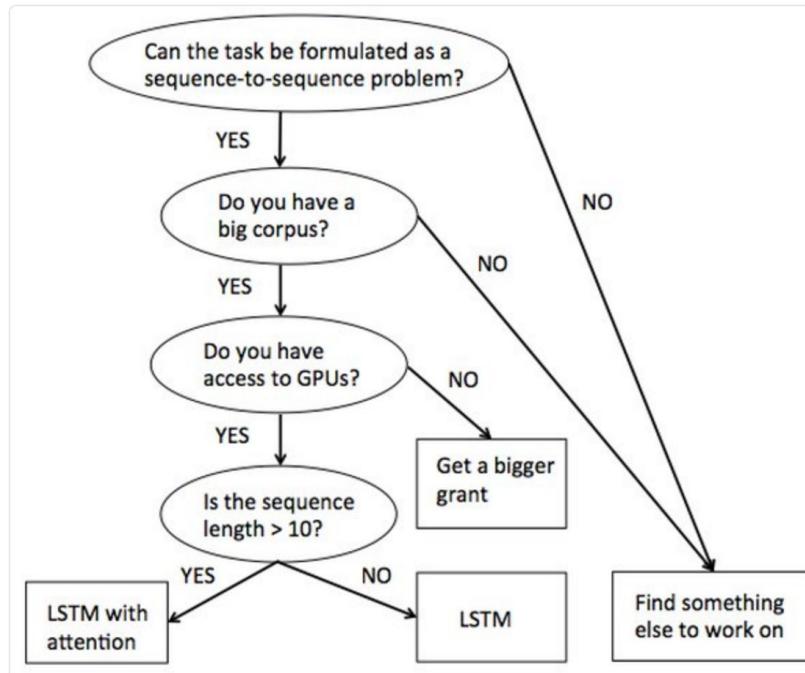
Isabelle Augenstein

@IAugenstein



Follow

Brainstorming at [@uclmr](#) : how to recommend "what method works best for NLP task X"? Our baseline is ready [#NLProc](#)



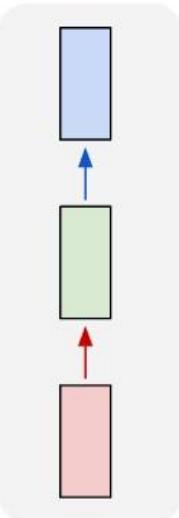
Sequence-to-Sequence

1. MT [Kalchbrenner et al, EMNLP 2013][Cho et al, EMLP 2014]**[Sutskever & Vinyals & Le, NIPS 2014]**[Luong et al, ACL 2015][Bahdanau et al, ICLR 2015]
2. Image captions [Mao et al, ICLR 2015][Vinyals et al, CVPR 2015][Donahue et al, CVPR 2015][Xu et al, ICML 2015]
3. Speech [Chorowsky et al, NIPS DL 2014][Chan et al, arxiv 2015]
4. Parsing [Vinyals & Kaiser et al, NIPS 2015]
5. Dialogue [Shang et al, ACL 2015][Sordoni et al, NAACL 2015][Vinyals & Le, ICML DL 2015]
6. Video Generation [Srivastava et al, ICML 2015]
7. Geometry [Vinyals & Fortunato & Jaitly, NIPS 2015]

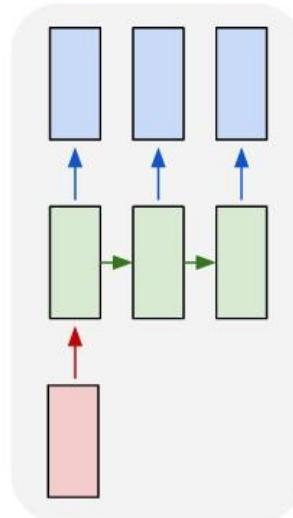
Sequence Models

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

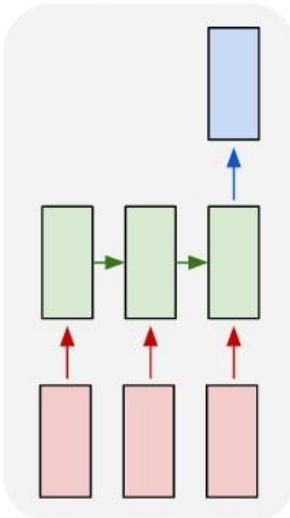
one to one



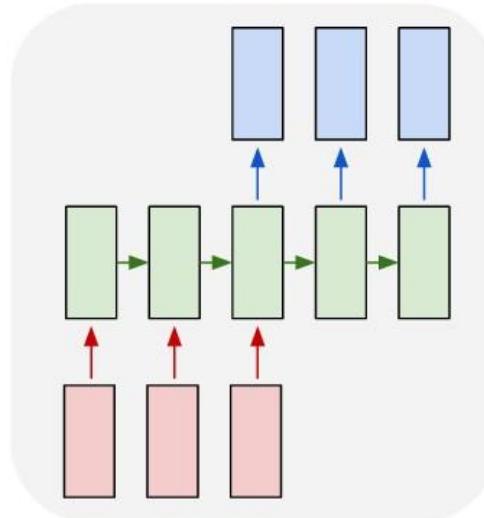
one to many



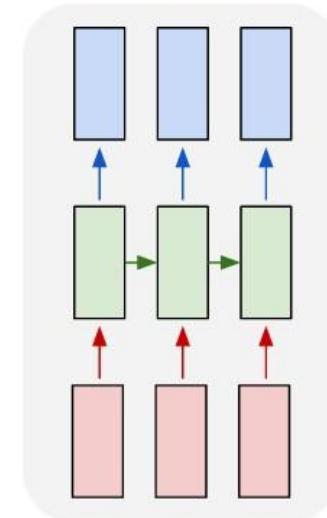
many to one



many to many



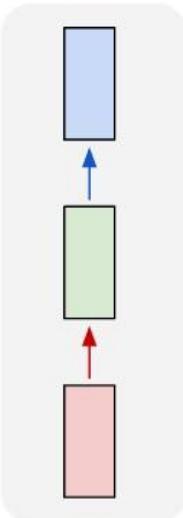
many to many



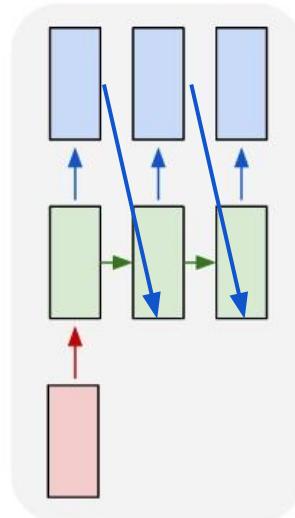
Sequence Models (Autoregressive)

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

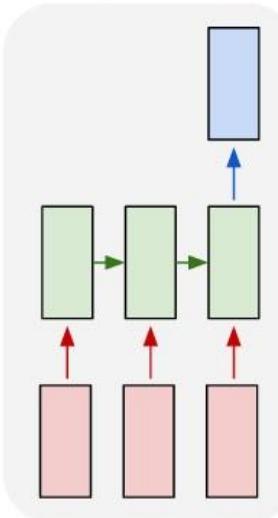
one to one



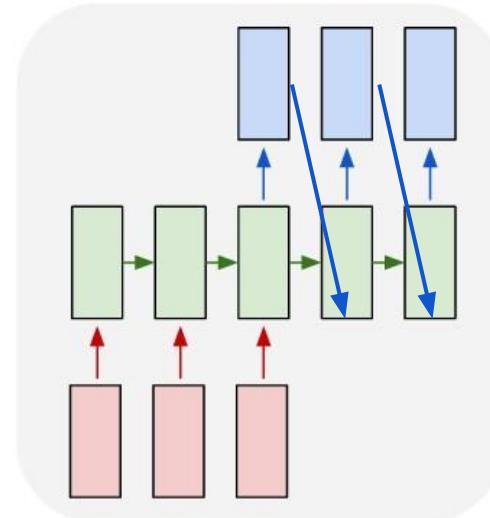
one to many



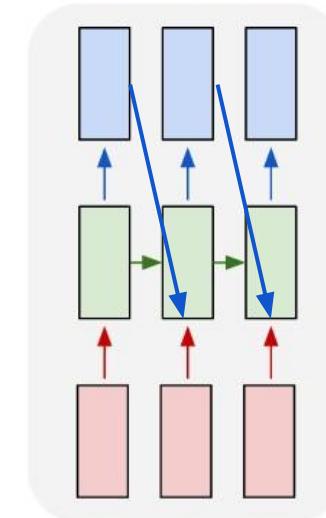
many to one



many to many



many to many



Training vs Inference

Training

$$\theta^* = \arg \max_{\theta} P(y_1, \dots, y_{T'} | x_1, \dots, x_T)$$

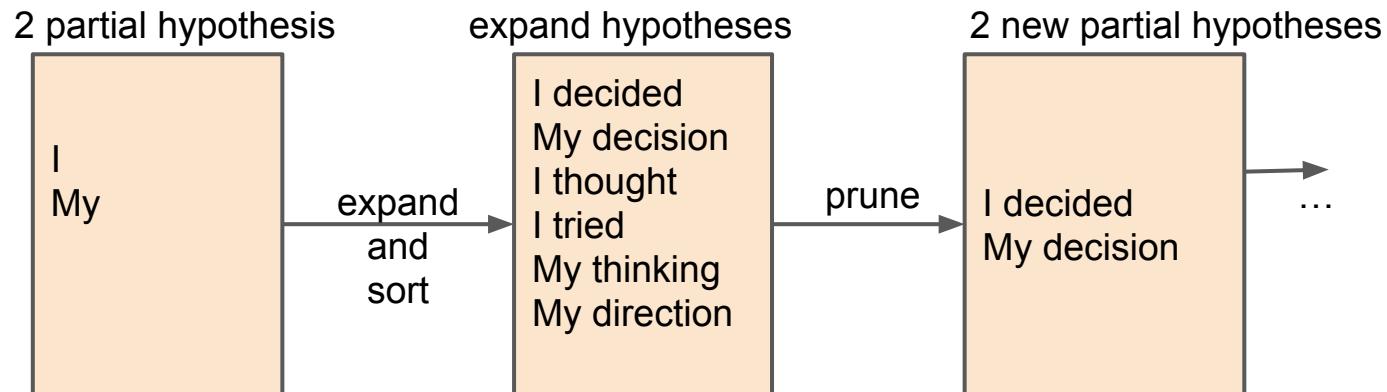
Inference

$$Y^* = \arg \max_Y P(Y | x_1, \dots, x_T)$$

- Used beam search decoding over Y
- Small beam (~20) works fine

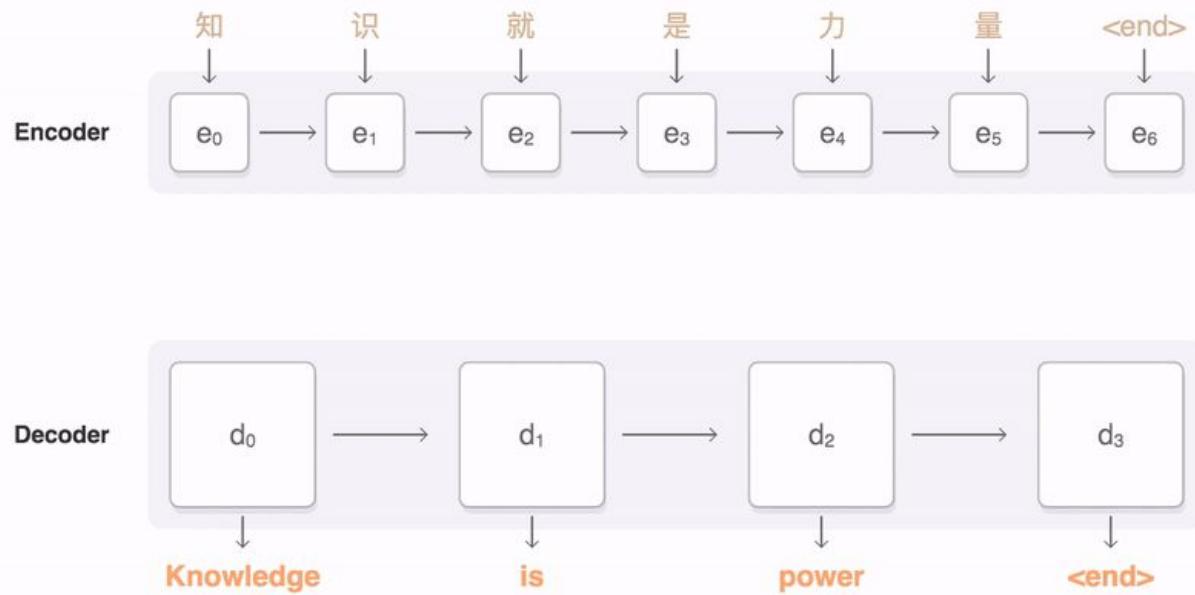
Decoding in a Nutshell (Beam Size 2)

$$y^* = \arg \max_{y_1, \dots, y_{T'}} P(y_1, \dots, y_{T'} | x_1, \dots, x_T)$$

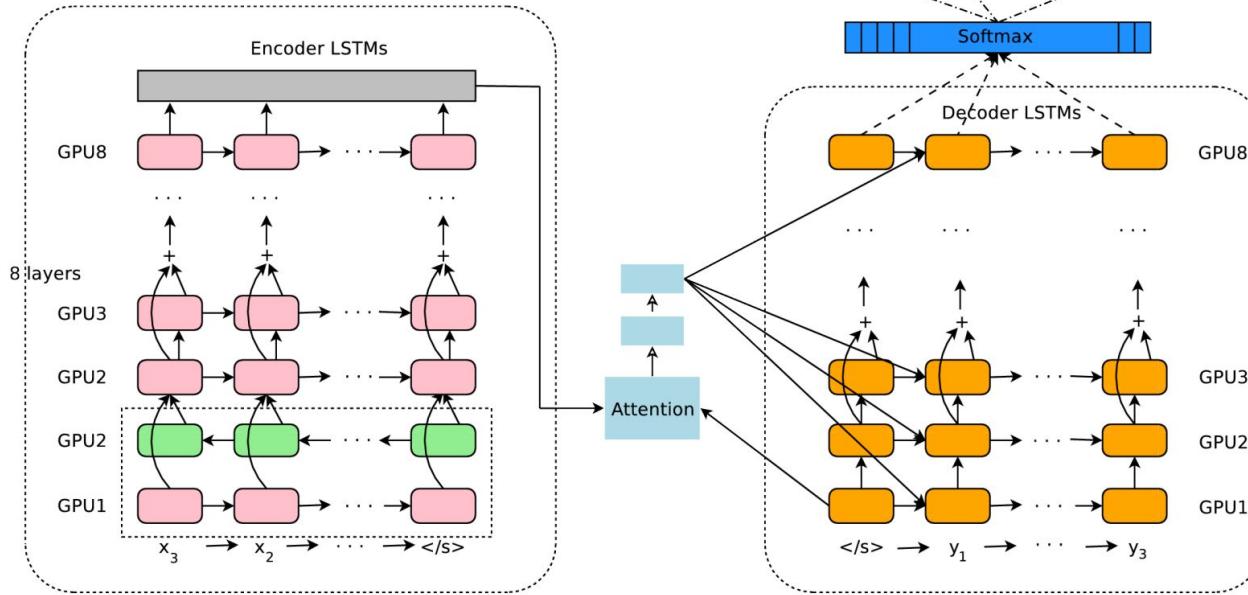


Google Neural MT

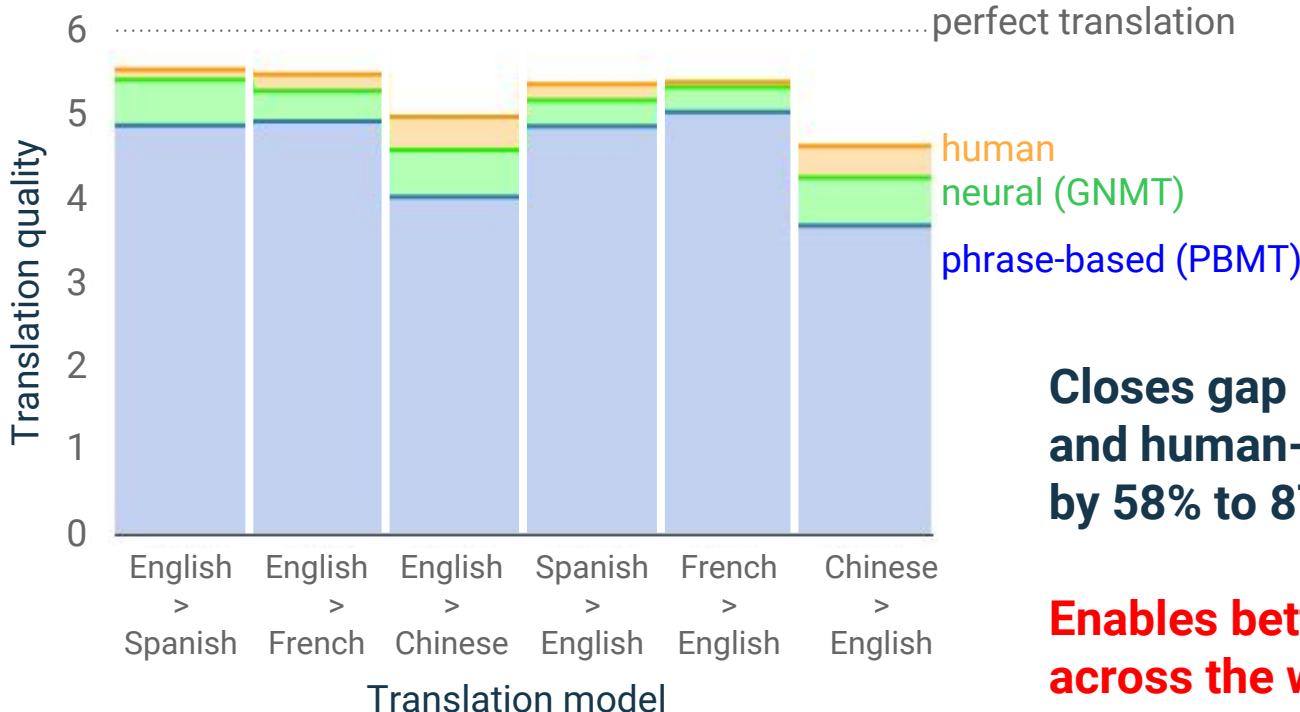
Wu et al, 2016 (Kalchbrenner et al, 2013; Sutskever et al, 2014;
Cho et al, 2014; Bahdanau et al, 2014; ...)



Google Neural MT



Google Neural MT



**Closes gap between old system
and human-quality translation
by 58% to 87%**

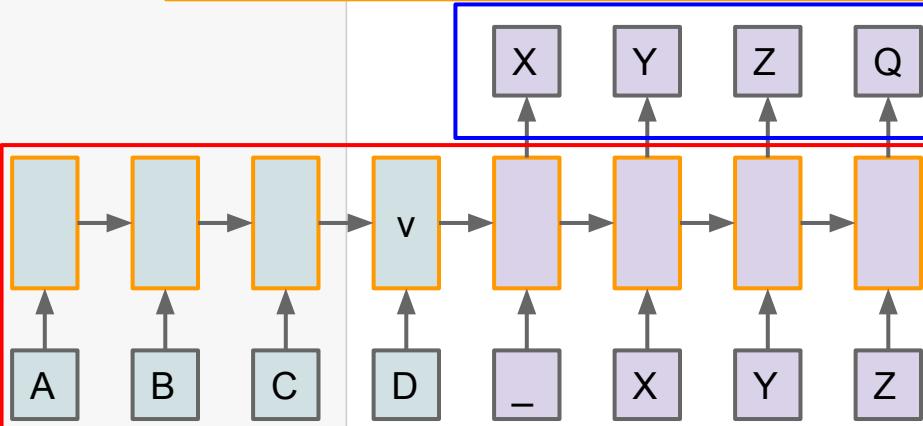
**Enables better communication
across the world**

Code

```
# words and targets are placeholders for [batch_size, num_steps]
# tensor of word and target ids
words = tf.placeholder(tf.int64, name='words')
targets = tf.placeholder(tf.int64, name='targets')
```

```
def model(batch_size, num_steps):
    output = [tf.zeros([batch_size, state_size])] * 4
    state = [tf.zeros([batch_size, state_size])] * 4
    preds = []
    cost = 0.0
    for i in range(num_steps):
        # Get the embedding for words
        embedding = tf.nn.embedding_lookup(embedding_params, words[:, i])
        # Run the LSTM cells
        output[0], state[0] = lstm[0](embedding, output[0], state[0])
        for d in range(1, 4):
            output[d], state[d] = lstm[d](output[d-1], output[d], state[d])
        # Get the logits
        logits = tf.matmul(output[-1], sm_w) + sm_b
        # Get the softmax predictions
        preds.append(tf.nn.softmax(logits))
        # Cost per step
        cost = cost + tf.reduce_mean(
            tf.nn.sparse_softmax_cross_entropy_with_logits(logits,
                                                          targets[:, i]))
    # Average cost across time steps
    cost = cost / np.float32(num_steps)
    return preds, cost
```

```
class LSTMCell(object):
    def __init__(self, state_size):
        self.state_size = state_size
        self.W_f = tf.Variable(self.initializer())
        self.W_i = tf.Variable(self.initializer())
        self.W_o = tf.Variable(self.initializer())
        self.W_C = tf.Variable(self.initializer())
        self.b_f = tf.Variable(tf.zeros((state_size)))
        self.b_i = tf.Variable(tf.zeros((state_size)))
        self.b_o = tf.Variable(tf.zeros((state_size)))
        self.b_C = tf.Variable(tf.zeros((state_size)))
    def __call__(self, x_t, h_t, c_t):
        X = tf.concat(1, [h_t, x_t])
        f_t = tf.sigmoid(tf.matmul(X, self.W_f) + self.b_f)
        i_t = tf.sigmoid(tf.matmul(X, self.W_i) + self.b_i)
        o_t = tf.sigmoid(tf.matmul(X, self.W_o) + self.b_o)
        Ctilde_t = tf.tanh(tf.matmul(X, self.W_C) + self.b_C)
        C_t = f_t * C_t + i_t * Ctilde_t
        h_t = o_t * tf.tanh(C_t)
        return h_t, C_t
    def initializer(self):
        return tf.random_uniform([2*self.state_size, self.state_size],
                               -0.1, 0.1)
```





Human: A young girl asleep on the sofa cuddling a stuffed bear.

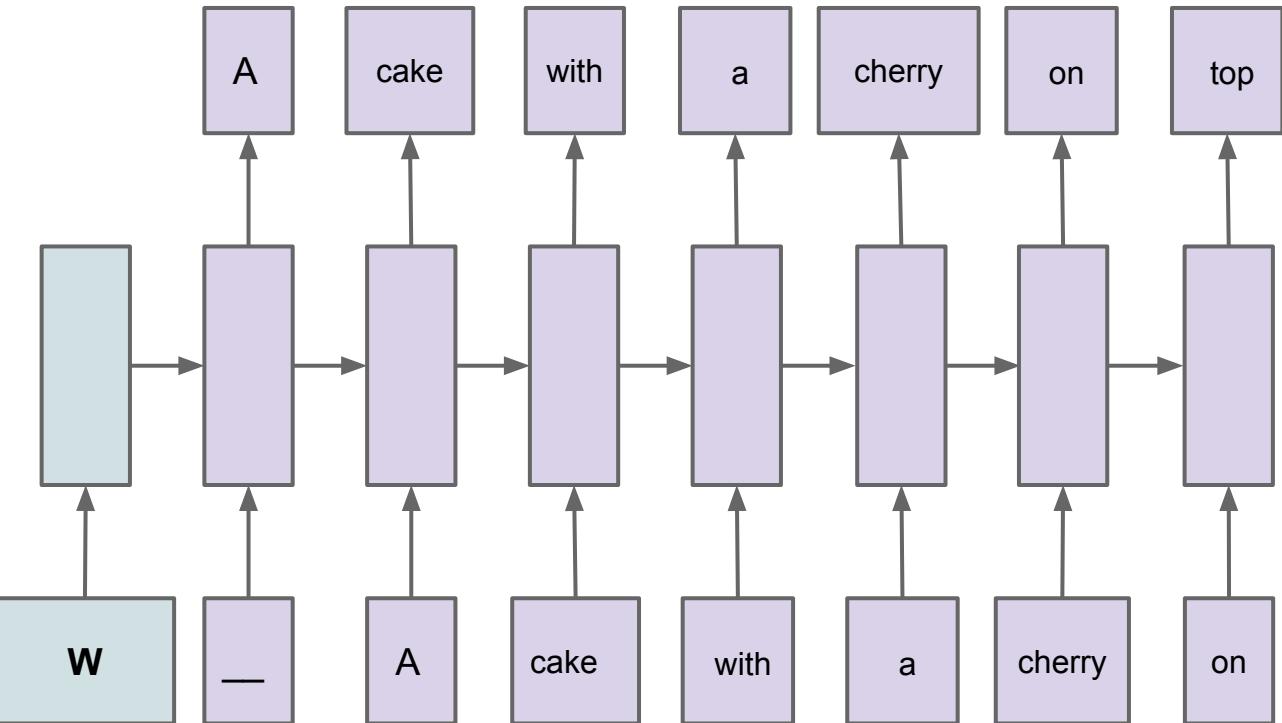
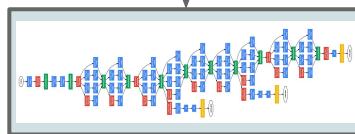
NIC: A close up of a child holding a stuffed animal.

NIC: A baby is asleep next to a teddy bear.

How to do Image Captions?

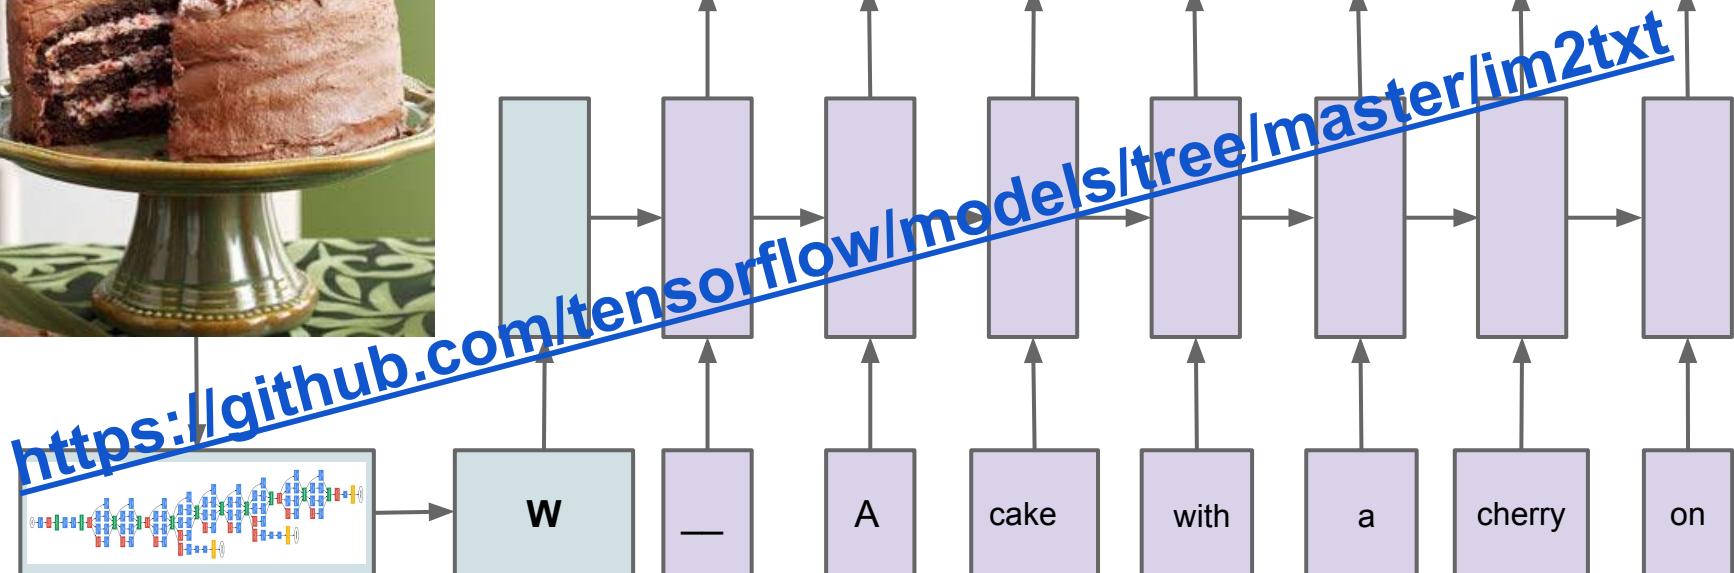
P(English | French)

Captioning



$$\theta^* = \arg \max_{\theta} p(S|I)$$

Captioning



$$\theta^* = \arg \max_{\theta} p(S|I)$$



Human: A close up of two bananas with bottles in the background.

BestModel: A bunch of bananas and a bottle of wine.

InitialModel: A close up of a plate of food on a table.



Human: A view of inside of a car where a cat is laying down.

BestModel: A cat sitting on top of a black car.

InitialModel: A dog sitting in the passenger seat of a car.



Human: A brown dog laying in a red wicker bed.

BestModel: A small dog is sitting on a chair.

InitialModel: A large brown dog laying on top of a couch.



Human: A man outside cooking with a sub in his hand.

BestModel: A man is holding a sandwich in his hand.

InitialModel: A man cutting a cake with a knife.



Human: Someone is using a small grill to melt his sandwich.

BestModel: A person is cooking some food on a grill.

InitialModel: A pizza sitting on top of a white plate.



Human: A woman holding up a yellow banana to her face.

BestModel: A woman holding a banana up to her face.

InitialModel: A close up of a person eating a hot dog.



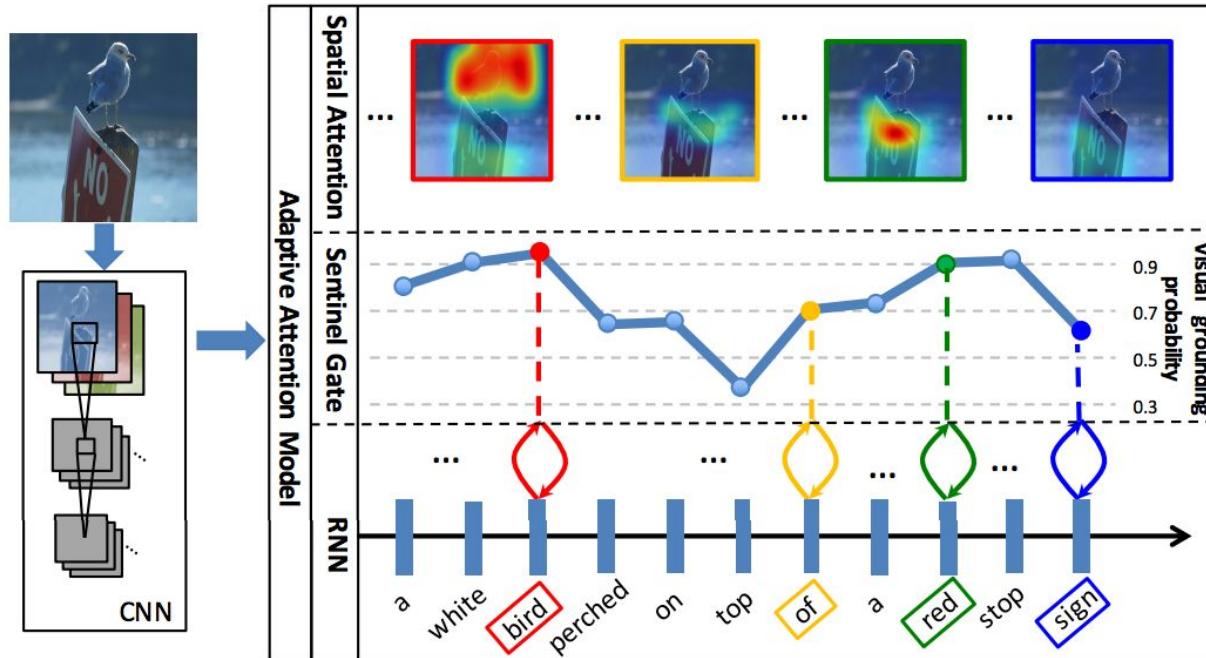
Human: A blue , yellow and red train travels across the tracks near a depot.

BestModel: A blue and yellow train traveling down train tracks.

InitialModel: A train that is sitting on the tracks.

Advanced Networks: Captioning

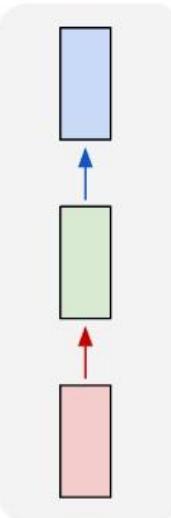
Figure from [Lu et al, 2016]



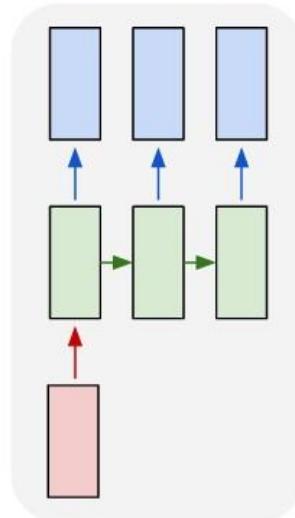
Sequence Models

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

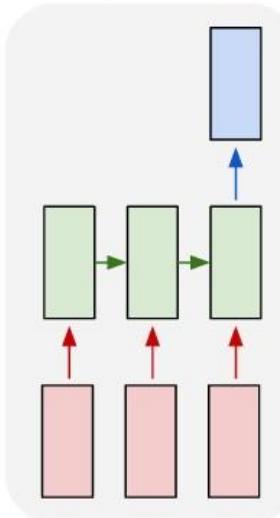
one to one



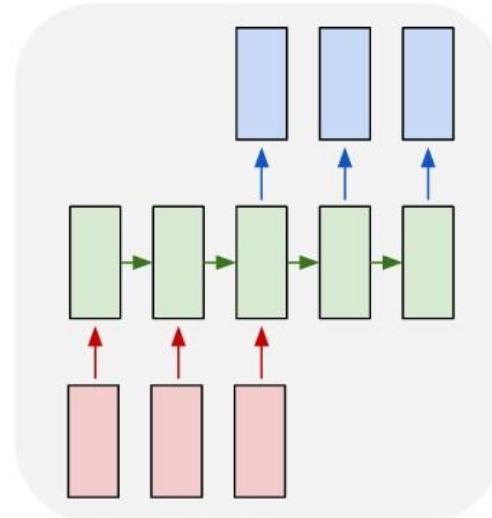
one to many



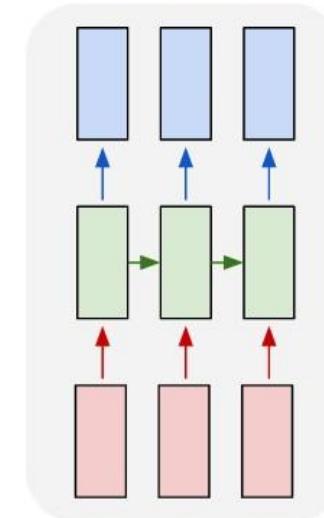
many to one



many to many



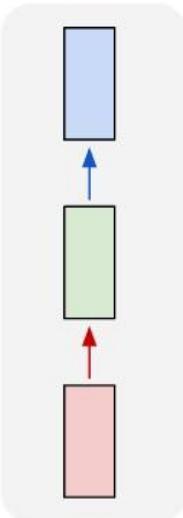
many to many



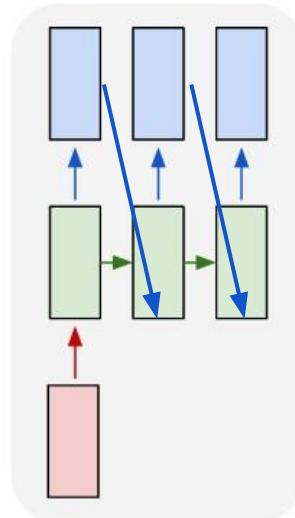
Sequence Models (Autoregressive)

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

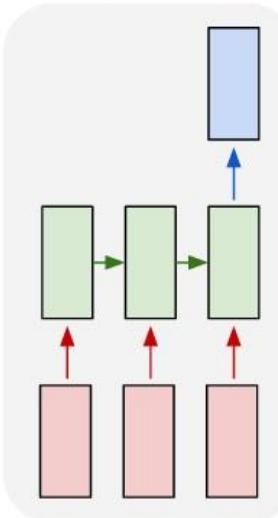
one to one



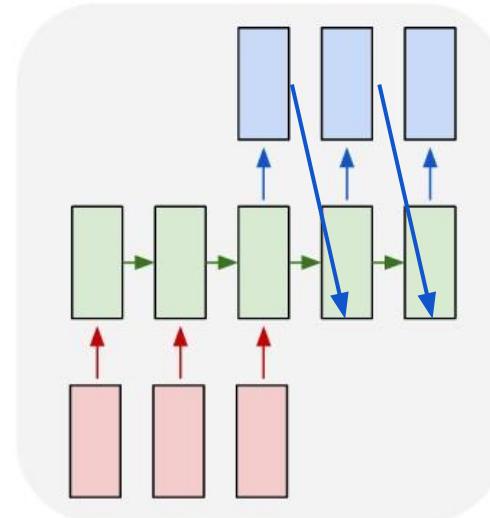
one to many



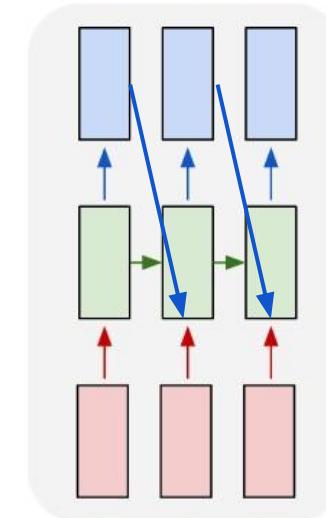
many to one



many to many



many to many



Conditional Pixel CNN

[van den Oord, Kalchbrenner, Vinyals, Espeholt, Graves, Kavukcuoglu, NIPS16]



Geyser



Hartebeest



Grey whale



Tiger

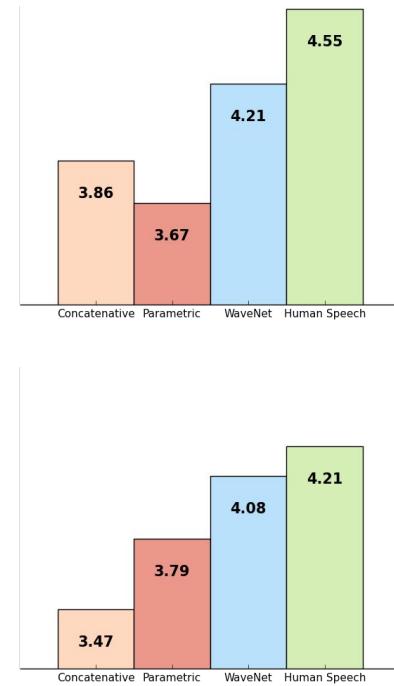
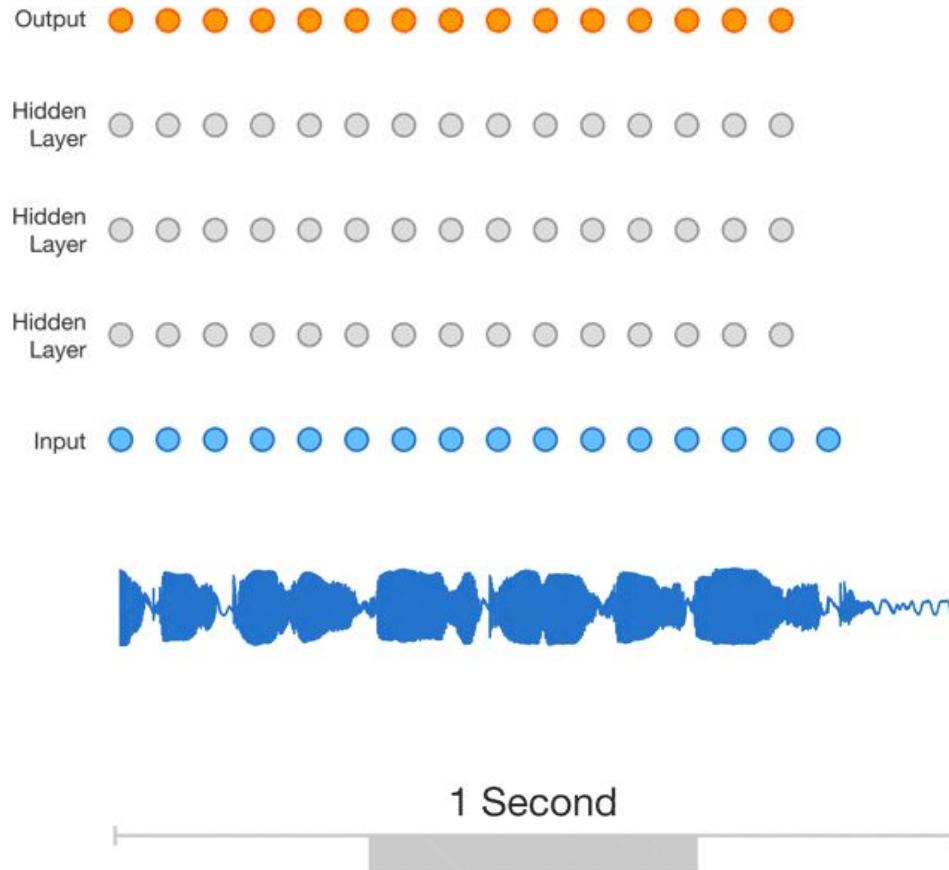


EntleBucher (dog)



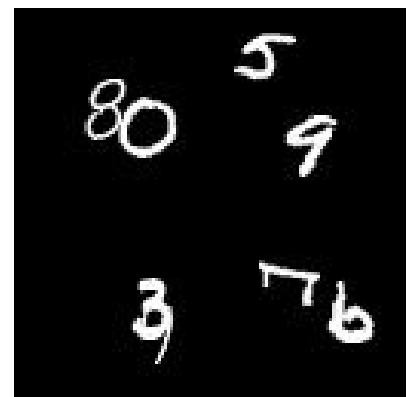
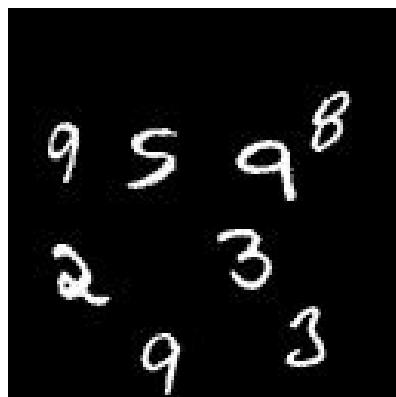
Yellow lady's slipper (flower)

WaveNets [van den Oord et al, 2016]



Video Pixel Network (VPN)

Model	Test
(Shi et al., 2015)	367.2
(Srivastava et al., 2015a)	341.2
(Brabandere et al., 2016)	285.2
(Patraucean et al., 2015)	179.8
Baseline model	110.1
VPN	87.6
Lower Bound	86.3



Neural Conversational Models

- Take movie subtitles (~900M words) or IT HelpDesk chats
- Predict the next dialog from history

i got to go .

no .

i get too emotional when i drink .

have another beer . i 've got to get up early .

no , you don 't . sit down .

i get too emotional when i drink .

will you have another beer ?

i 've got to go !

why ?

i got to get up early in the morning .

you 're drunk .

and emotional !

you got to go .

Conditional Language [Vinyals, Le, ICML DL WS 2015]

Human: what do you think about tesla ? Machine: he 's a good conductor .

Human: What do you think about messi ? Machine: he 's a great player .

Human: what do you think about cleopatra ? Machine: oh , she 's very regal .

Human: who is skywalker ? Machine: he is a hero .

Human: is sky blue or black ? Machine: blue .

Human: does a cat have a tail ? Machine: yes .

Human: does a cat have a wing ? Machine: no

Conditionals

Human: what do

Human: What do

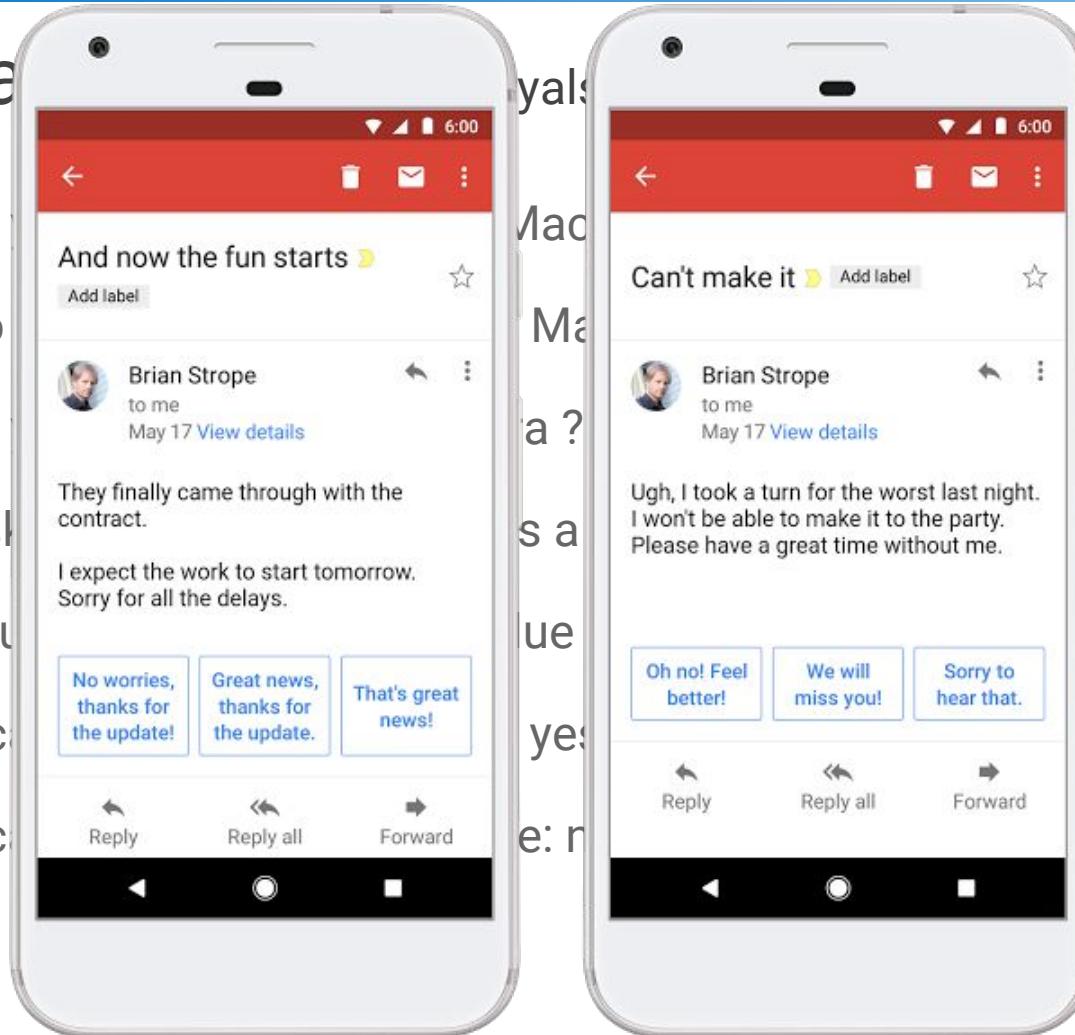
Human: what do

Human: who is sh

Human: is sky blu

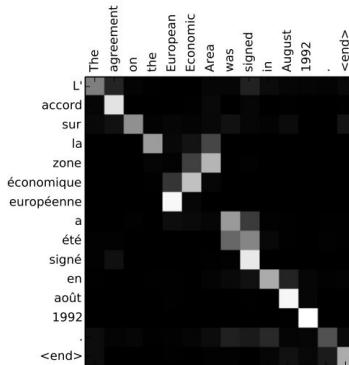
Human: does a c

Human: does a c

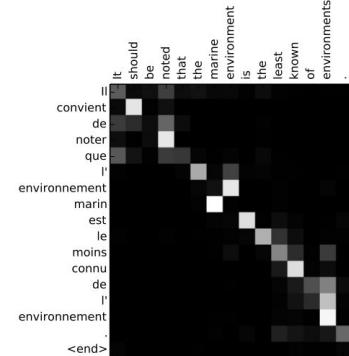


Part V: Beyond RNNs

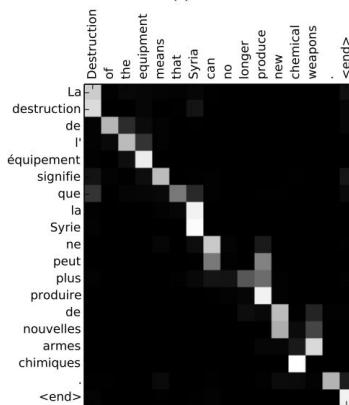
Attention (Bahdanau et al)



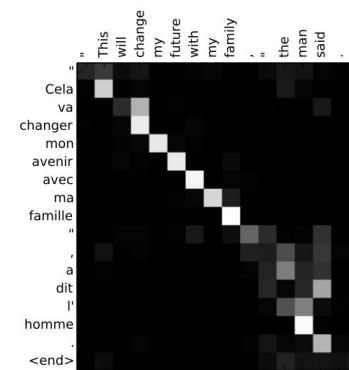
(a)



(b)



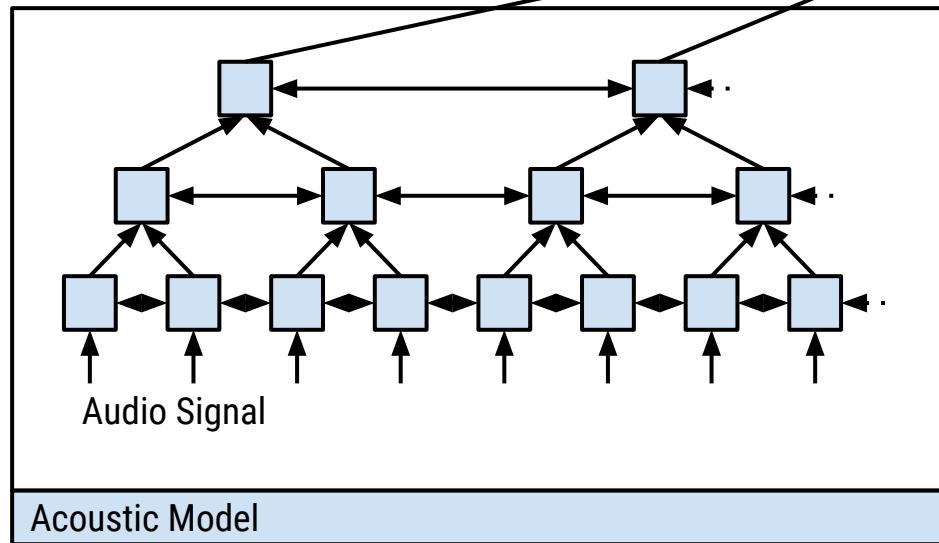
(c)



(d)

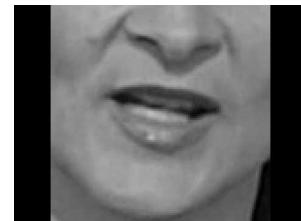
Listen Attend and Spell (LAS)

- Reducing time resolution with a pyramidal encoder



Lip Reading

Channel	Series name	# hours	# sent.
BBC 1 HD	News [†]	1,584	50,493
BBC 1 HD	Breakfast	1,997	29,862
BBC 1 HD	Newsnight	590	17,004
BBC 2 HD	World News	194	3,504
BBC 2 HD	Question Time	323	11,695
BBC 4 HD	World Today	272	5,558
All		4,960	118,116

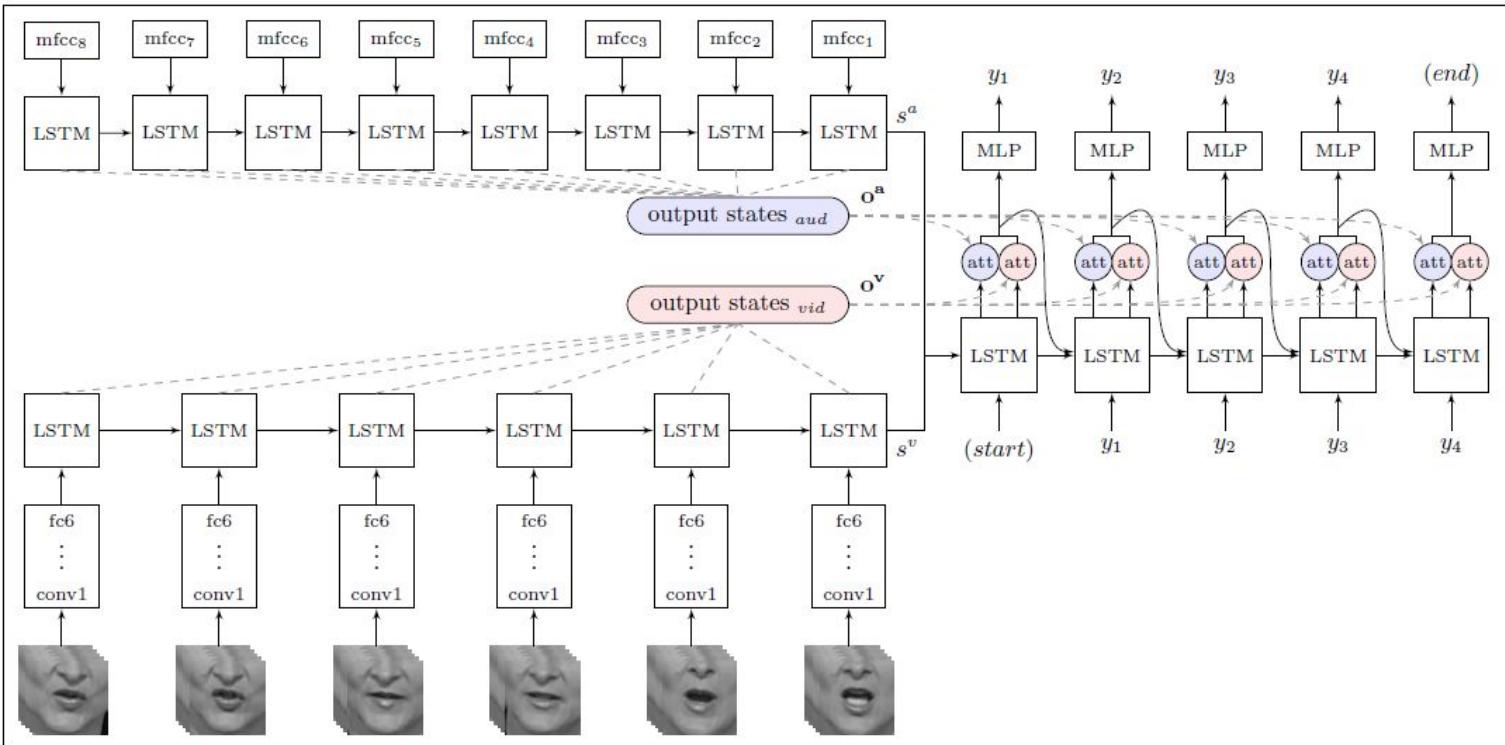


http://www.robots.ox.ac.uk/~vgg/data/lip_reading/

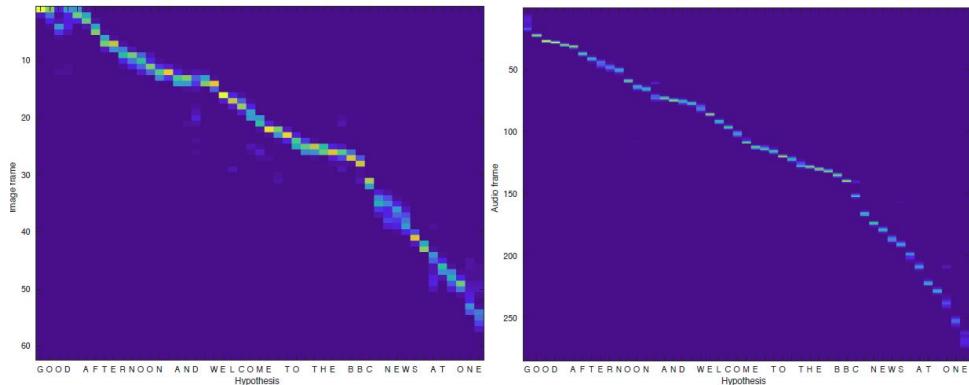
1. Chung, J., et al. "Lip reading sentences in the wild." *CVPR* (2017).
2. Assael, Y., et al. "Lipnet: Sentence-level lipreading." *arxiv* (2016).

Lip Reading

Separate
embedding
and attention
for audio and
visual
streams



Lip Reading



Method	SNR	CER	WER	BLEU [†]
Lips only				
Professional [‡]	-	58.7%	73.8%	23.8
WAS	-	59.9%	76.5%	35.6
WAS+CL	-	47.1%	61.1%	46.9
WAS+CL+SS	-	42.4%	58.1%	50.0
WAS+CL+SS+BS	-	39.5%	50.2%	54.9

Visual Attention (Xu et al)

Figure 3. Examples of attending to the correct object (*white* indicates the attended regions, *underlines* indicate the corresponding word)



A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.

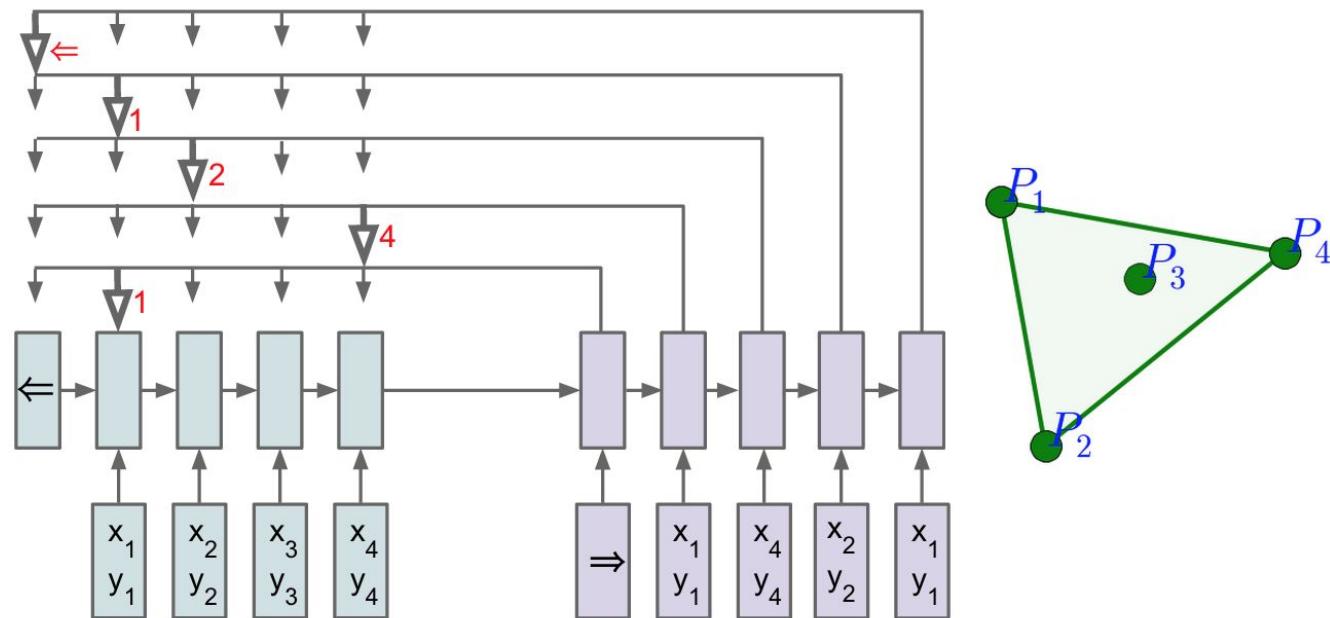


A group of people sitting on a boat in the water.



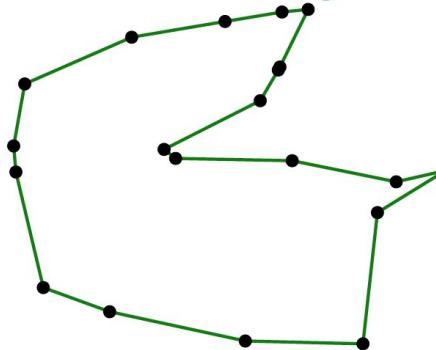
A giraffe standing in a forest with trees in the background.

Pointer Networks

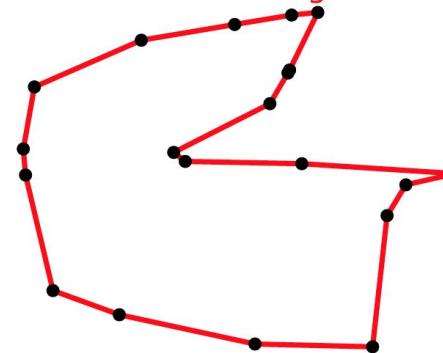


TSP: NP hard!

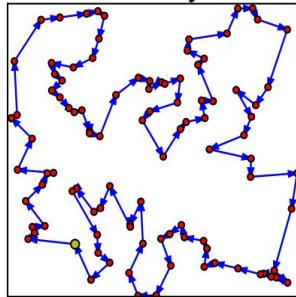
Ground Truth: tour length is 3.518



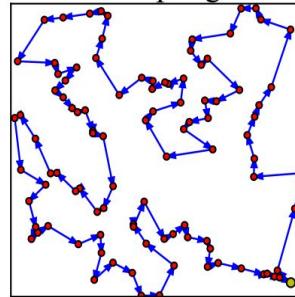
Predictions: tour length is 3.523



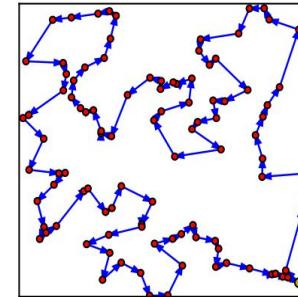
RL pretraining
-Greedy



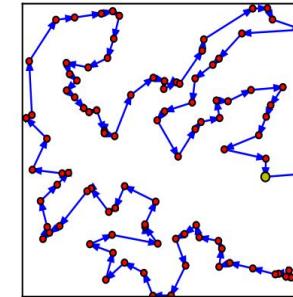
RL pretraining
-Sampling



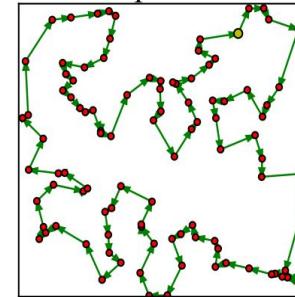
RL pretraining
-Active Search



Active Search

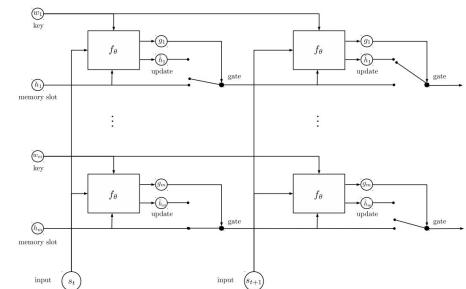
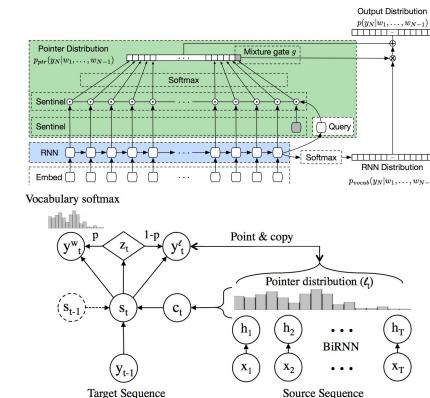


Optimal



Frontiers: Natural Language Understanding

- Good progress on dealing with rarely seen / unknown words
[Luong et al, 2015][Gulcehre et al, 2016][Merity et al, 2016][Wu et al, 2016]...
- Limited domain (e.g. BABI) tasks increasingly solved
[Kadlec et al, 2016][Dhingra et al, 2016][Trischler et al, 2016][Henaff et al, 2016]...
- We can't generate long / coherent text
[Liu et al, 2018]
- We can't summarize (or write :)) a book



Frontiers: Complex Decision Making in Games



<https://github.com/deepmind/pysc2>

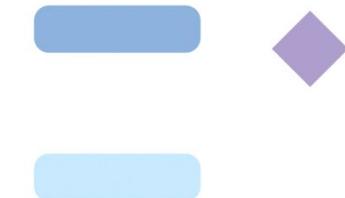
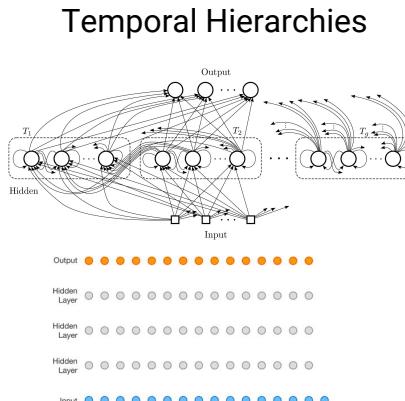
MiniGames

AGENT	METRIC	MOVE TO BEACON	COLLECT MINERAL SHARDS	FIND AND DEFEAT ZERGLINGS	DEFEAT ROACHES	DEFEAT ZERGLINGS AND BANELINGS	COLLECT MINERALS AND GAS	BUILD MARINES
RANDOM POLICY	MEAN	1	17	4	1	23	12	< 1
	MAX	6	35	19	46	118	750	5
RANDOM SEARCH	MEAN	25	32	21	51	55	2318	8
	MAX	29	57	33	241	159	3940	46
DEEPMIND HUMAN PLAYER	MEAN	26	133	46	41	729	6880	138
	MAX	28	142	49	81	757	6952	142
STARCRAFT GRANDMASTER	MEAN	28	177	61	215	727	7566	133
	MAX	28	179	61	363	848	7566	133
ATARI-NET	BEST MEAN	25	96	49	101	81	3356	< 1
	MAX	33	131	59	351	352	3505	20
FULLYCONV	BEST MEAN	26	103	45	100	62	3978	3
	MAX	45	134	56	355	251	4130	42
FULLYCONV LSTM	BEST MEAN	26	104	44	98	96	3351	6
	MAX	35	137	57	373	444	3995	62

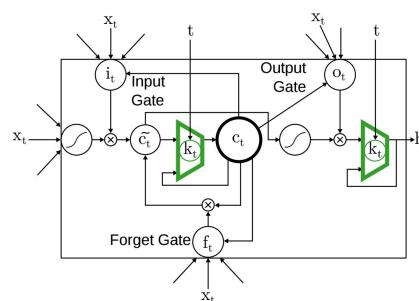


Common Challenge: Long Term Dependencies

- Thousands / Millions of steps
 - Recursive programs
 - Learning to Learn
 - Complex Decision Making
 - NLU



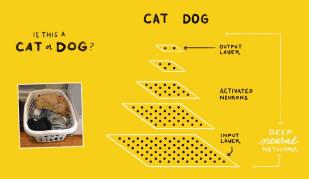
Jaderberg et al, 2016



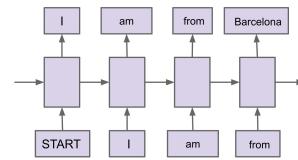
Neil et al, 2016

The Deep Learning Toolbox

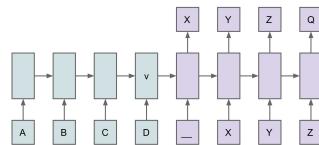
Feed forward models



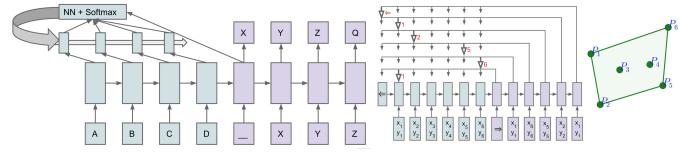
Sequence Prediction



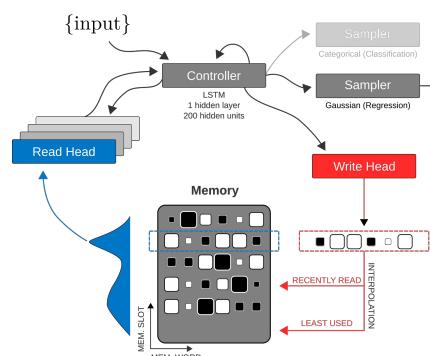
Seq2Seq



Attention/Pointers

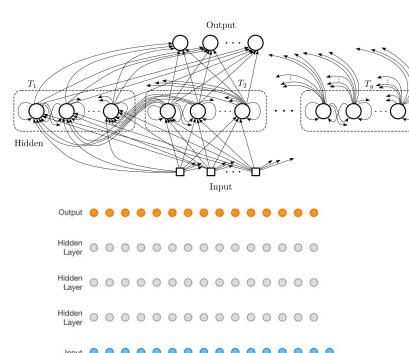


Read/Write memories

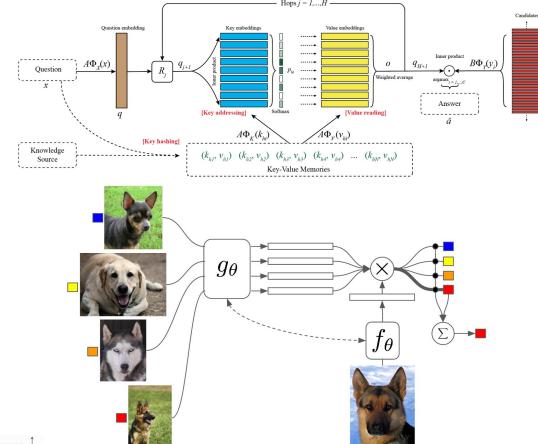


$$\begin{aligned} \text{COSINE DISTANCE} \\ K(\mathbf{k}_t, \mathbf{M}_t(i)) = \frac{\mathbf{k}_t \cdot \mathbf{M}_t(i)}{\|\mathbf{k}_t\| \|\mathbf{M}_t(i)\|} \\ \text{READ WEIGHTS} \\ w_t^r(i) \leftarrow \frac{\exp(K(\mathbf{k}_t, \mathbf{M}_t(i)))}{\sum_j \exp(K(\mathbf{k}_t, \mathbf{M}_t(j)))} \\ \text{READ VECTOR} \\ \mathbf{r}_t \leftarrow \sum_i w_t^r(i) \mathbf{M}_t(i) \\ \text{WRITE WEIGHTS} \\ \mathbf{w}_t^w \leftarrow \sigma(\alpha) \mathbf{w}_{t-1}^r + (1 - \sigma(\alpha)) \mathbf{w}_{t-1}^{lu} \\ \text{USAGE WEIGHTS} \\ \mathbf{w}_t^{lu}(i) \leftarrow \gamma \mathbf{w}_{t-1}^r + \mathbf{w}_t^r + \mathbf{w}_t^w \\ \text{LEAST-USED WEIGHTS} \\ w_t^{lu}(i) = \begin{cases} 0 & \text{if } w_t^r(i) > m(\mathbf{w}_t^{lu}, n) \\ 1 & \text{if } w_t^r(i) \leq m(\mathbf{w}_t^{lu}, n) \end{cases} \\ \text{MEMORY UPDATING} \\ \mathbf{M}_t(i) \leftarrow \mathbf{M}_{t-1}(i) + w_t^r(i) \mathbf{k}_t, \forall i \end{aligned}$$

Temporal Hierarchies



Key,Value memories



Recurrent Architectures



Figure credits: Jeff Dean, Chris Olah, Santoro et al 2016, Koutella 2014, van den Oord et al 2016, Miller et al 2016, Vinyals et al 2016

Summary

- Sequences are important, everything is a sequence
- Recurrent Nets are a way forward, but they pose some challenges
 - Vanishing/Exploding gradients
 - Expensive to compute
- Conditional models are very powerful, and many applications “just worked”
when seq2seq and similar were trained on large amounts of data
- Attention/memory, longer term memory... many challenges remain!

Questions?

@OriolVinyalsML // vinyals@google.com

Extra Resources:

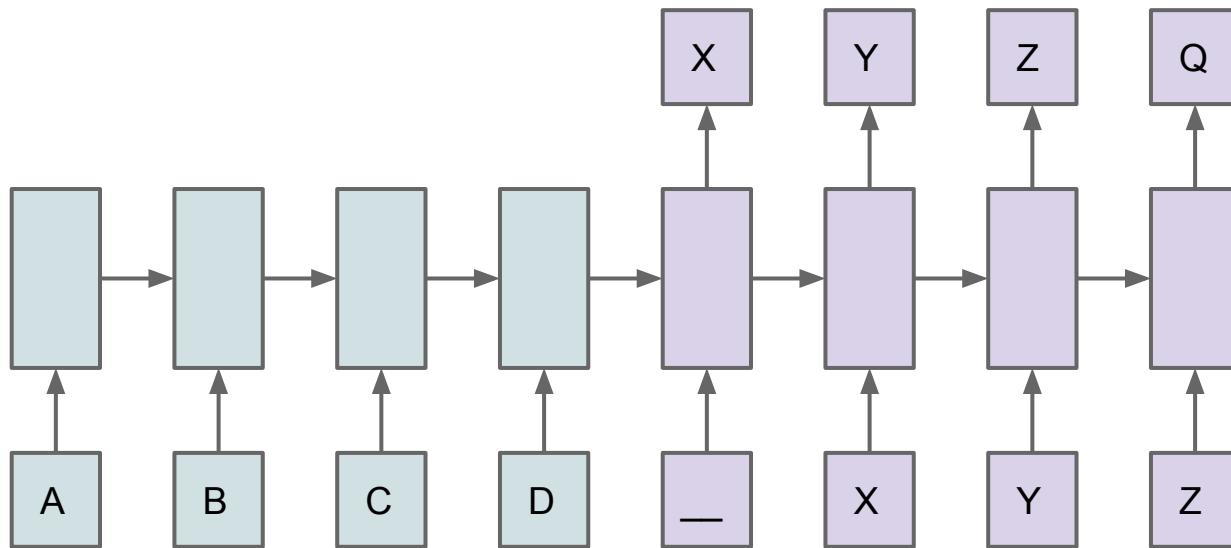
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Vinyals & Jaitly ICML 2017 Tutorial (Seq2Seq)

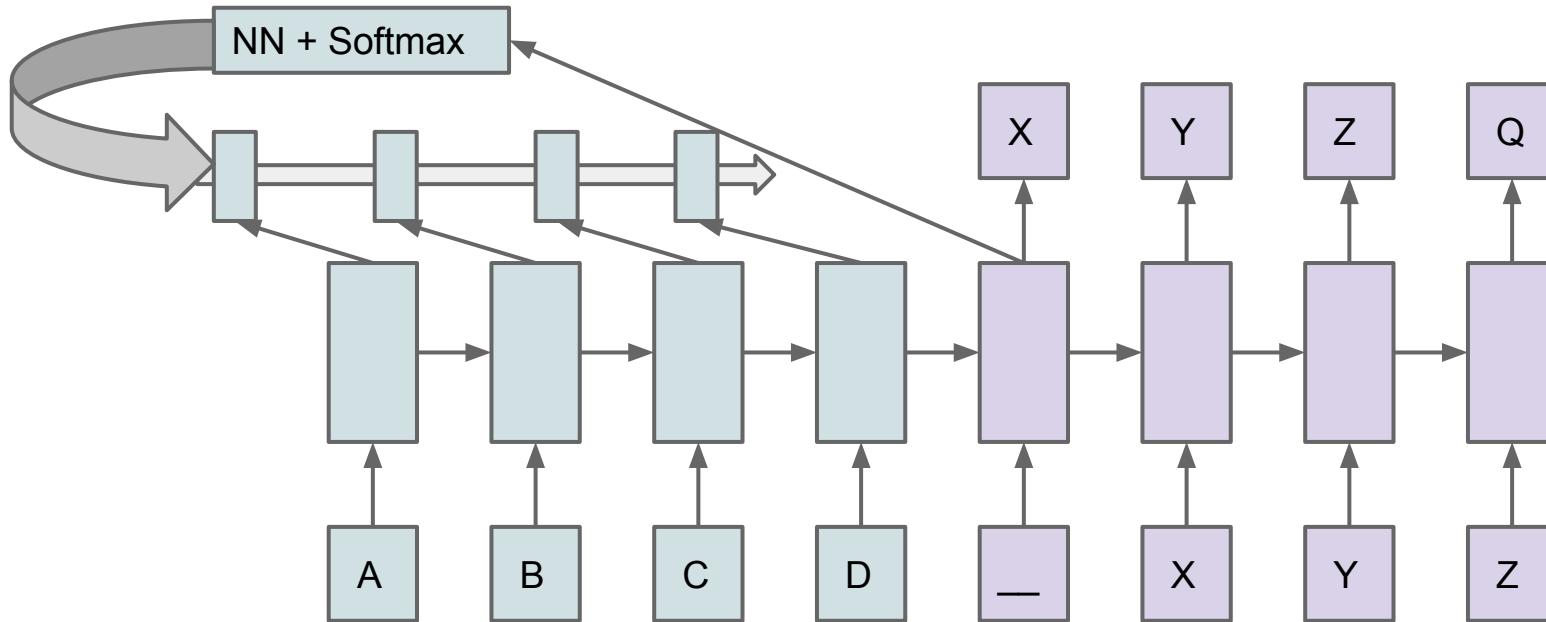
Next lecture...

Beyond image recognition, end-to-end learning, embeddings - Raia Hadsell

Attention



Attention



Position Based Attention

Inputs: "I am a cat."

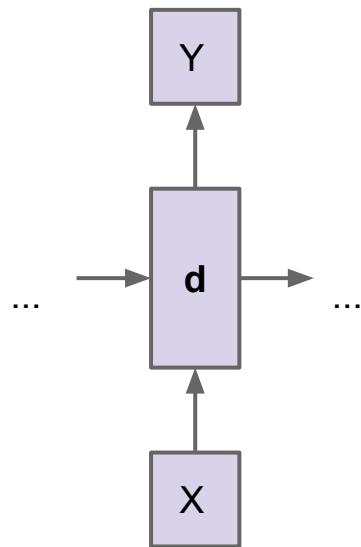
Input RNN states: $\mathbf{e}_1 \mathbf{e}_2 \mathbf{e}_3 \mathbf{e}_4$

Decoder RNN state: \mathbf{d}

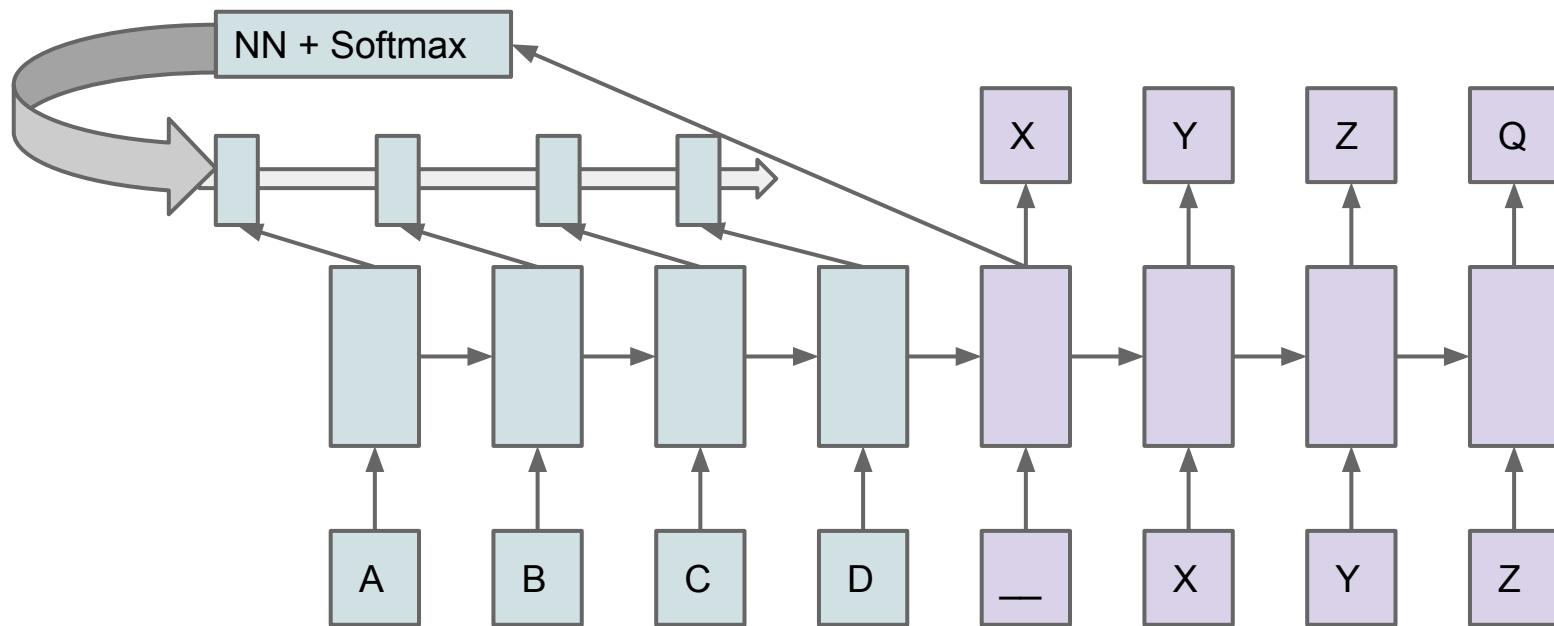
Map \mathbf{d} to $\mathbf{a} = [0.1 \ 0.9 \ 0.0 \ 0.0]$ E.g. $\mathbf{a} = s(\mathbf{Wd})$

Compute: $\mathbf{d}' = 0.1*\mathbf{e}_1 + 0.9*\mathbf{e}_2 + 0*\mathbf{e}_3 + 0*\mathbf{e}_4$

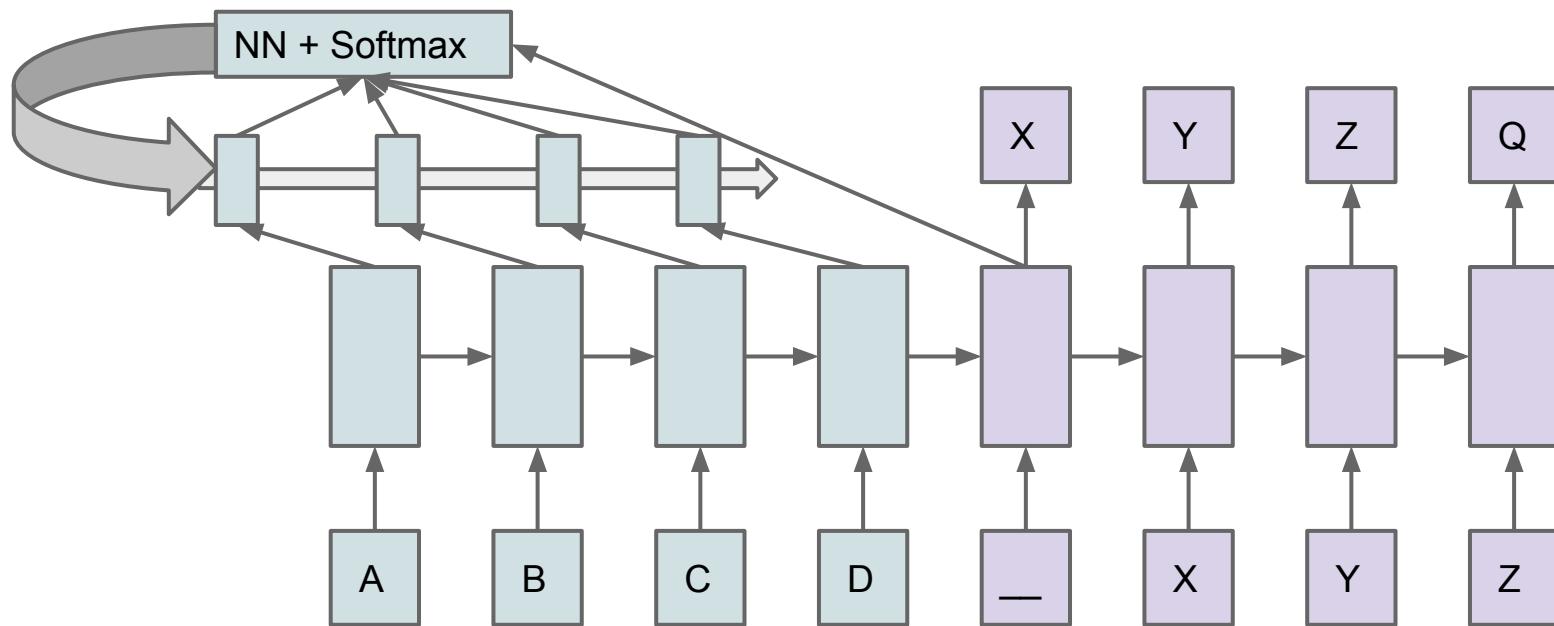
Set $\mathbf{d} = [\mathbf{d} \ \mathbf{d}']$



Position Based Attention



Content Based Attention



Content Based Attention

Inputs: "I am a cat."

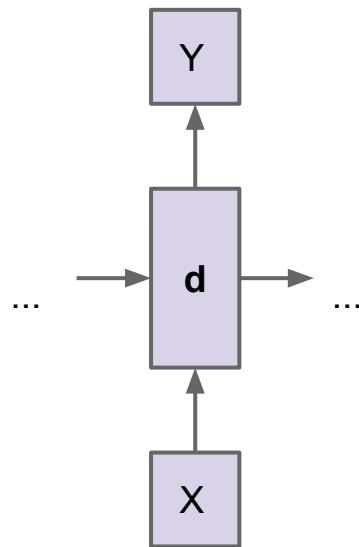
Input RNN states: $\mathbf{e}_1 \mathbf{e}_2 \mathbf{e}_3 \mathbf{e}_4$

Decoder RNN state: \mathbf{d}

Compute $\mathbf{d}^T \mathbf{e}_i : [0.0 \ 0.05 \ 0.9 \ 0.05]$

Compute: $\mathbf{d}' = 0*\mathbf{e}_1 + 0.05*\mathbf{e}_2 + 0.9*\mathbf{e}_3 + 0.05*\mathbf{e}_4$

Set $\mathbf{d} = [\mathbf{d} \ \mathbf{d}']$



Content Based Attention

Attention [Bahdanau, Cho and Bengio, 2014]

$$u_j = v^T \tanh(W_1 e_j + W_2 d) \quad j \in (1, \dots, n)$$

$$a_j = \text{softmax}(u_j) \quad j \in (1, \dots, n)$$

$$d' = \sum_{j=1}^n a_j e_j$$

Neural Turing Machine & Memory Networks:

Given a **key** vector \mathbf{k} and a **strength** scalar s emitted by the controller, get a weighting (w_1, \dots, w_N) over memory locations where

$$w_i = \frac{e^{\hat{w}_i}}{\sum_j e^{\hat{w}_j}} \quad \hat{w}_i = \log(1 + e^s) C(\mathbf{k}, \mathbf{m}_i)$$
$$C(\mathbf{k}, \mathbf{m}_i) = \frac{\mathbf{k} \cdot \mathbf{m}_i}{\|\mathbf{k}\| \|\mathbf{m}_i\|}$$

and $C(\mathbf{k}, \mathbf{m}_i)$ is the **cosine similarity** between the key and the **word** \mathbf{m}_i at memory location i

Frontiers: Learning Programs

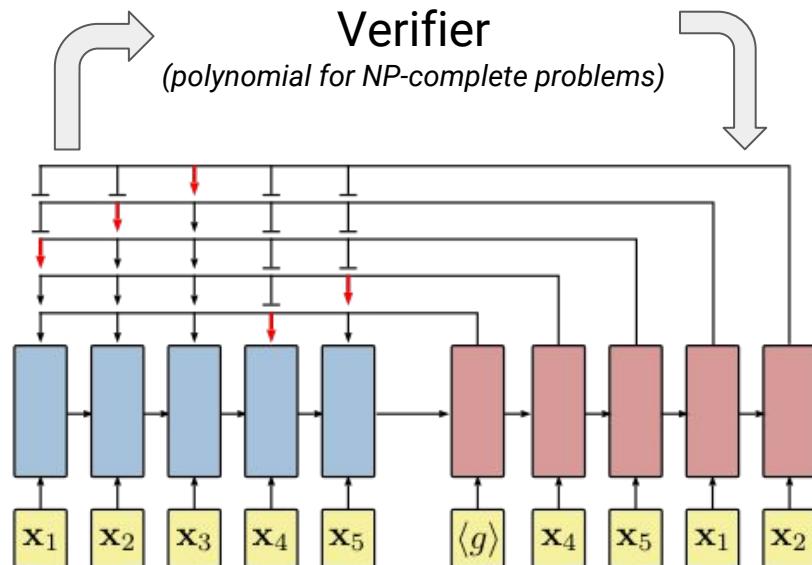
LOTS of related work: [Zaremba et al, 2014][Reed et al, 2015][Neelakantan et al, 2015][Vinyals et al, 2015]...

RL-PtrNets [Bello, Pham, et al 2016]

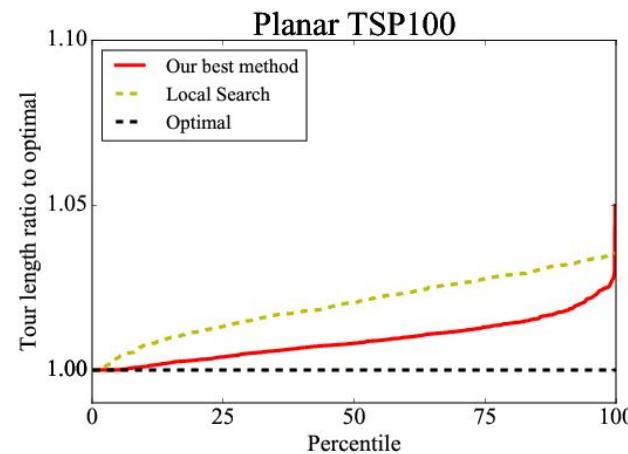
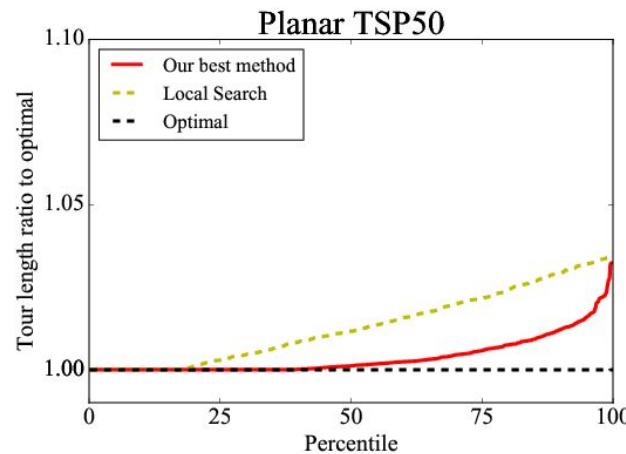
Actor-critic training using the expected combinatorial optimization objective

Different configurations:

- *Train* on training set
- *Search* at inference time
- *Refine* parameters at inference time



Frontiers: Learning Programs



Average tour lengths on TSP100 ($\sim 10^{157}$ solutions)

Christofides: 9.18 | Local search: 7.99 | RL-PtrNets: 7.83 | Optimal: 7.77